


Syscall Tracing für libiotrace

Gerrit Klein - HPC Projekt (Prof. Keller)



Gliederung

Gliederung

1. Vorstellung libiotrace

Gliederung

1. Vorstellung libiotrace
2. Filename Resolution

Gliederung

1. Vorstellung libiotrace
2. Filename Resolution
3. Motivation - Syscall Tracing

Gliederung

1. Vorstellung libiotrace
2. Filename Resolution
3. Motivation - Syscall Tracing
4. `ministrace`

Gliederung

1. Vorstellung libiotrace
2. Filename Resolution
3. Motivation - Syscall Tracing
4. `ministrace`
5. Todo: Integration

Vorstellung libiotrace

Motivation

Motivation

- Problem: File-I/O = Bottleneck

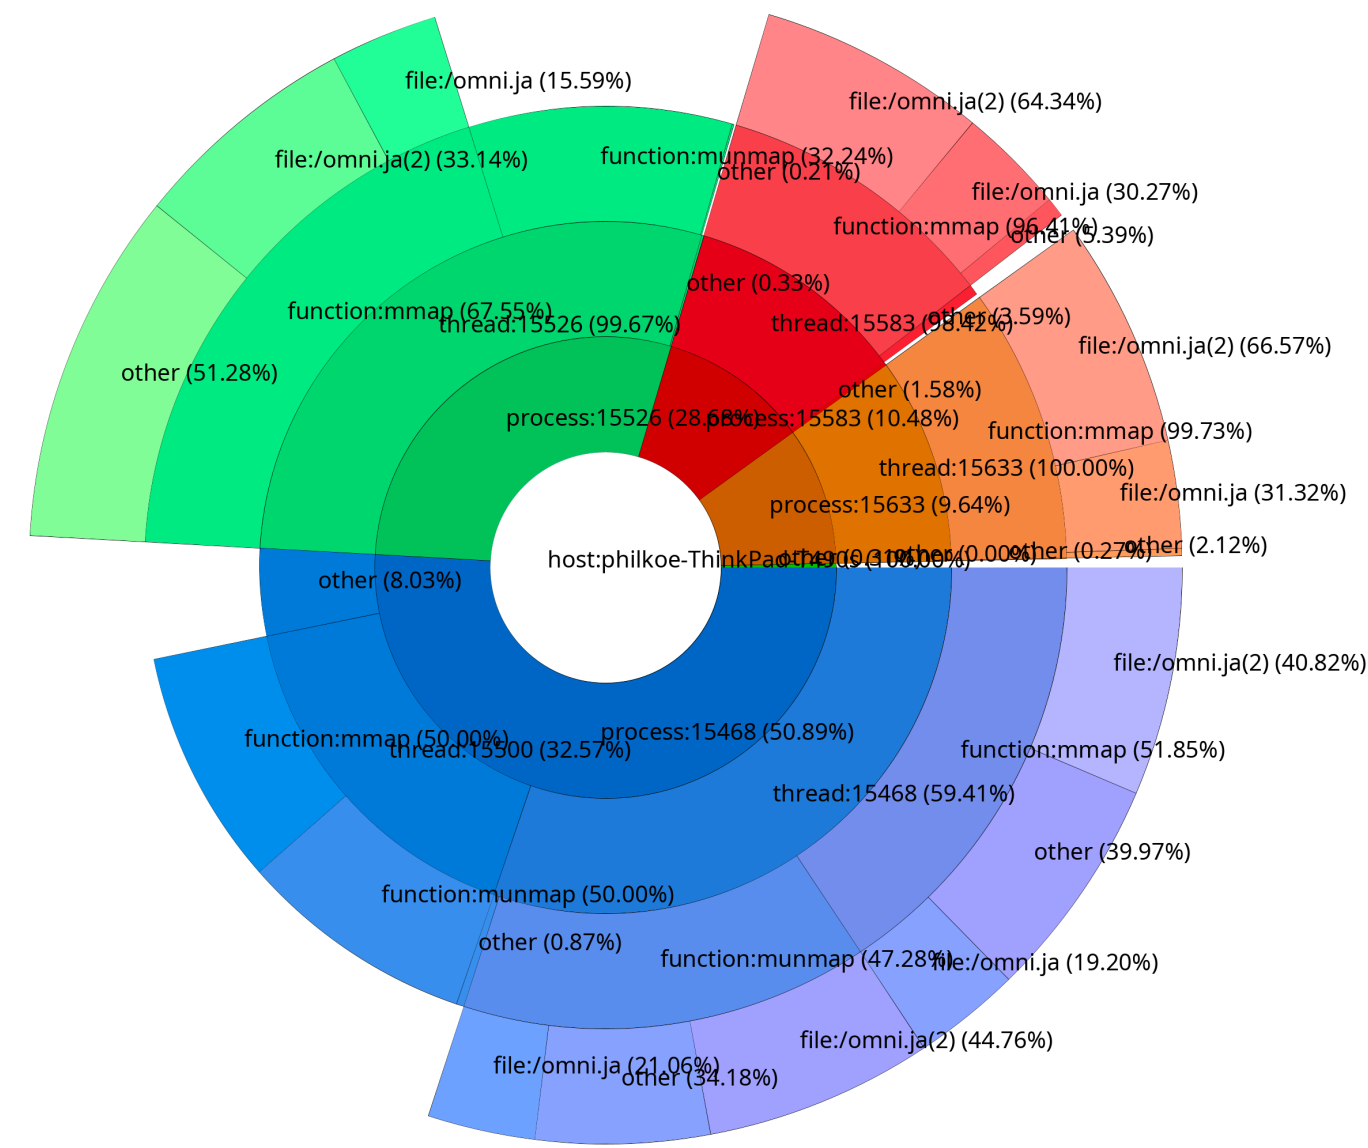
Motivation

- Problem: File-I/O = Bottleneck
 - **Tracing** (POSIX-/ MPI-I/O): libiotrace

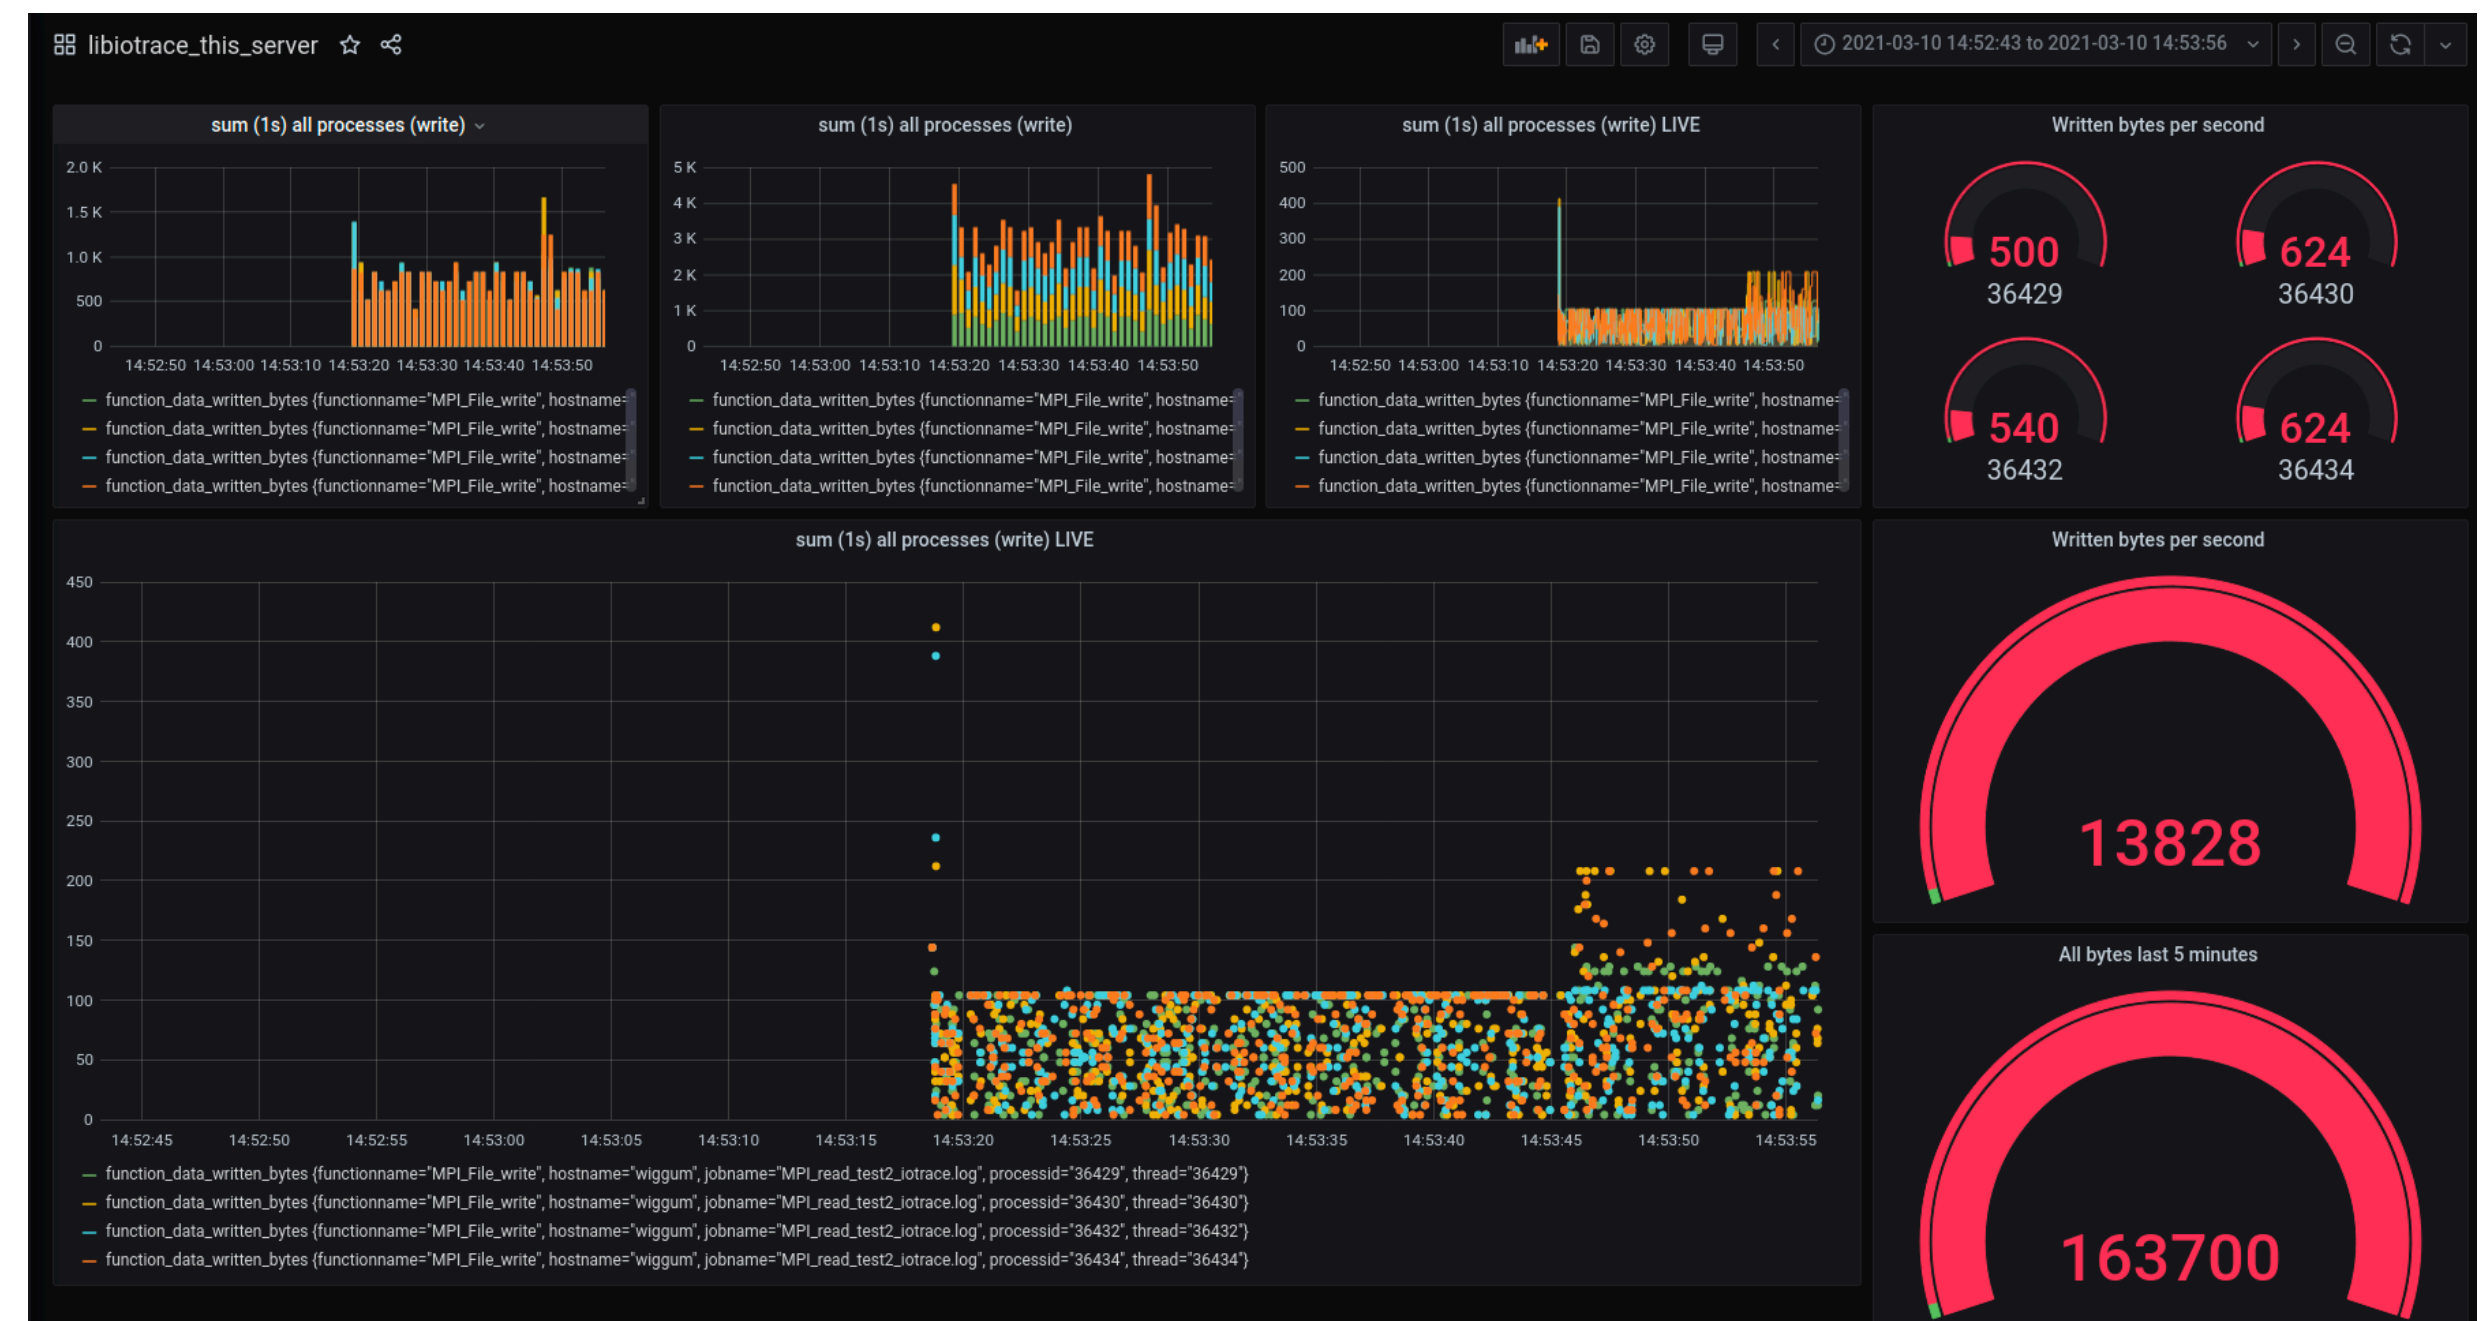
Motivation

- Problem: File-I/O = Bottleneck
- **Tracing** (POSIX-/ MPI-I/O): libiotrace
- **Auswertung:** Java Tool (Post mortem) / Grafana (Live Tracing)

Read and written bytes for I/O (with components which read or write more than 3.0%)



Quelle: GitHub

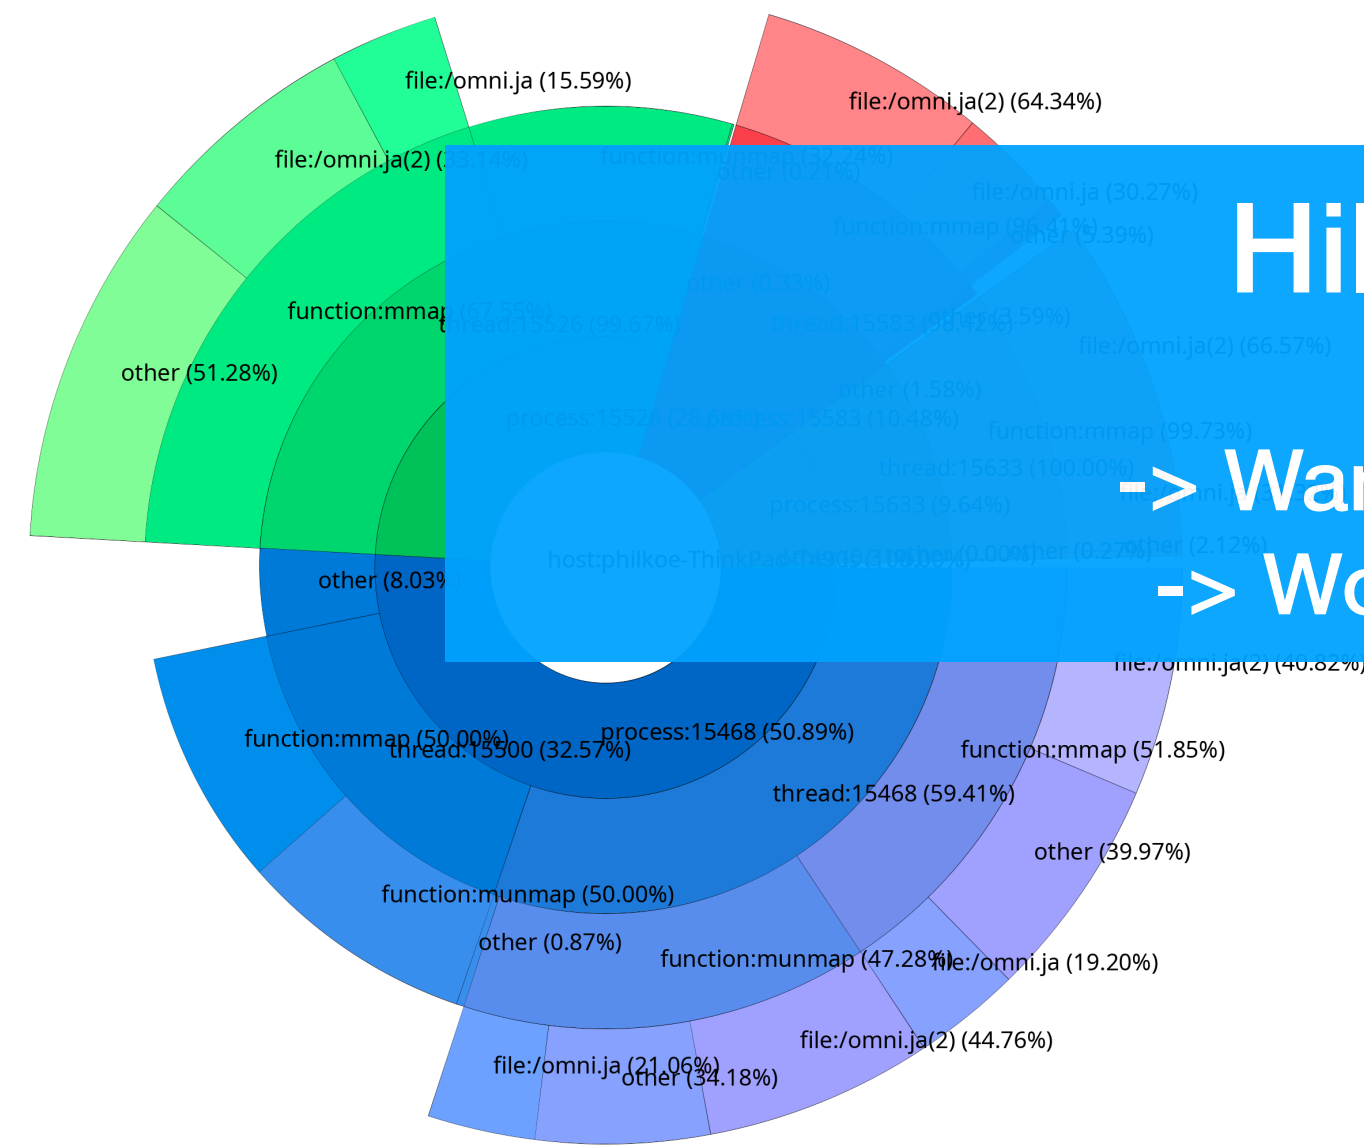


Quelle: GitHub

Motivation

- Problem: File-I/O = Bottleneck
 - **Tracing** (POSIX-/ MPI-I/O): libiotrace
 - **Auswertung:** Java Tool (Post mortem) / Grafana (Live Tracing)

Read and written bytes for I/O (with components which read or write more than 3.0%)



Quelle: GitHub



Quelle: GitHub

Funktionsweise

→ ~ **LD_PRELOAD**=path/to/libiotrace_shared.so ./path/to/exec

Funktionsweise

→ ~ **LD_PRELOAD**=path/to/libiotrace_shared.so ./path/to/exec

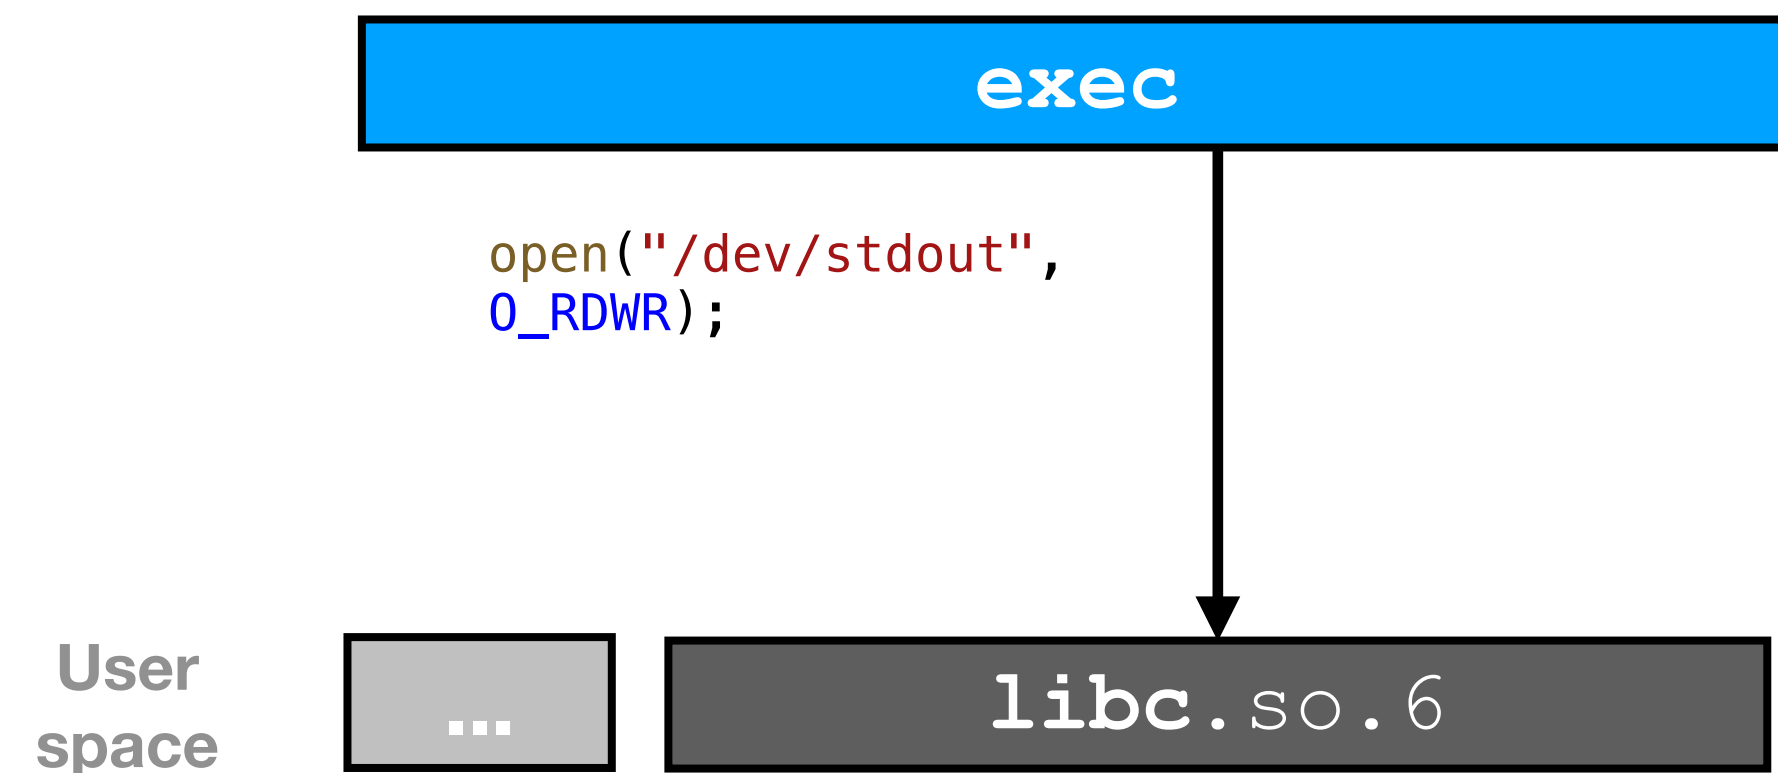


exec

User
space

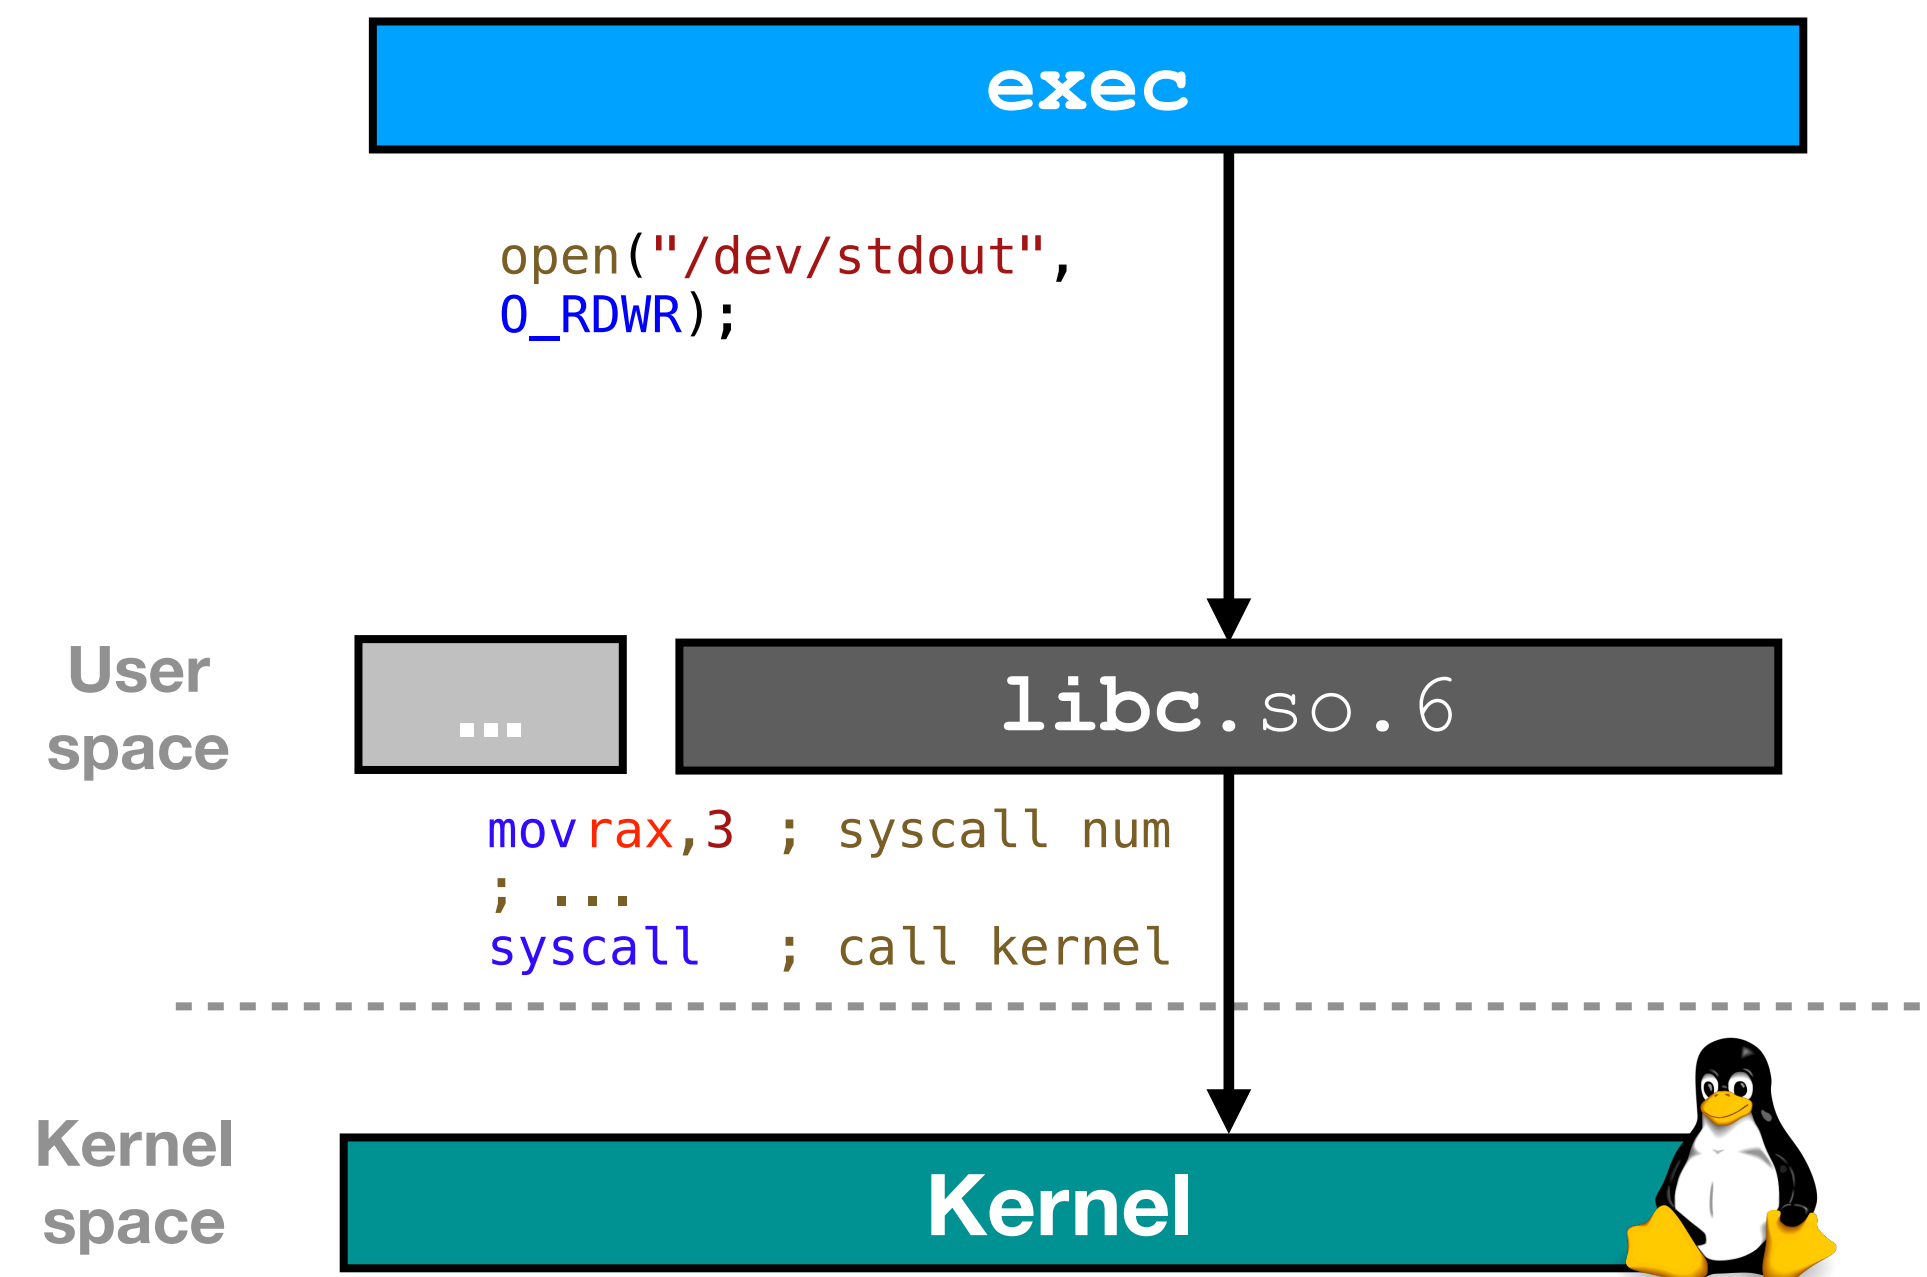
Funktionsweise

→ ~ **LD_PRELOAD**=path/to/libiotrace_shared.so ./path/to/exec



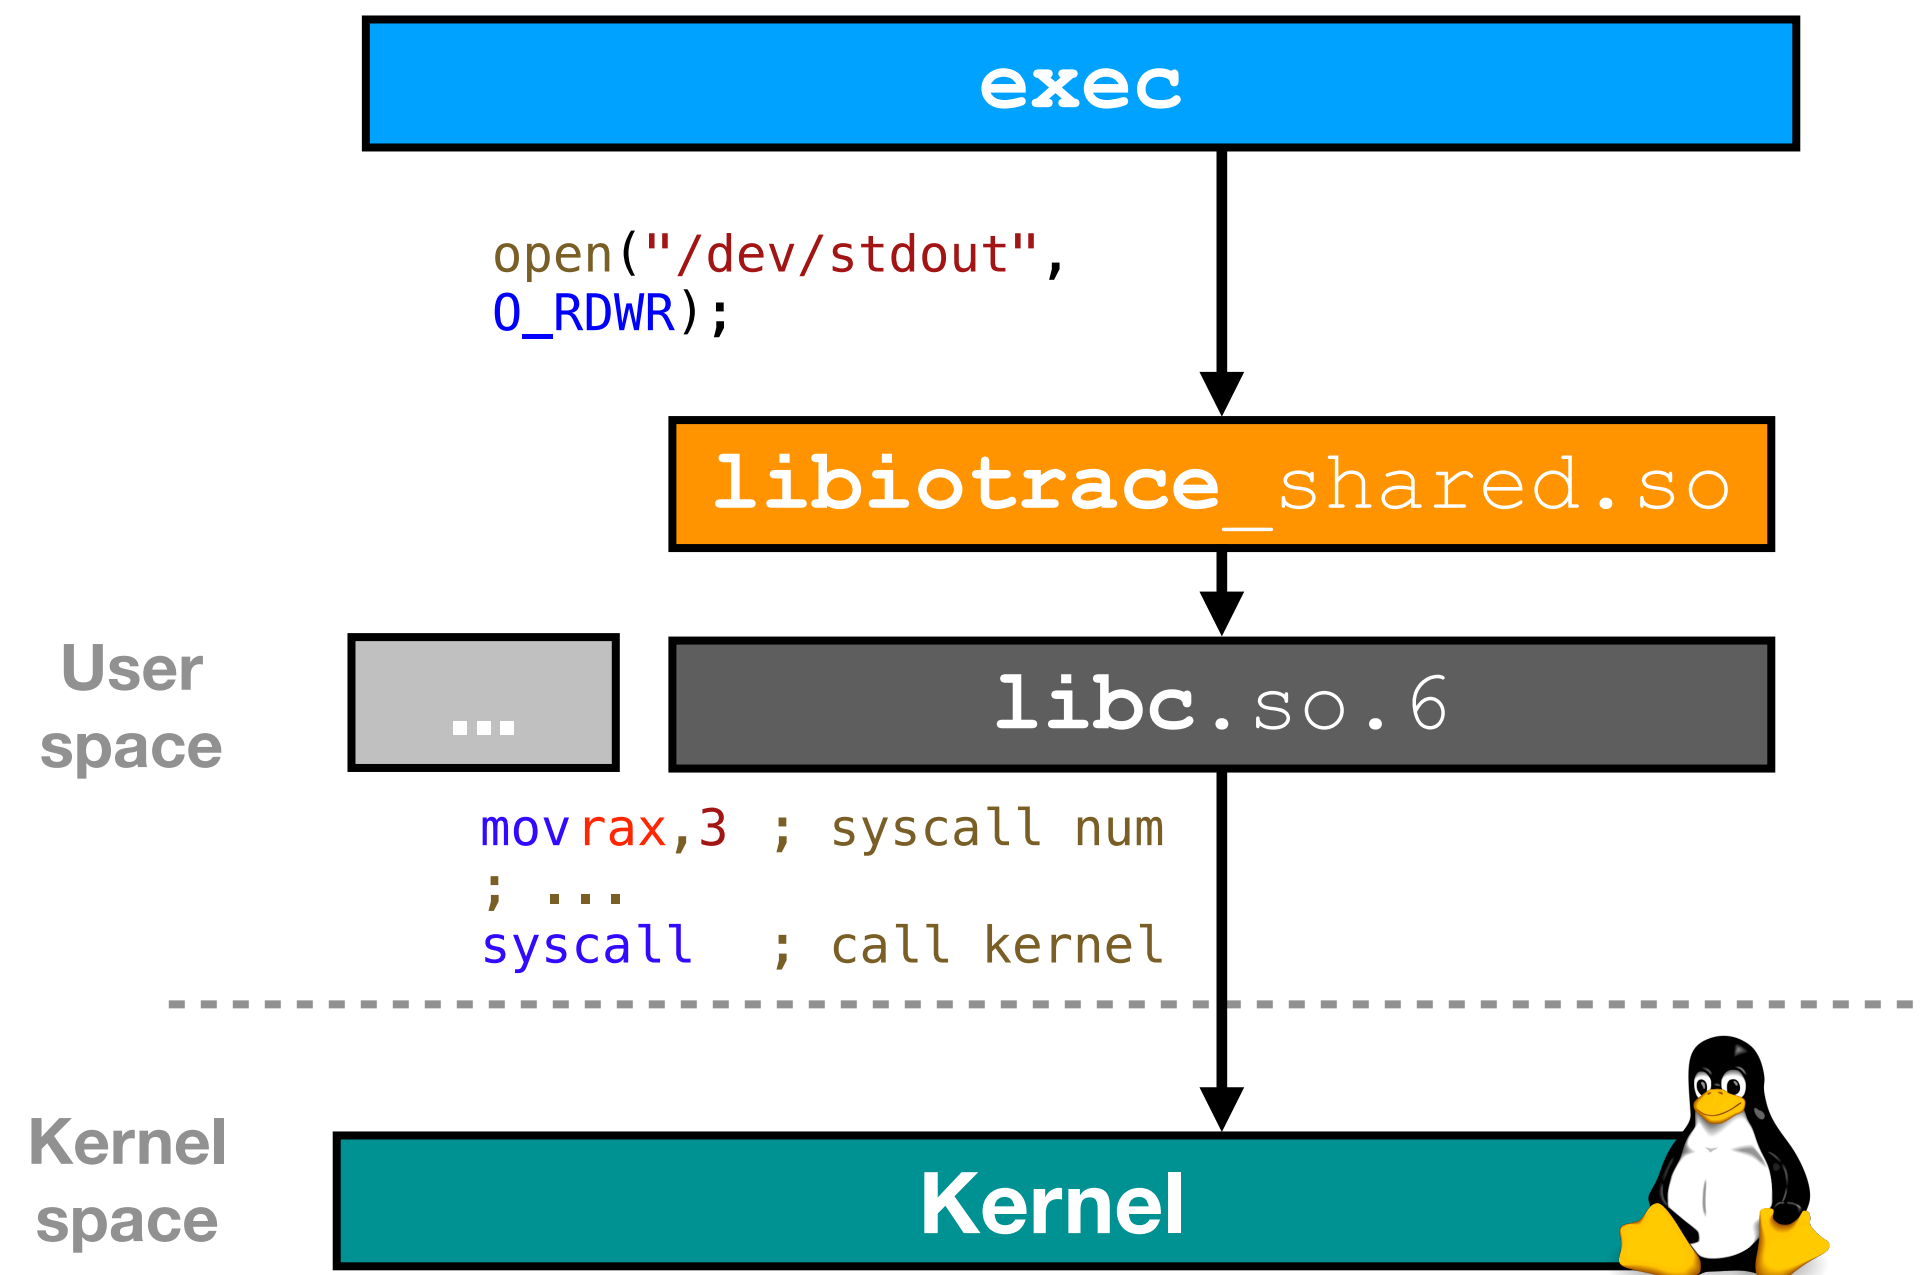
Funktionsweise

→ ~ **LD_PRELOAD**=path/to/libiotrace_shared.so ./path/to/exec



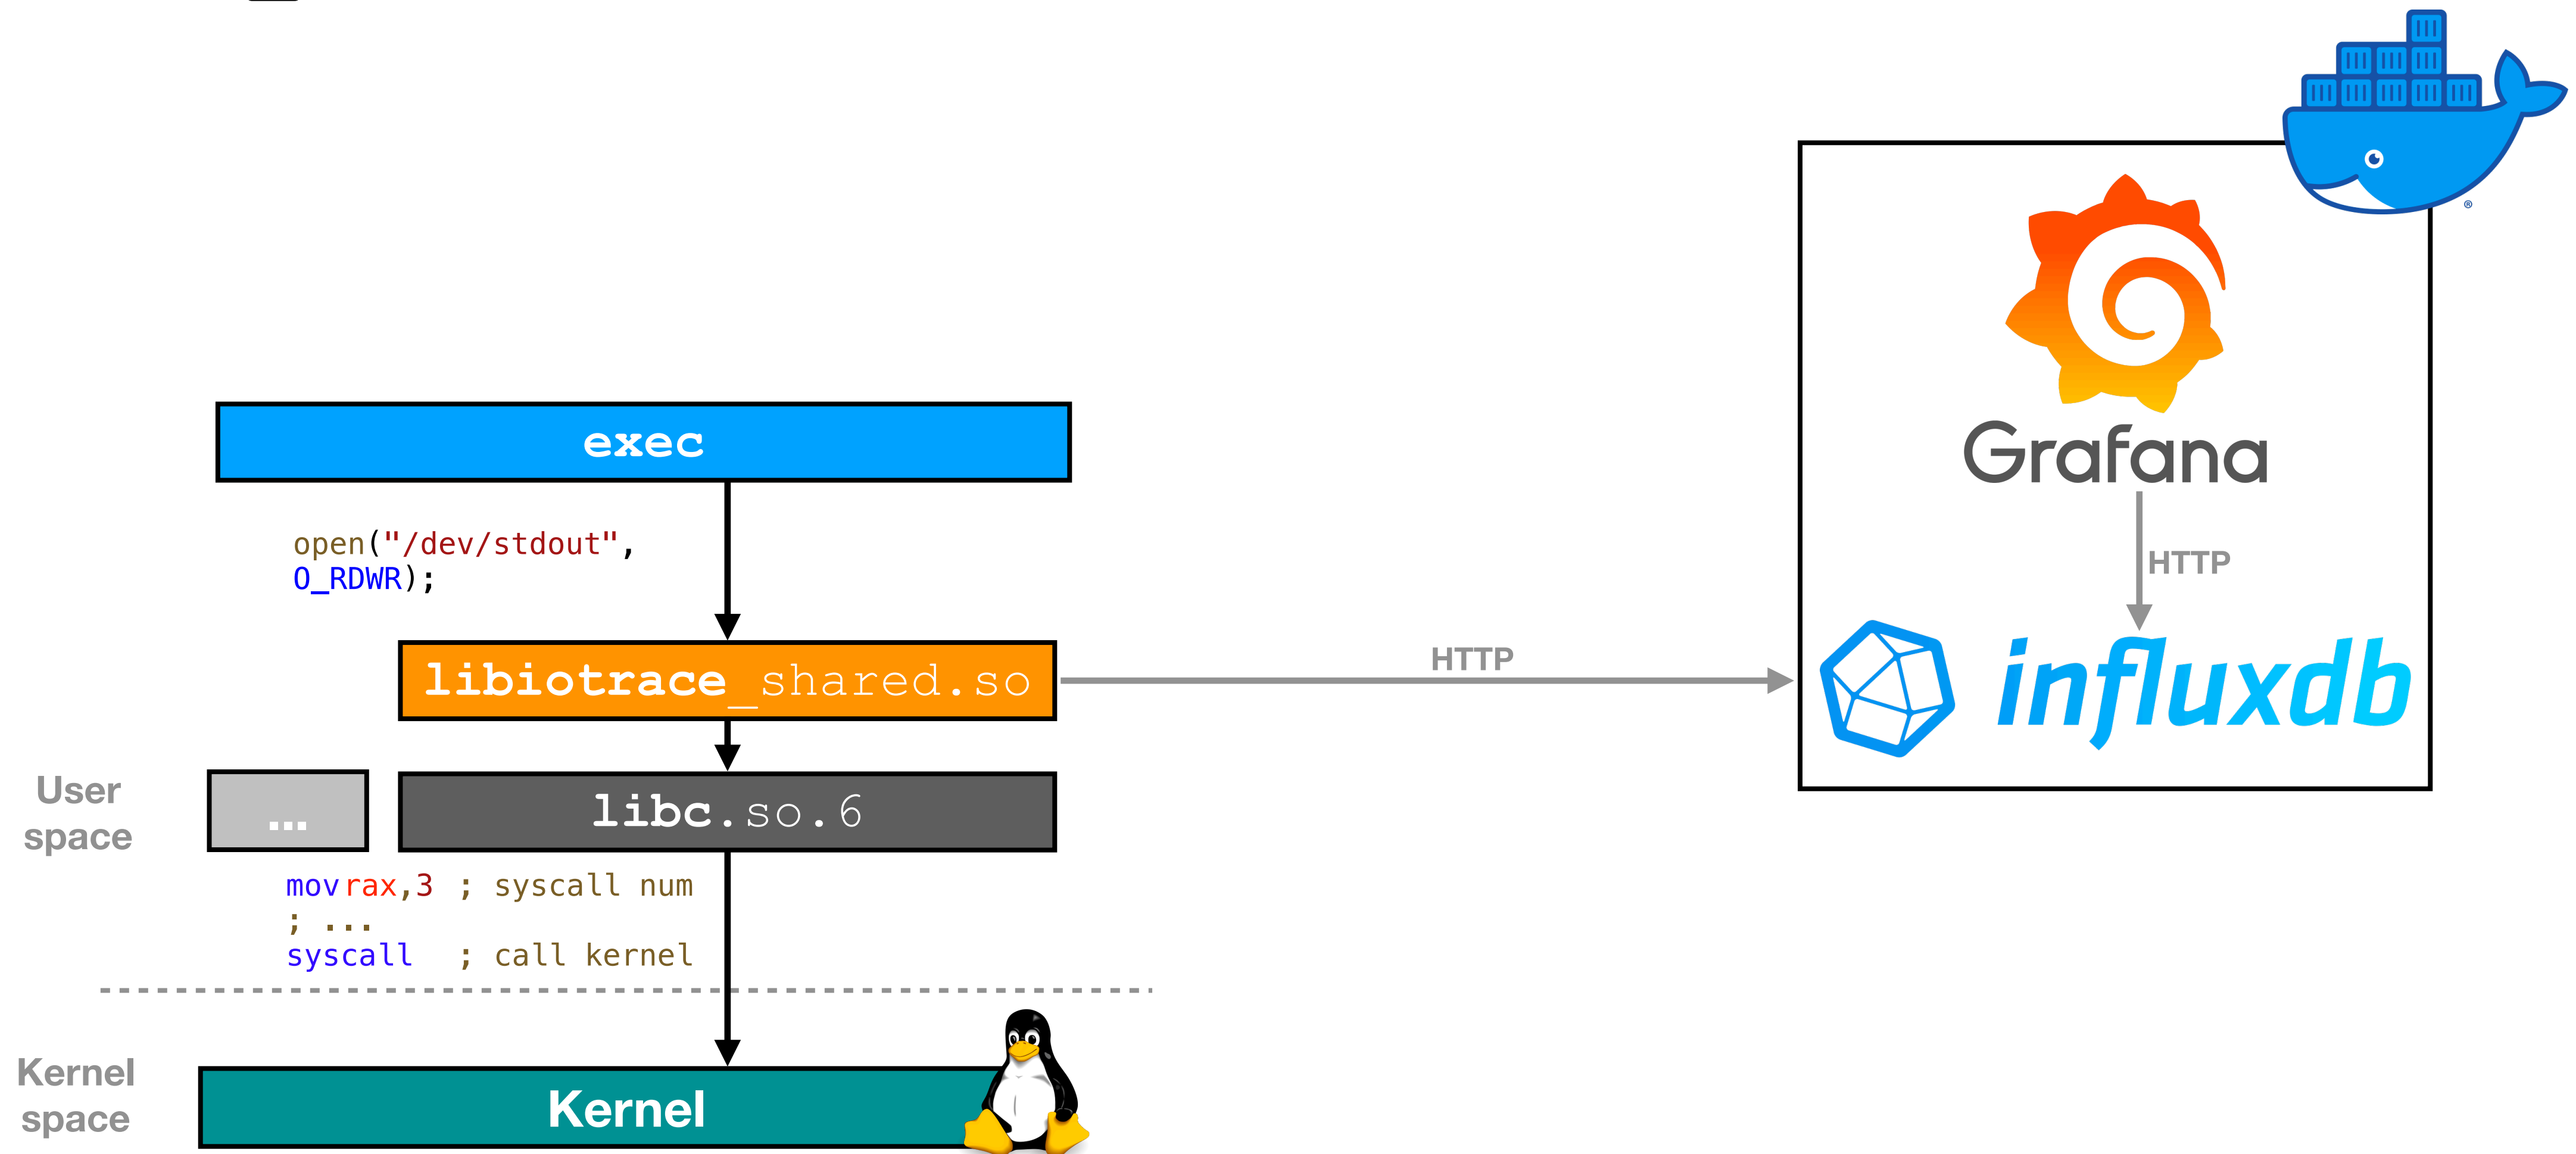
Funktionsweise

→ ~ **LD_PRELOAD**=path/to/libiotrace_shared.so ./path/to/exec



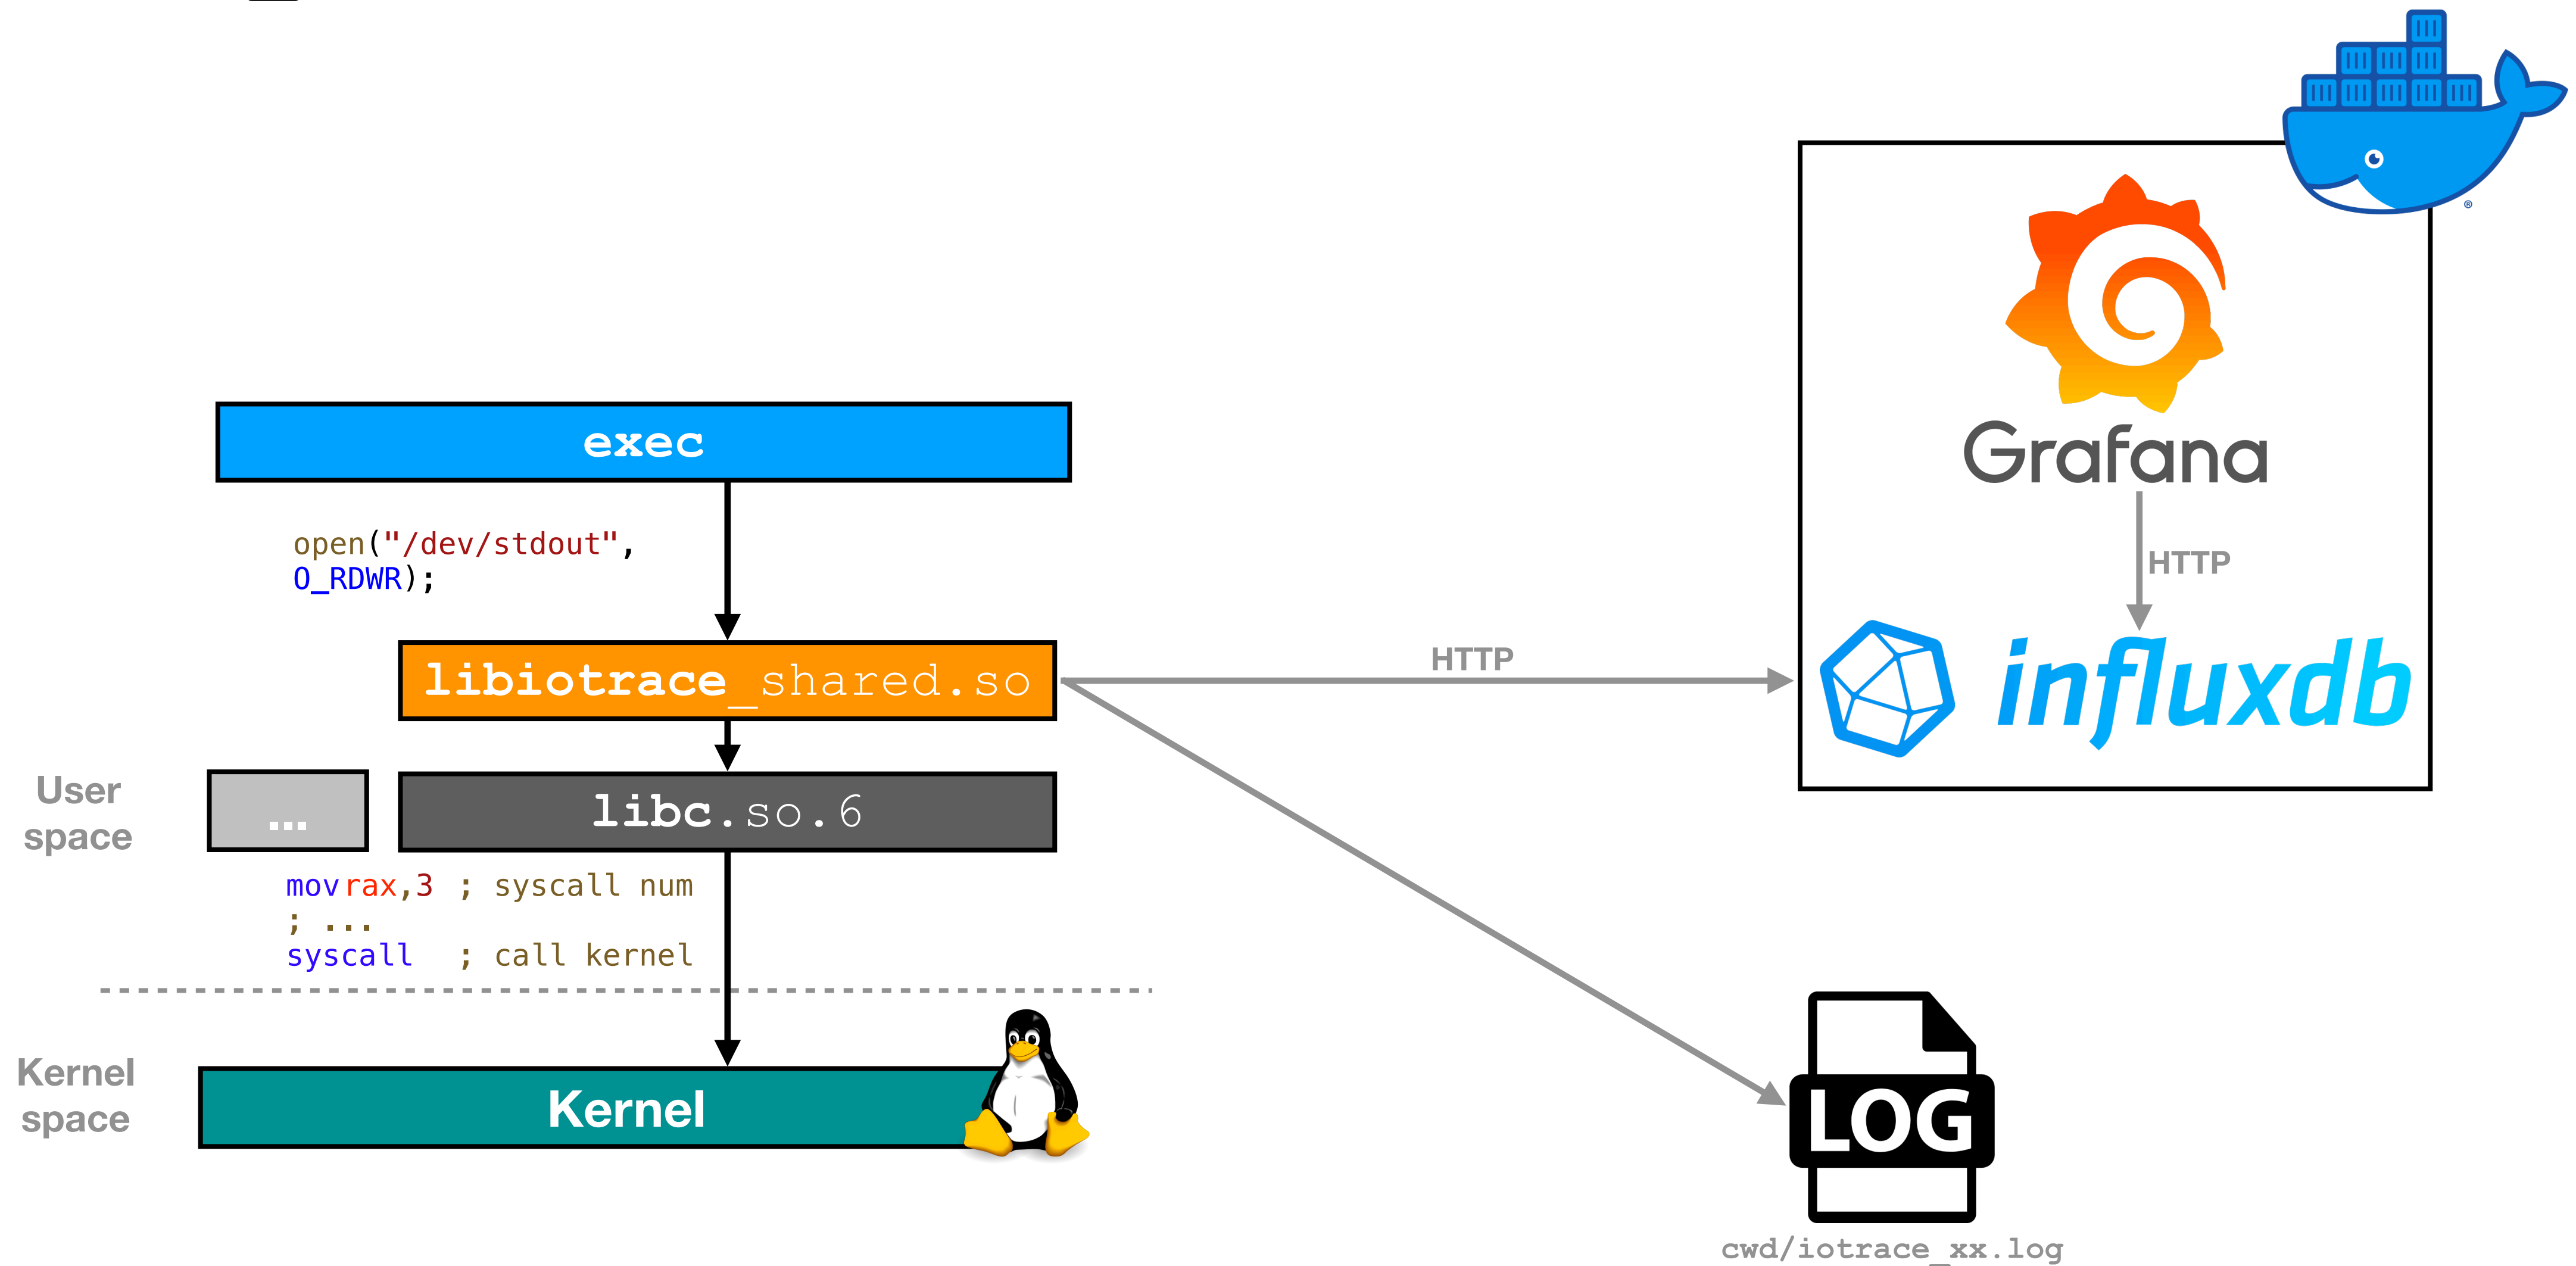
Funktionsweise

→ ~ **LD_PRELOAD**=path/to/libiotrace_shared.so ./path/to/exec



Funktionsweise

→ ~ **LD_PRELOAD**=path/to/libiotrace_shared.so ./path/to/exec



Filename Resolution

Motivation

Motivation

- Problem für Tracing:

Motivation

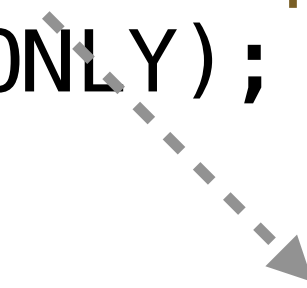
- Problem für Tracing:
- Initiales öffnen: **Filename**

```
int fd = open("path/to/file",  
0_RDONLY);
```


Motivation

- Problem für Tracing:
- Initiales öffnen: **Filename**
- Folgende Operationen: **Handles**

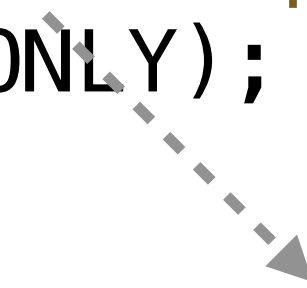
```
int fd = open("path/to/file",  
0_RDONLY);  
  
(void)read(fd, buf, sizeof buf  
-1);
```



Motivation

- Problem für Tracing:
- Initiales öffnen: **Filename**
- Folgende Operationen: **Handles**
- Auswertung - Situation:

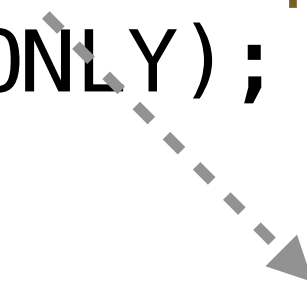
```
int fd = open("path/to/file",  
0_RDONLY);  
  
(void)read(fd, buf, sizeof buf  
-1);
```



Motivation

- Problem für Tracing:
- Initiales öffnen: **Filename**
- Folgende Operationen: **Handles**
- Auswertung - Situation:
- Post mortem: Java-Programm

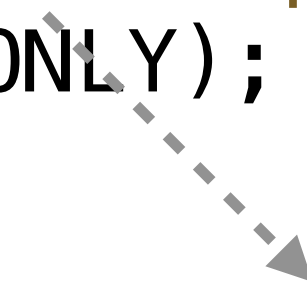
```
int fd = open("path/to/file",  
0_RDONLY);  
  
(void)read(fd, buf, sizeof buf  
-1);
```



Motivation

- Problem für Tracing:
- Initiales öffnen: **Filename**
- Folgende Operationen: **Handles**
- Auswertung - Situation:
- Post mortem: Java-Programm

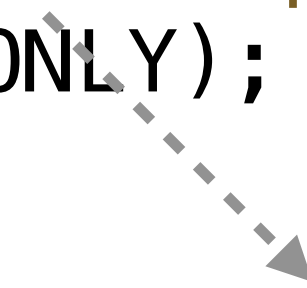
```
int fd = open("path/to/file",  
O_RDONLY);  
  
(void)read(fd, buf, sizeof buf  
-1);
```



Motivation

- Problem für Tracing:
- Initiales öffnen: **Filename**
- Folgende Operationen: **Handles**
- Auswertung - Situation:
- Post mortem: Java-Programm
- Live Tracing

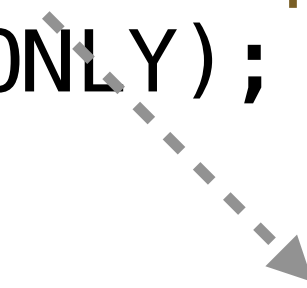
```
int fd = open("path/to/file",  
O_RDONLY);  
  
(void)read(fd, buf, sizeof buf  
-1);
```



Motivation

- Problem für Tracing:
- Initiales öffnen: **Filename**
- Folgende Operationen: **Handles**
- Auswertung - Situation:
- Post mortem: Java-Programm
- Live Tracing

```
int fd = open("path/to/file",  
O_RDONLY);  
  
(void)read(fd, buf, sizeof buf  
-1);
```



Funktionsweise

Handle	Filename
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

Handle	Filename
...	...

"Filename Map"

Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

Handle	Filename
3	"/etc/fstab"
...	...

"Filename Map"

Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

```
{... ,"function_name":"mmap","file_type":  
{"descriptor":3},"function_data":  
{"address":"0xffff9b6f3000","length":100,"protection_flags":  
["read"],... }
```

Handle	Filename
3	"/etc/fstab"
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

```
{... ,"function_name":"mmap","file_type":  
{"descriptor":3},"function_data":  
{"address":"0xffff9b6f3000","length":100,"protection_flags":  
["read"],... }, "traced_filename":"/etc/fstab" }
```

Handle	Filename
3	"/etc/fstab"
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

```
{... ,"function_name":"mmap","file_type":  
{"descriptor":3},"function_data":  
{"address":"0xffff9b6f3000","length":100,"protection_flags":  
["read"],... }, "traced_filename":"/etc/fstab" }
```

Handle	Filename
3	"/etc/fstab"
"0xffff9b6f3000"	"/etc/fstab"
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":
{"descriptor":3},"function_data":{"mode":"read_only","creation":
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":
{"device_id":64768,"inode_nr":799128}}}}

{... ,"function_name":"mmap","file_type":
{"descriptor":3},"function_data":
{"address":"0xffff9b6f3000","length":100,"protection_flags":
["read"],... }, "traced_filename":"/etc/fstab" }

{... ,"function_name":"munmap","wrapper":{"... ,"file_type":
{"address":"0xffff9b6f3000"}
}
```

Handle	Filename
3	"/etc/fstab"
"0xffff9b6f3000"	"/etc/fstab"
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

```
{... ,"function_name":"mmap","file_type":  
{"descriptor":3},"function_data":  
{"address":"0xffff9b6f3000","length":100,"protection_flags":  
["read"],... }, "traced_filename":"/etc/fstab" }
```

```
{... ,"function_name":"munmap","wrapper":{"... ,"file_type":  
{"address":"0xffff9b6f3000"},  
"traced_filename":"/etc/fstab" }
```

Handle	Filename
3	"/etc/fstab"
"0xffff9b6f3000"	"/etc/fstab"
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

```
{... ,"function_name":"mmap","file_type":  
{"descriptor":3},"function_data":  
{"address":"0xffff9b6f3000","length":100,"protection_flags":  
["read"],... }, "traced_filename":"/etc/fstab" }
```

```
{... ,"function_name":"munmap","wrapper":{"... ,"file_type":  
{"address":"0xffff9b6f3000"},  
"traced_filename":"/etc/fstab" }
```

Handle	Filename
3	"/etc/fstab"
"0xffff9b6f3000"	"/etc/fstab"
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}  
  
{... ,"function_name":"mmap","file_type":  
{"descriptor":3},"function_data":  
{"address":"0xffff9b6f3000","length":100,"protection_flags":  
["read"],... }, "traced_filename":"/etc/fstab" }  
  
{...,"function_name":"munmap","wrapper":{"...,"file_type":  
{"address":"0xffff9b6f3000"},  
"traced_filename":"/etc/fstab"}  
  
{...,"function_name":"close","return_state":"ok","file_type":  
{"descriptor":3}}
```

Handle	Filename
3	/etc/fstab
0xffff9b6f3000	/etc/fstab
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

```
{... ,"function_name":"mmap","file_type":  
{"descriptor":3},"function_data":  
{"address":"0xffff9b6f3000","length":100,"protection_flags":  
["read"],... }, "traced_filename":"/etc/fstab" }
```

```
{... ,"function_name":"munmap","wrapper":{"... ,"file_type":  
{"address":"0xffff9b6f3000"},  
"traced_filename":"/etc/fstab"}
```

```
{... ,"function_name":"close","return_state":"ok","file_type":  
{"descriptor":3}, "traced_filename":"/etc/fstab"}
```

Handle	Filename
3	/etc/fstab
0xffff9b6f3000	/etc/fstab
...	...

"Filename Map"
Handle -> Filename

Funktionsweise

```
{... "function_name":"open","file_type":  
{"descriptor":3},"function_data":{"mode":"read_only","creation":  
[],"status":[],"file_mode":[],"file_name":"/etc/fstab","id":  
{"device_id":64768,"inode_nr":799128}}}
```

```
{... ,"function_name":"mmap","file_type":  
{"descriptor":3},"function_data":  
{"address":"0xffff9b6f3000","length":100,"protection_flags":  
["read"],... }, "traced_filename":"/etc/fstab" }
```

```
{... ,"function_name":"munmap","wrapper":{"... ,"file_type":  
{"address":"0xffff9b6f3000"},  
"traced_filename":"/etc/fstab" }
```

```
{... ,"function_name":"close","return_state":"ok","file_type":  
{"descriptor":3}, "traced_filename":"/etc/fstab" }
```

Handle	Filename
3	/etc/fstab
0xffff9b6f3000	/etc/fstab
...	...

"Filename Map"
Handle -> Filename

Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL  
function_name: "open";  
function_data: {...}  
...
```

Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL
```

```
function_name: "open";
```

```
function_data: {...}
```

```
...
```

```
fnres_trace_fc
```

```
tevent(&b);
```



`fctevent.c`

Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL  
function_name: "open";  
function_data: {...}  
...
```

```
fnres_trace_fc  
tevent(&b);
```



`fctevent.c`

```
1. switch-case(hashcd_fct_name)
```

Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL  
function_name: "open";  
function_data: {...}  
...
```

```
fnres_trace_fc  
tevent(&b);
```



`fctevent.c`

1. `switch-case` (`hashed_fct_name`)
2. Map Key ableiten von `basic` struct

Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL  
function_name: "open";  
function_data: {...}  
...
```

```
fnres_trace_fc  
tevent(&b);
```



`fctevent.c`

1. `switch-case` (`hashed_fct_name`)
2. Map Key ableiten von `basic` struct

```
typedef struct {  
    union id handle;  
    enum id_type handle_type;  
} fnmap_key;
```

Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL  
function_name: "open";  
function_data: {...}  
...
```

```
fnres_trace_fctevent(&b);
```



fctevent.c

1. **switch-case** (hashed_fct_name)
2. Map Key ableiten von `basic` struct
3. Key hashen + Pointer zu filename C String speichern

```
typedef struct {  
    union id handle;  
    enum id_type handle_type;  
} fnmap_key;
```


Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL  
function_name: "open";  
function_data: {...}  
...
```

```
fnres_trace_fctevent(&b);
```



`fctevent.c`

1. **switch-case** (`hashed_fct_name`)
2. Map Key ableiten von `basic` struct
3. Key hashen + Pointer zu filename C String speichern
4. `traced_filename` setzen

```
typedef struct {  
    union id handle;  
    enum id_type handle_type;  
} fnmap_key;
```

Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL  
function_name: "open";  
function_data: {...}  
...
```

```
fnres_trace_fctevent(&b);
```



fctevent.c

1. **switch-case** (hashed_fct_name)
2. Map Key ableiten von `basic` struct
3. Key hashen + Pointer zu filename C String speichern
4. `traced_filename` setzen

AFTER:

```
struct basic b;
```

```
traced_filename: 0x0300..  
function_name: "open";  
function_data: {...}  
...
```

```
typedef struct {  
    union id handle;  
    enum id_type handle_type;  
} fnmap_key;
```

Integration: Modul `fnres`

BEFORE:

```
struct basic b;
```

```
traced_filename: NULL  
function_name: "open";  
function_data: {...}  
...
```

```
fnres_trace_fctevent(&b);
```



fctevent.c

1. **switch-case** (hashed_fct_name)
2. Map Key ableiten von `basic` struct
3. Key hashen + Pointer zu filename C String speichern
4. `traced_filename` setzen

AFTER:

```
struct basic b;
```

```
traced_filename: 0x0300..  
function_name: "open";  
function_data: {...}  
...
```

```
typedef struct {  
    union id handle;  
    enum id_type handle_type;  
} fnmap_key;
```

- `atomic_hash`: Lock-free, C
- Map: Hashed `fnmap_key -> char* filename`

Motivation - Syscall Tracing

Limitierungen `fnres`

Limitierungen `fnres`

```
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "mmap",  
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "close",
```

Limitierungen `fnres`

```
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "mmap",  
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "close",
```

Limitierungen fnres

```
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "mmap",  
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "close",
```

- Fortified functions `open -> __open_2`

Limitierungen fnres

```
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "mmap",  
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "close",
```

- Fortified functions `open -> __open_2`
- Kein libc call ?!

Limitierungen fnres

```
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "mmap",  
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "close",
```

- Fortified functions `open -> __open_2`
- Kein libc call ?!

```
[pid 35846] openat(AT_FDCWD, "/dev/shm/UC8GnG", O_RDWR|O_CREAT|O_EXCL, 0644) = 17  
/lib/x86_64-linux-gnu/libpthread-2.31.so(__open64+0xd4) [0x14ad4]  
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x2b5) [0x12df5]  
...  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]  
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]  
...  
[pid 35846] mmap(NULL, 32, PROT_READ|PROT_WRITE, MAP_SHARED, 17, 0) = 0x7f57ac8ee000  
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]  
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]  
...  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]  
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```

Limitierungen fnres

```
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "mmap",  
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "close",
```

- Fortified functions `open -> __open_2`
- Kein libc call ?!

```
[pid 35846] openat(AT_FDCWD, "/dev/shm/UC8GnG", O_RDWR|O_CREAT|O_EXCL, 0644) = 17  
/lib/x86_64-linux-gnu/libpthread-2.31.so(__open64+0xd4) [0x14ad4] <- Kein "libc call" = Kein libiotrace wrapper  
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x2b5) [0x12df5]  
...  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]  
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]  
...  
[pid 35846] mmap(NULL, 32, PROT_READ|PROT_WRITE, MAP_SHARED, 17, 0) = 0x7f57ac8ee000  
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]  
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]  
...  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]  
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```


Limitierungen fnres

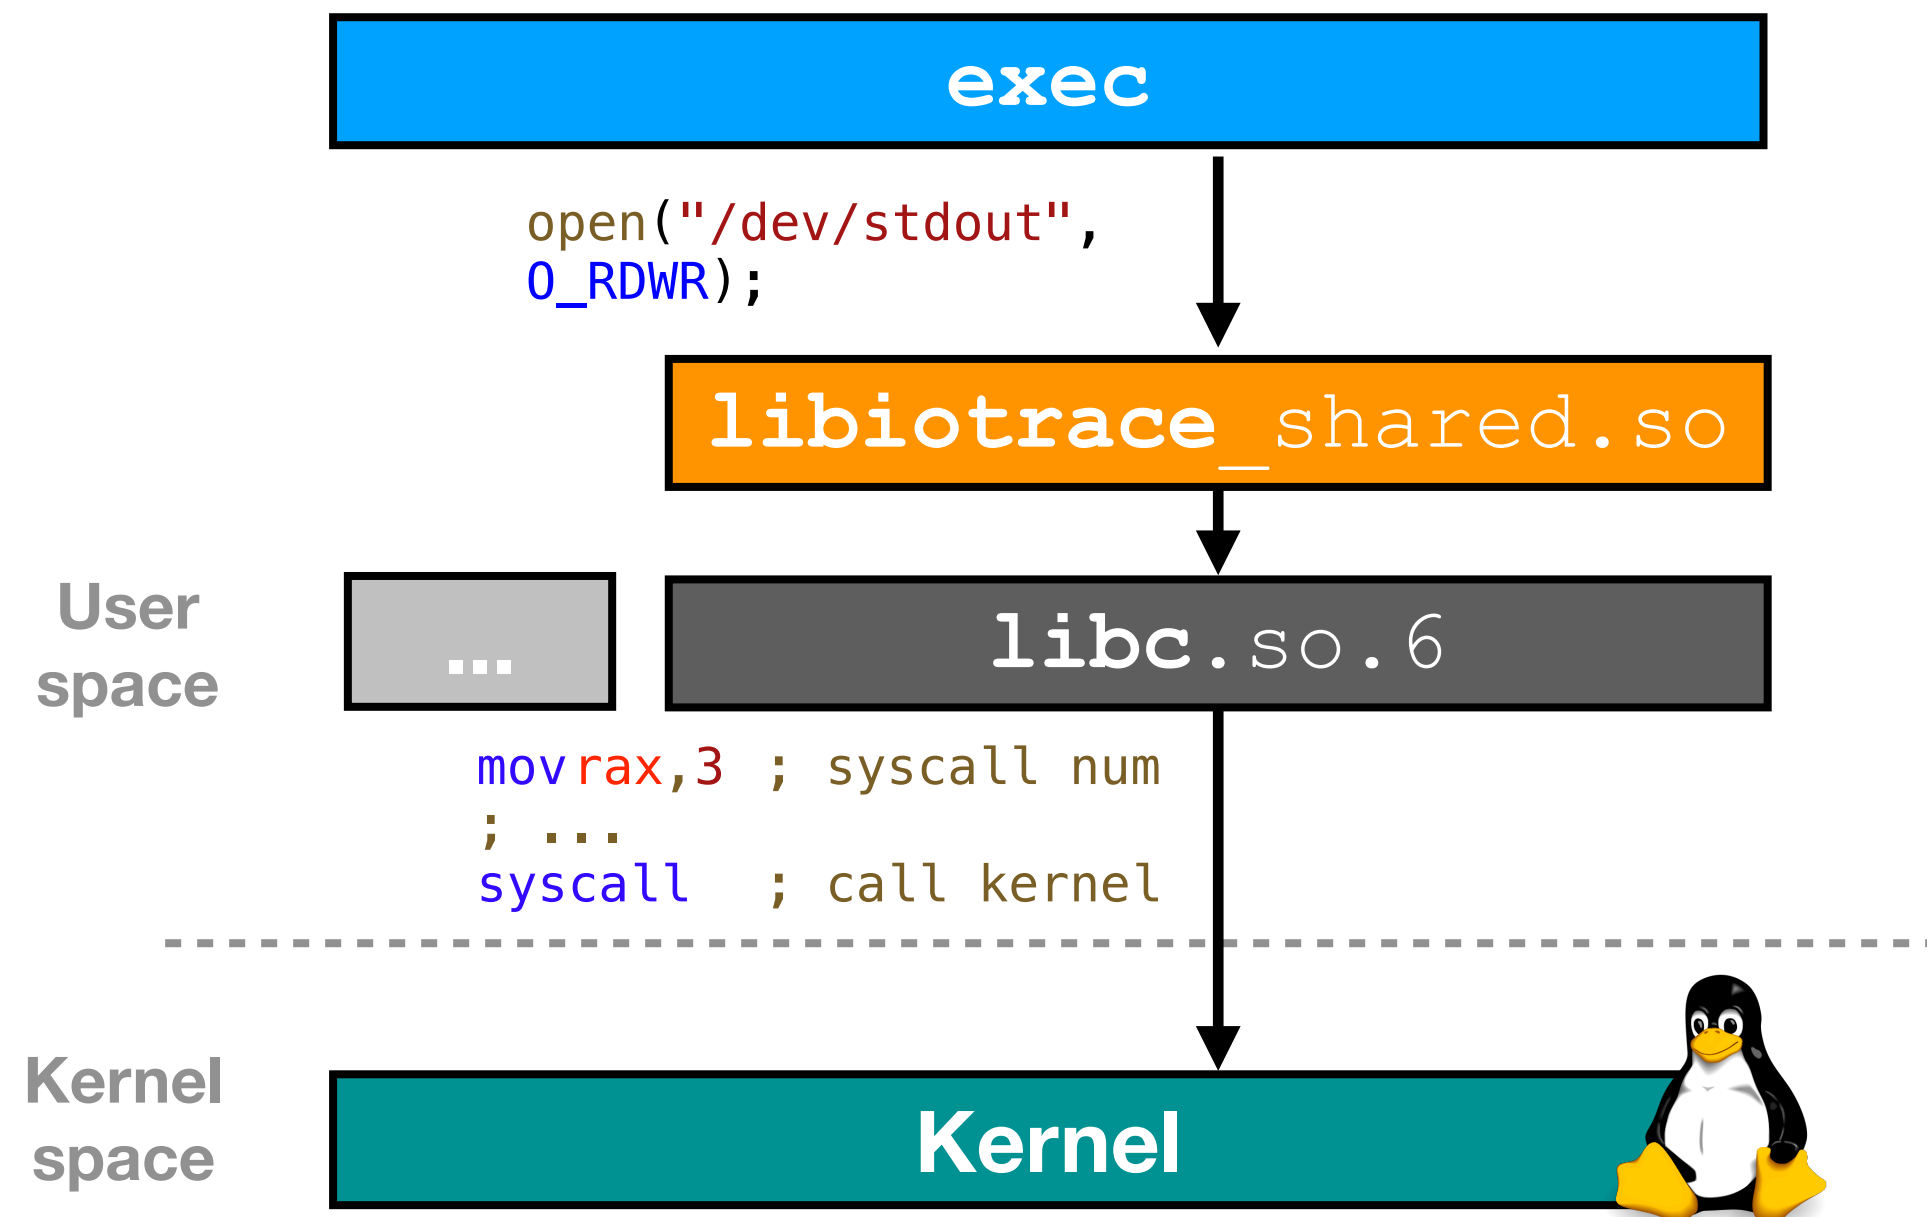
```
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "mmap",  
filename": "_ NOT FOUND _", "hostname": "fpj-vm", "process_id": 16215, "thread_id": 16215, "function_name": "close",
```

- Fortified functions `open -> __open_2`
- Kein libc call ?!

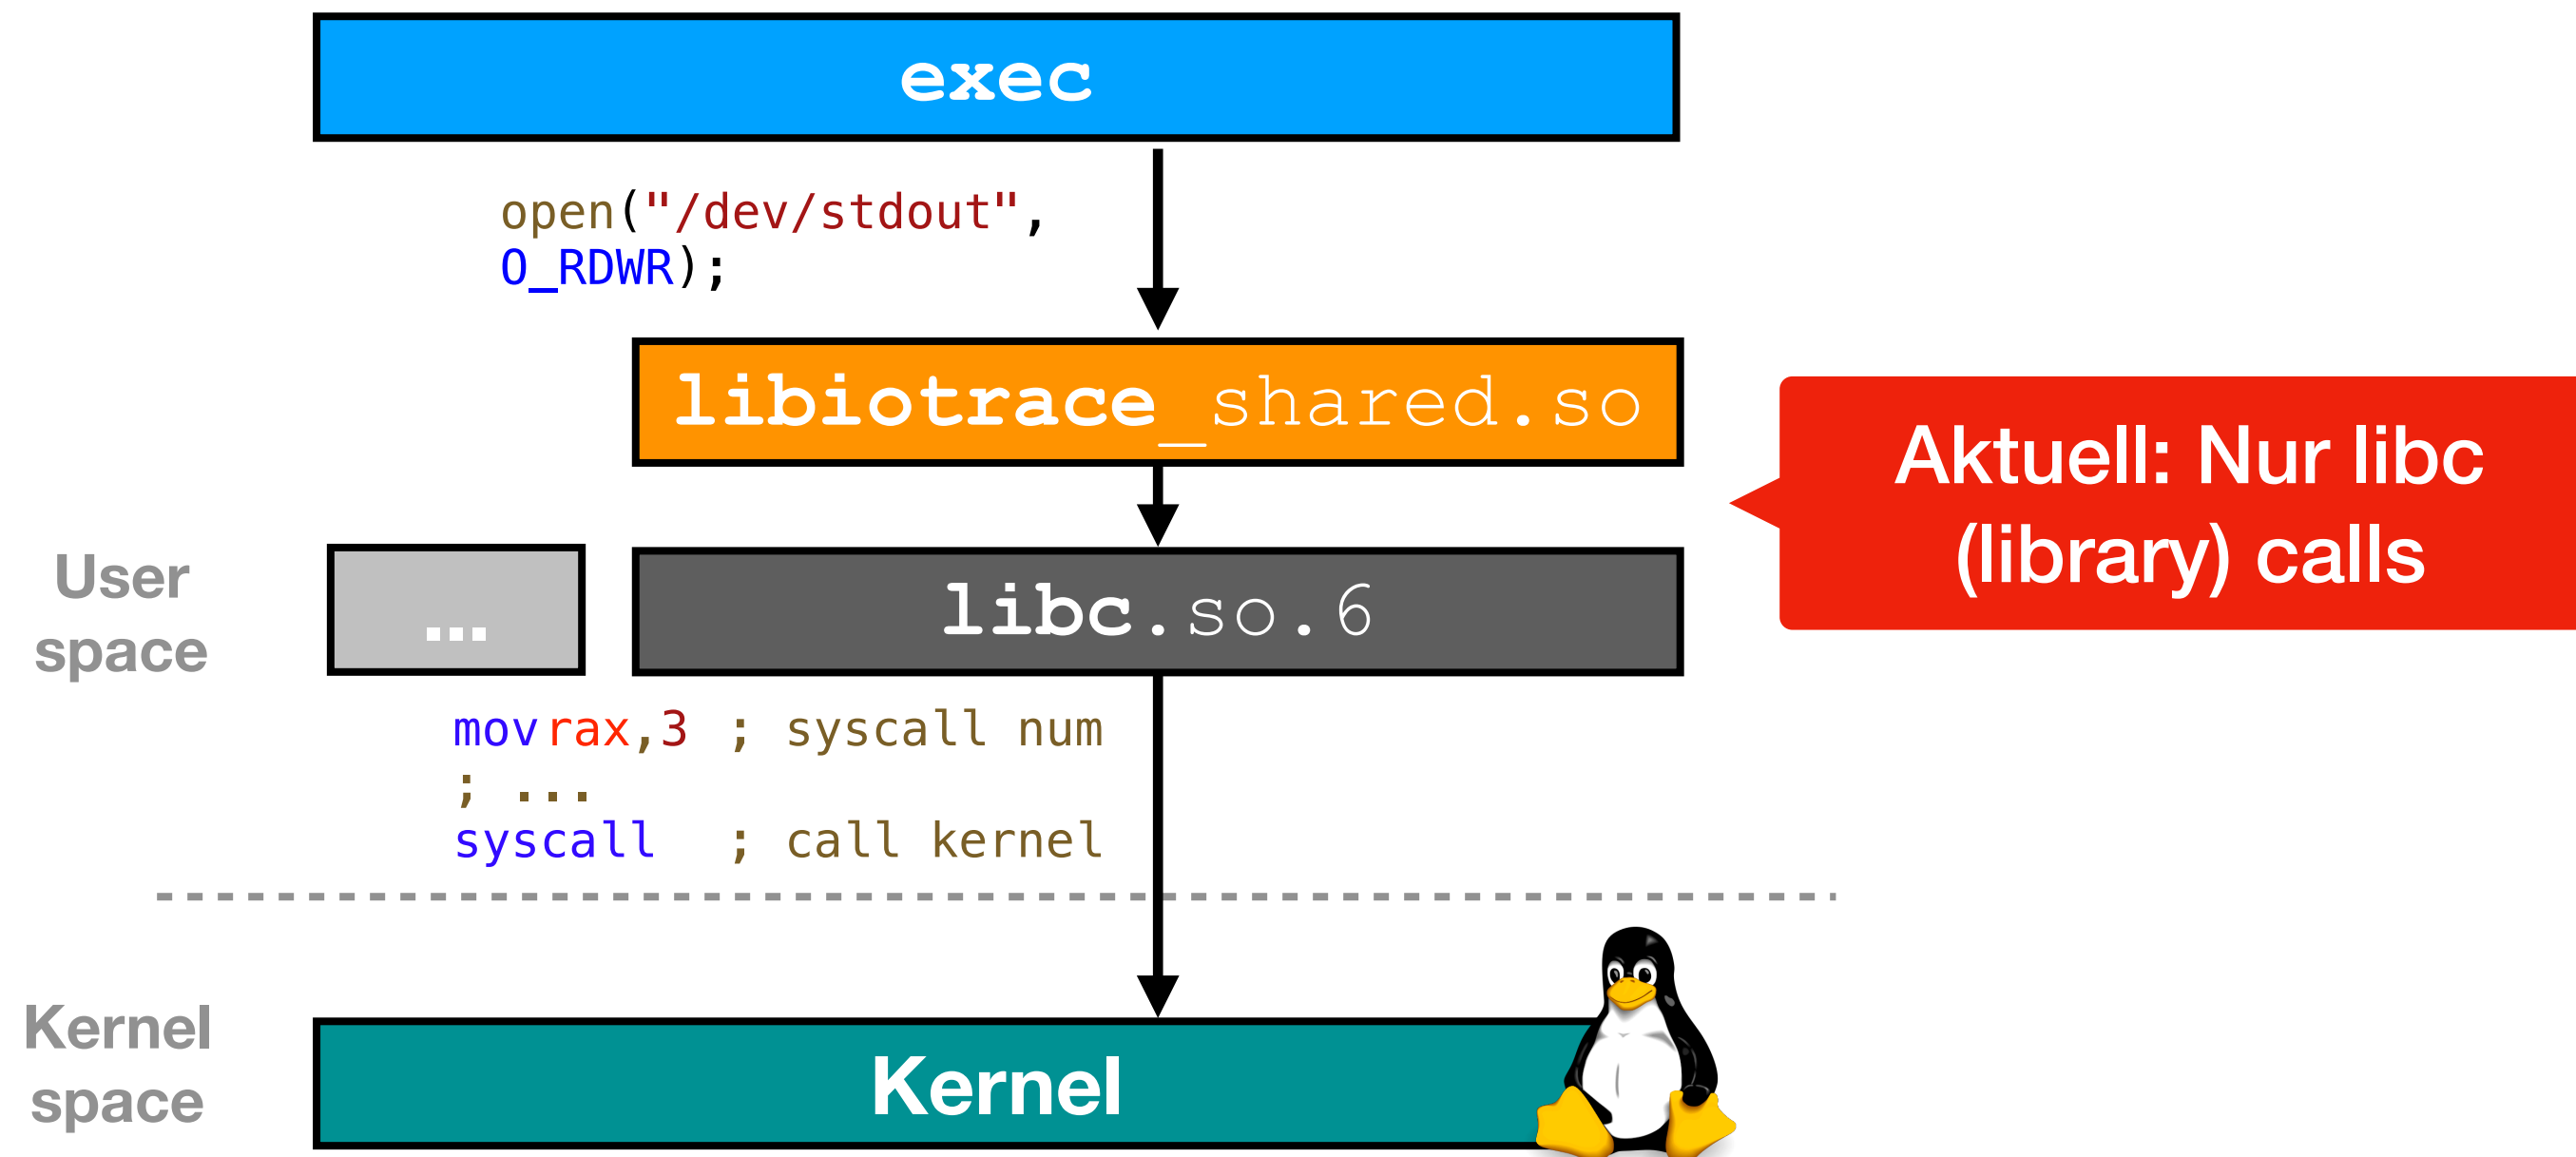
```
[pid 35846] openat(AT_FDCWD, "/dev/shm/UC8GnG", O_RDWR|O_CREAT|O_EXCL, 0644) = 17  
/lib/x86_64-linux-gnu/libpthread-2.31.so(__open64+0xd4) [0x14ad4]  
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x2b5) [0x12df5]  
...  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]  
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]  
...  
[pid 35846] mmap(NULL, 32, PROT_READ|PROT_WRITE, MAP_SHARED, 17, 0) = 0x7f57ac8ee000  
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]  
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]  
...  
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]  
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]  
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```

<- Kein "libc call" = Kein libiotrace wrapper

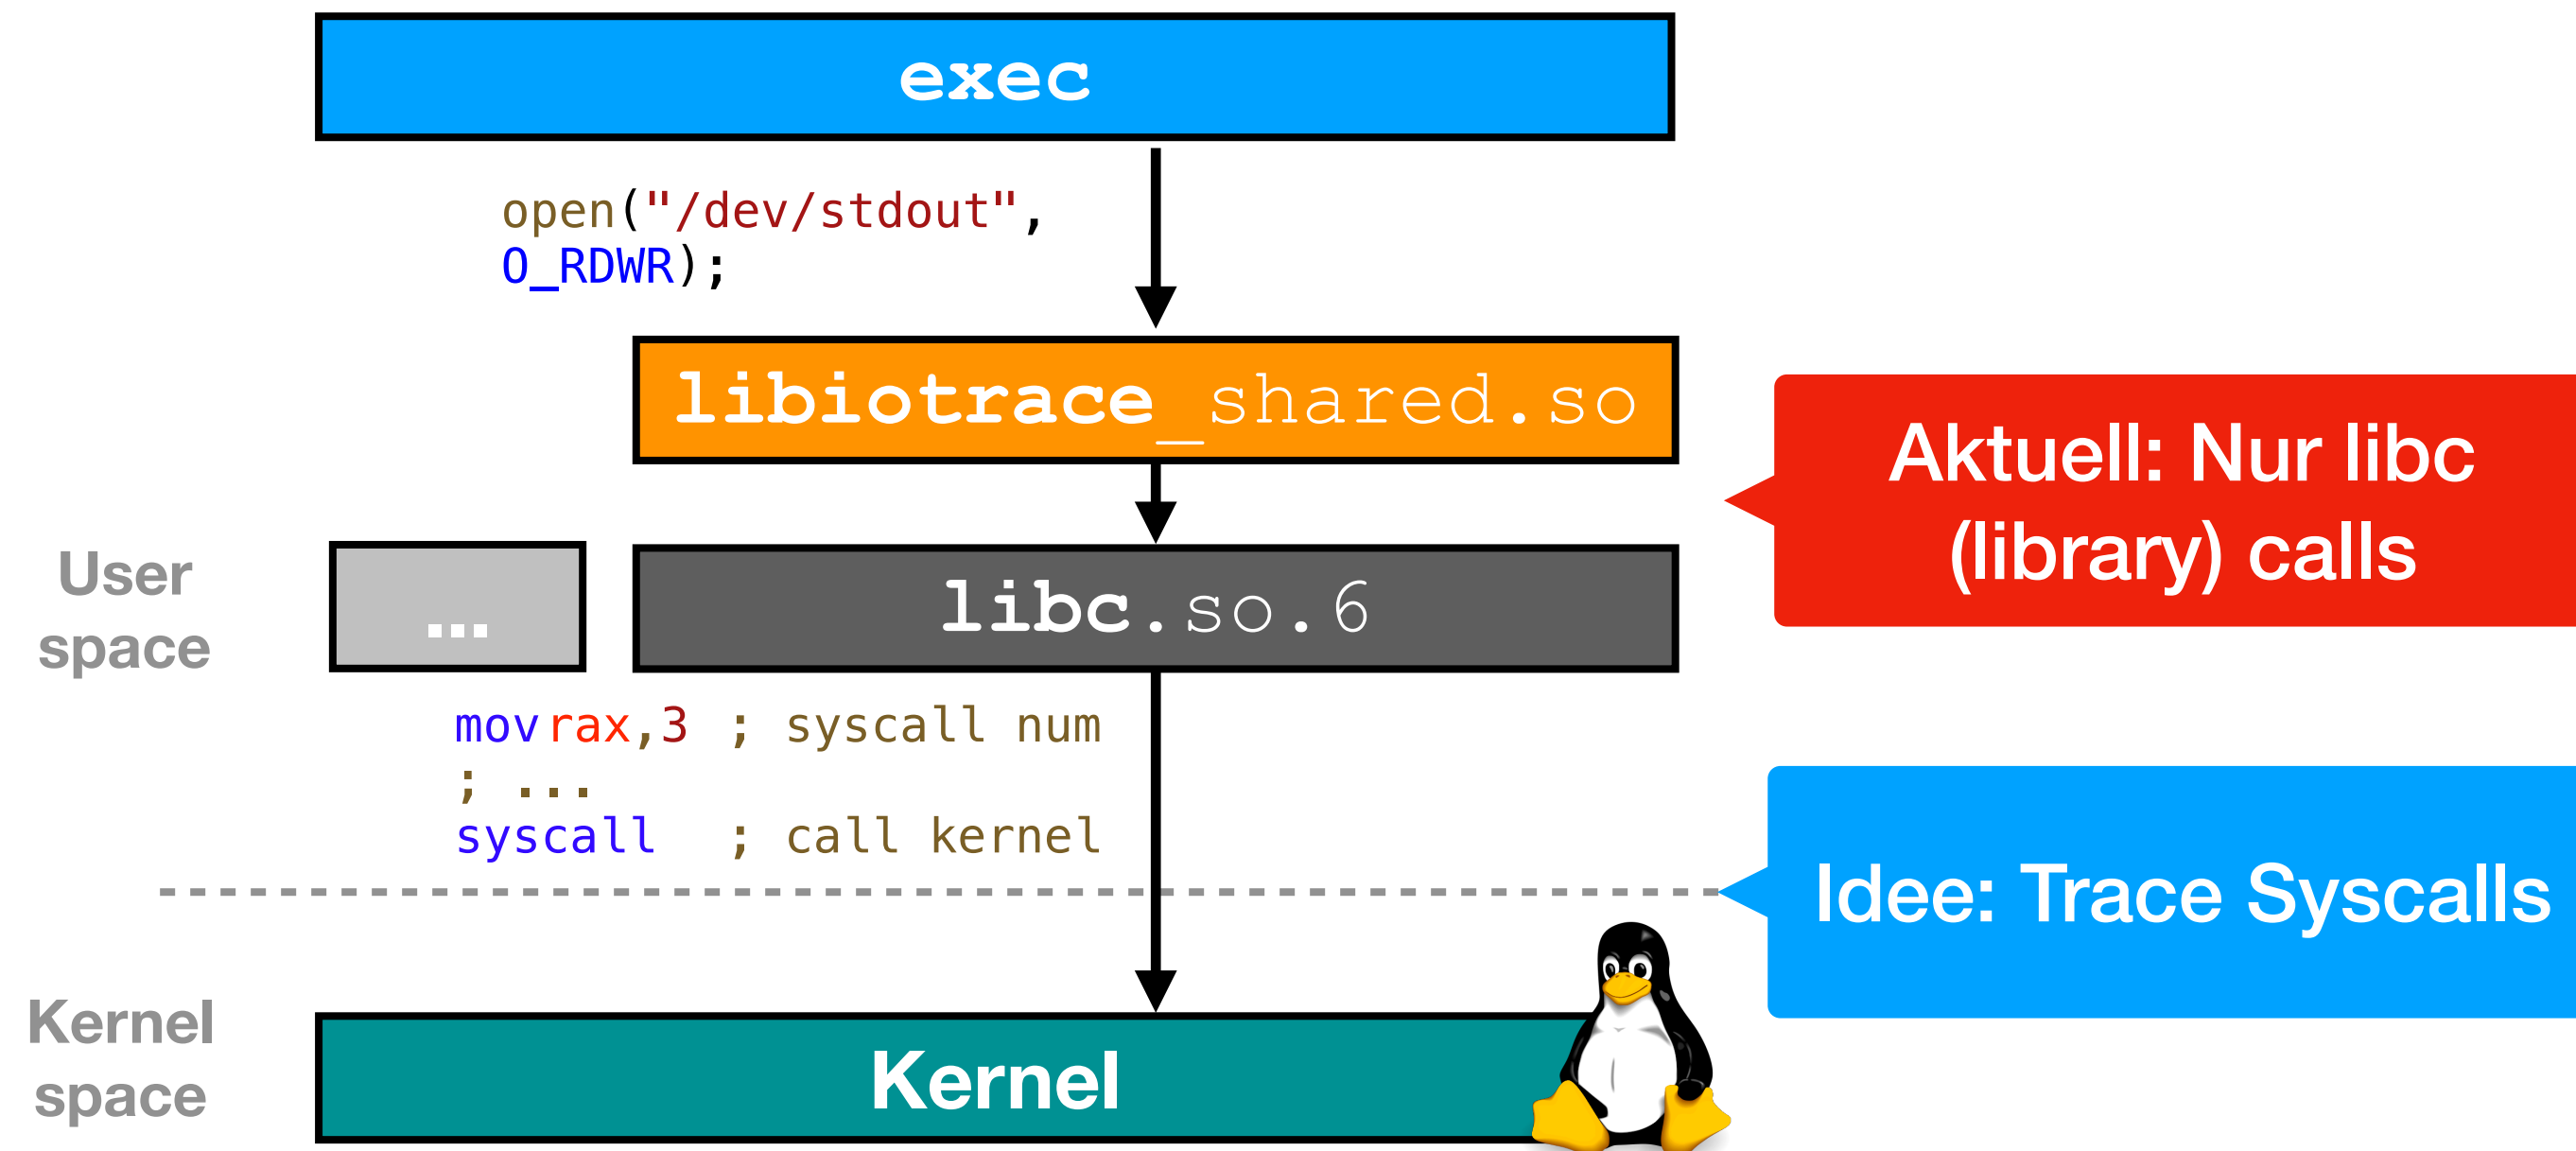
"Remedy"



"Remedy"




"Remedy"



"Remedy"

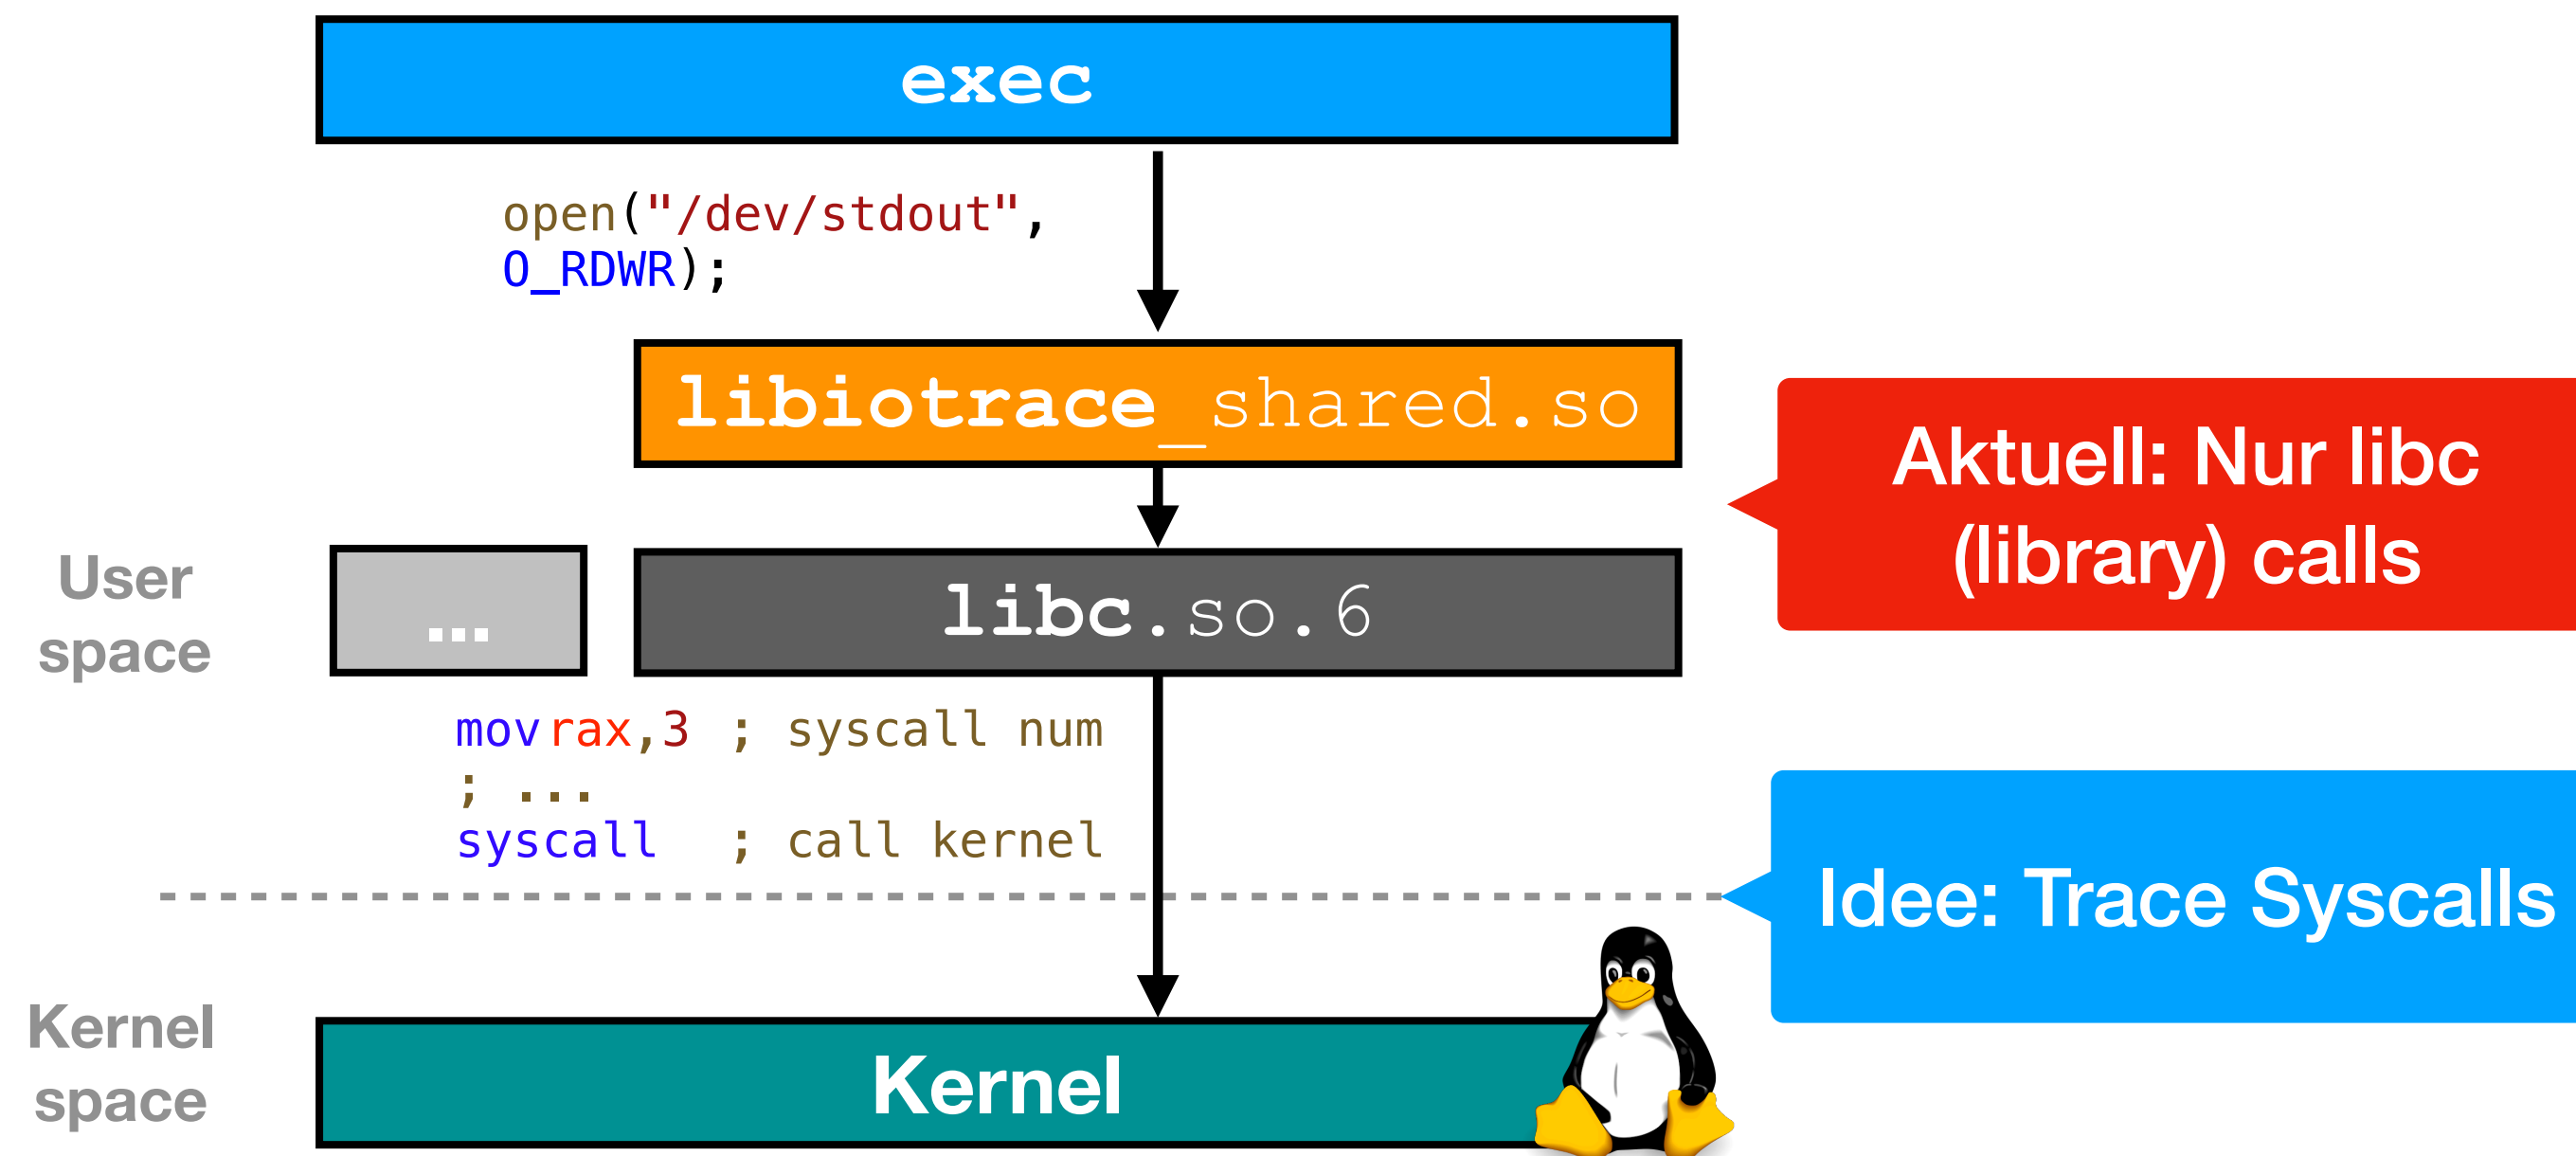
- Warum?



"Remedy"


- **Warum?**

- Jeder `open` -> `Syscall`

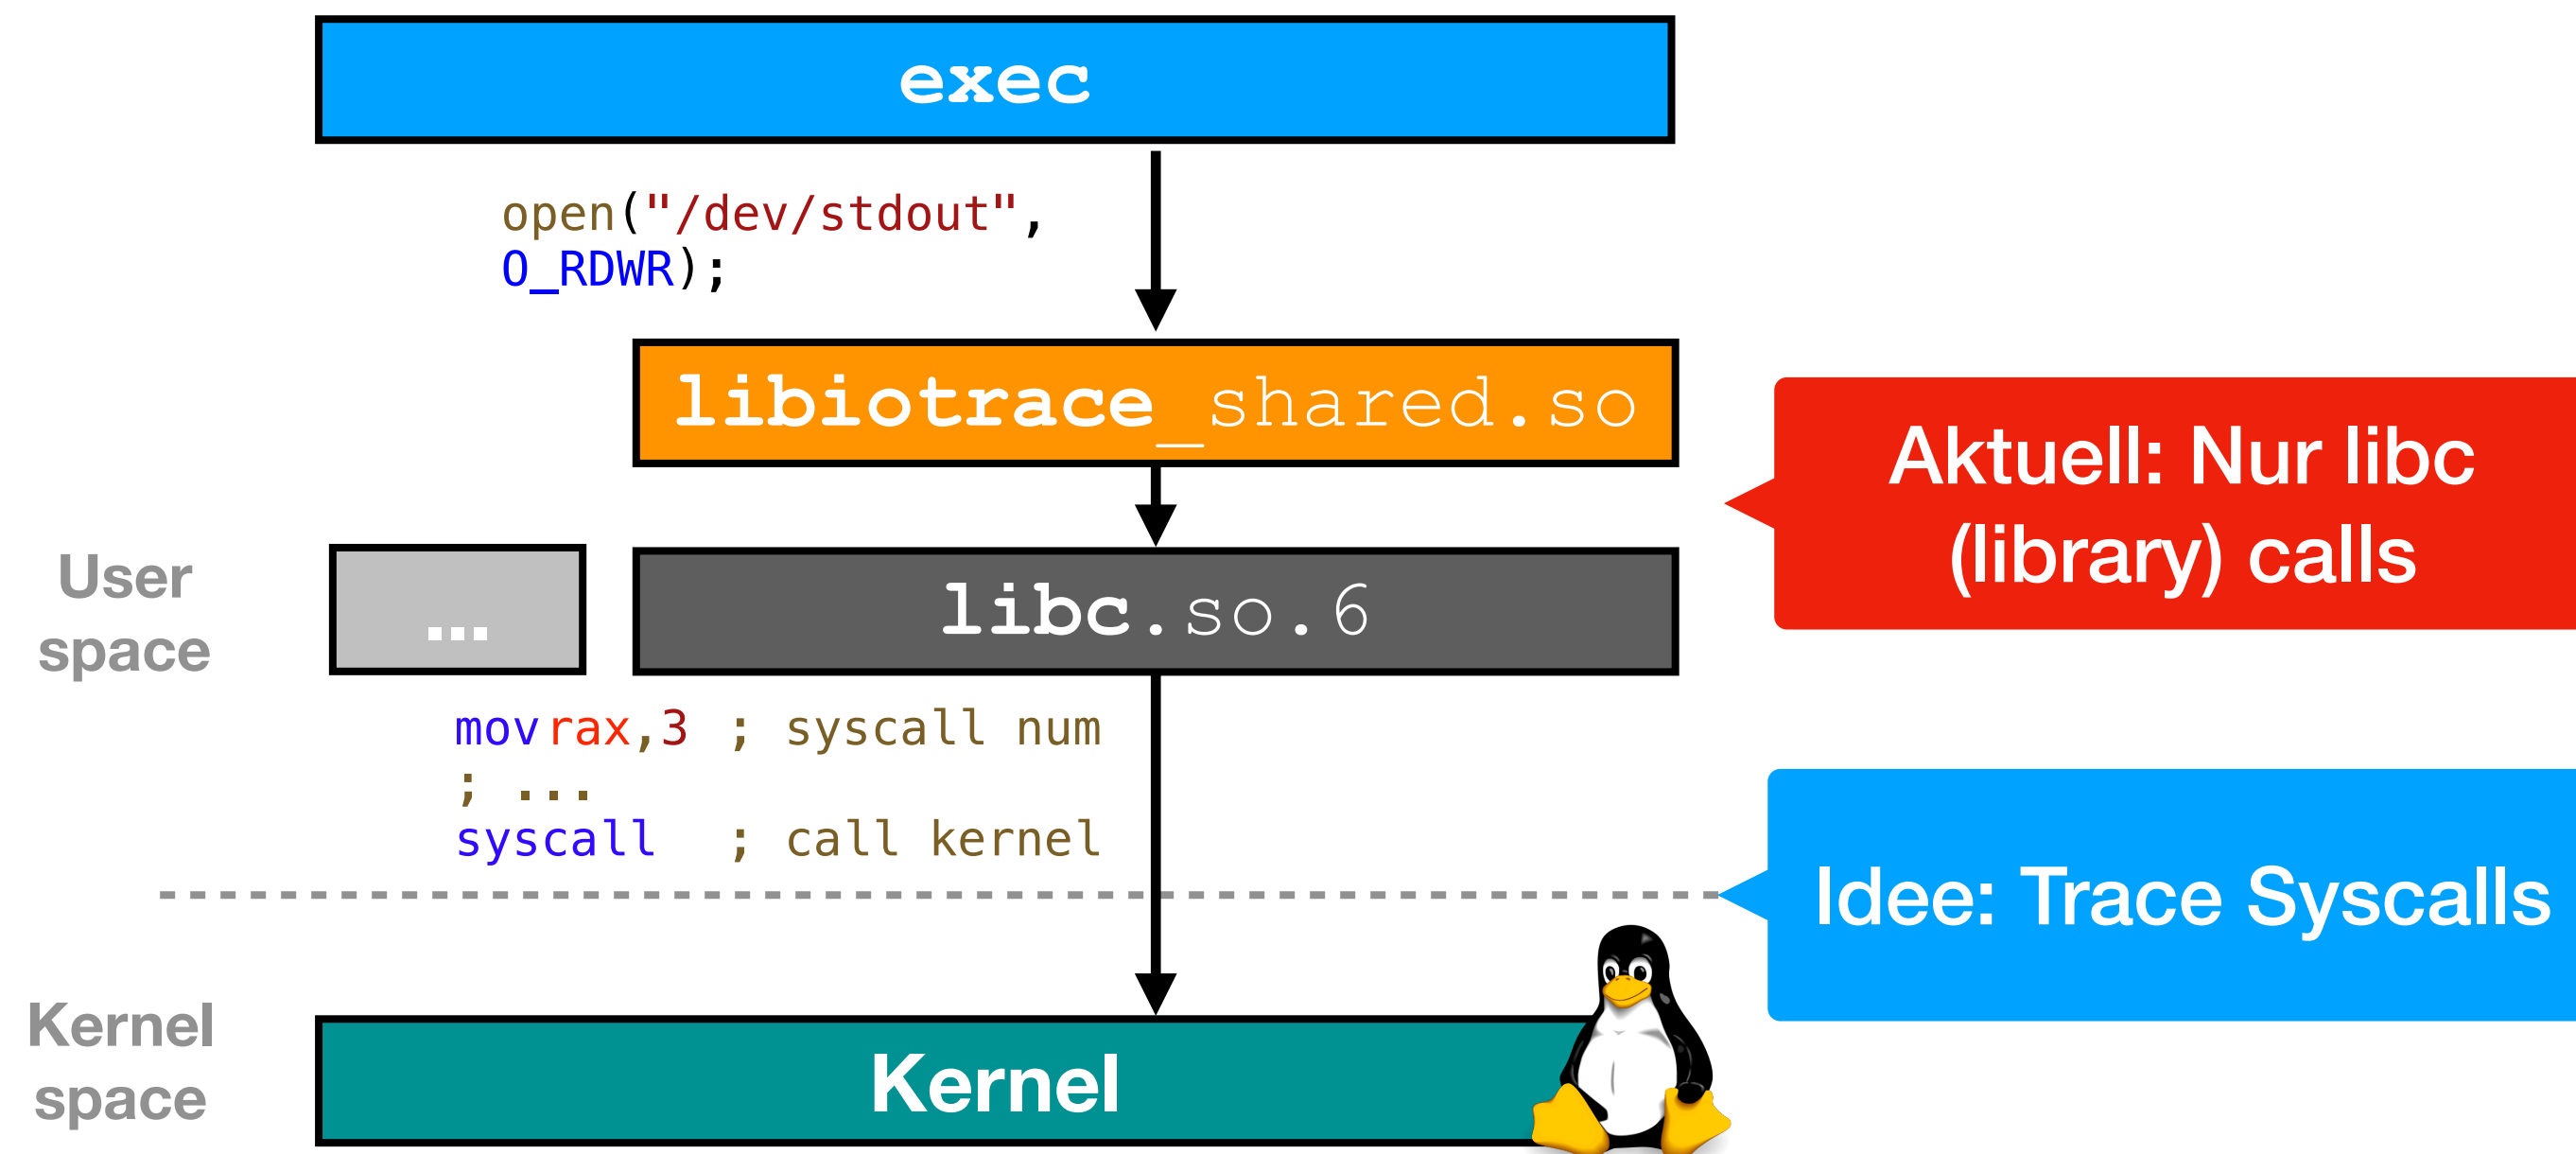


"Remedy"

- Warum?
- Jeder open -> Syscall
- Wie?



"Remedy"



- **Warum?**

- Jeder `open` -> `Syscall`

- **Wie?**

- `ptrace(2)`

Ansatz für Umsetzung

Ansatz für Umsetzung

- **Execution Stack unwinding**

Ansatz für Umsetzung

- **Execution Stack unwinding**
- Von Remote Target !!

Ansatz für Umsetzung

- **Execution Stack unwinding**
- **Von Remote Target !!**

```
[pid 35846] mmap(NULL, 32, PROT_READ|PROT_WRITE, MAP_SHARED, 17, 0) = 0x7f57ac8ee000
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```


Ansatz für Umsetzung

- **Execution Stack unwinding**
- **Von Remote Target !!**

```
[pid 35846] mmap(NULL, 32, PROT READ|PROT WRITE, MAP SHARED, 17, 0) = 0x7f57ac8ee000
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```


Ansatz für Umsetzung


- **Execution Stack unwinding**
- **Von Remote Target !!**



```
[pid 35846] mmap(NULL, 32, PROT_READ|PROT_WRITE, MAP_SHARED, 17, 0) = 0x7f57ac8ee000
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```

Ansatz für Umsetzung

- **Execution Stack unwinding**
- **Von Remote Target !!**




```
[pid 35846] mmap(NULL, 32, PROT READ|PROT WRITE, MAP SHARED, 17, 0) = 0x7f57ac8ee000
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```

```
[pid 35846] openat(AT_FDCWD, "/dev/shm/UC8GnG", O_RDWR|O_CREAT|O_EXCL, 0644) = 17
/lib/x86_64-linux-gnu/libpthread-2.31.so(__open64+0xd4) [0x14ad4]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x2b5) [0x12df5]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```

Ansatz für Umsetzung

- **Execution Stack unwinding**
- **Von Remote Target !!**




```
[pid 35846] mmap(NULL, 32, PROT READ|PROT WRITE, MAP SHARED, 17, 0) = 0x7f57ac8ee000
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```


```
[pid 35846] openat(AT_FDCWD, "/dev/shm/UC8GnG", O_RDWR|O_CREAT|O_EXCL, 0644) = 17
/lib/x86_64-linux-gnu/libpthread-2.31.so(__open64+0xd4) [0x14ad4]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x2b5) [0x12df5]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```


Ansatz für Umsetzung

- **Execution Stack unwinding**
- **Von Remote Target !!**



```
[pid 35846] mmap(NULL, 32, PROT_READ|PROT_WRITE, MAP_SHARED, 17, 0) = 0x7f57ac8ee000
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```



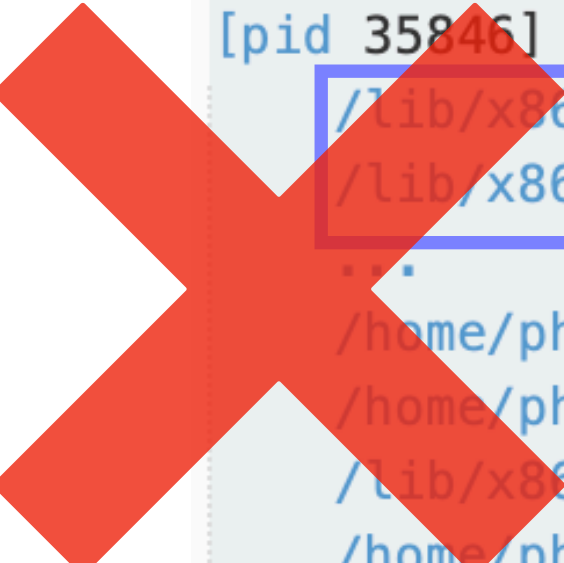
```
[pid 35846] openat(AT_FDCWD, "/dev/shm/UC8GnG", O_RDWR|O_CREAT|O_EXCL, 0644) = 17
/lib/x86_64-linux-gnu/libpthread-2.31.so(__open64+0xd4) [0x14ad4]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x2b5) [0x12df5]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```

Ansatz für Umsetzung

- Execution Stack unwinding
- Von Remote Target !!



```
[pid 35846] mmap(NULL, 32, PROT_READ|PROT_WRITE, MAP_SHARED, 17, 0) = 0x7f57ac8ee000
/lib/x86_64-linux-gnu/libc-2.31.so(mmap64+0x26) [0x11ba46]
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(mmap+0x197) [0x943d7]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x41e) [0x12f5e]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```



```
[pid 35846] openat(AT_FDCWD, "/dev/shm/UC8GnG", O_RDWR|O_CREAT|O_EXCL, 0644) = 17
/lib/x86_64-linux-gnu/libpthread-2.31.so(__open64+0xd4) [0x14ad4]
/lib/x86_64-linux-gnu/libpthread-2.31.so(sem_open+0x2b5) [0x12df5]
...
/home/philkoe/git/fsprj2/libiotrace/build/src/libiotrace_shared.so(MPI_File_open+0x20a) [0xa8899]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(main+0xc8) [0x1391]
/lib/x86_64-linux-gnu/libc-2.31.so(__libc_start_main+0xf3) [0x270b3]
/home/philkoe/git/fsprj2/libiotrace/build/test/MPI_read(_start+0x2e) [0x120e]
```



**Warnung,
Eintrag ergänzen,
....**

ptrace(2)

ptrace(2)

- **Syscall** (`sys_ptrace`)

`ptrace`(2)


- **Syscall** (`sys_ptrace`)
- Unix-like OS (Linux, macOS, ...)

`ptrace`(2)

- `Syscall (sys_ptrace)`
- Unix-like OS (Linux, macOS, ...)
- "**process trace**"

ptrace(2)


- Syscall (`sys_ptrace`)
- Unix-like OS (Linux, macOS, ...)
- "process trace"
- Rollen: **Tracer & Tracee**



Quelle: fadeevab.com

ptrace(2)


- Syscall (`sys_ptrace`)
- Unix-like OS (Linux, macOS, ...)
- "process trace"
- Rollen: **Tracer & Tracee**
- State **inspection** & modification



Quelle: fadeevab.com

ptrace(2)

- Syscall (`sys_ptrace`)
- Unix-like OS (Linux, macOS, ...)
- "process trace"
- Rollen: **Tracer & Tracee**
- State **inspection** & modification
- Use-cases: Breakpoint Debugging -> `gdb(1)` & **Syscall Tracing** -> `strace(1)`



Quelle: fadeevab.com

Syscall Tracing

Syscall Tracing

- Syscalls verwenden CPU Register:

%rax	System call	%rdi	%rsi	%rdx	%r10	%r8	%r9
0	sys_read	unsigned int fd	char *buf	size_t count			
1	sys_write	unsigned int fd	const char *buf	size_t count			
2	sys_open	const char *filename	int flags	int mode			
3	sys_close	unsigned int fd					
4	sys_stat	const char *filename	struct stat *statbuf				
5	sys_fstat	unsigned int fd	struct stat *statbuf				
6	sys_lstat	const char *filename	struct stat *statbuf				

Quelle: blog.rchapman.org

Syscall Tracing

- Syscalls verwenden CPU Register:

%rax	System call	%rdi	%rsi	%rdx	%r10	%r8	%r9
0	sys_read	unsigned int fd	char *buf	size_t count			
1	sys_write	unsigned int fd	const char *buf	size_t count			
2	sys_open	const char *filename	int flags	int mode			
3	sys_close	unsigned int fd					
4	sys_stat	const char *filename	struct stat *statbuf				
5	sys_fstat	unsigned int fd	struct stat *statbuf				
6	sys_lstat	const char *filename	struct stat *statbuf				

Quelle: blog.rchapman.org

- Für tracee auslesen: `PTRACE_GETREGSET`

ministrace

strace**(1)**

strace(1)

- man: "trace system calls and signals"

strace(1)

- man: "trace system calls and signals"
- Verwendet ptrace(2)

strace(1)


- man: "trace system calls and signals"
- Verwendet ptrace(2)

```
→ ~ strace sleep 5
execve("/usr/bin/sleep", ["sleep", "5"], 0xffffffff1e5f048 /* 25 vars */) = 0
brk(NULL)                                = 0xaaaaad0698000
faccessat(AT_FDCWD, "/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=29405, ...}) = 0
mmap(NULL, 29405, PROT_READ, MAP_PRIVATE, 3, 0) = 0xfffff88deb000
close(3)                                = 0
...
clock_nanosleep(CLOCK_REALTIME, 0, {tv_sec=5, tv_nsec=0},
```

strace(1)

- man: "trace system calls and signals"
- Verwendet ptrace(2)

Syscall-Name (NR)

→  `strace sleep 5`
`execve("/usr/bin/sleep", ["sleep", "5"], 0xffffffff1e5f048 /* 25 vars */) = 0`
`brk(NULL) = 0xaaaad0698000`
`faccessat(AT_FDCWD, "/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)`
`openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3`
`fstat(3, {st_mode=S_IFREG|0644, st_size=29405, ...}) = 0`
`mmap(NULL, 29405, PROT_READ, MAP_PRIVATE, 3, 0) = 0xfffff88deb000`
`close(3) = 0`
`...`
`clock_nanosleep(CLOCK_REALTIME, 0, {tv_sec=5, tv_nsec=0},`

strace(1)

- man: "trace system calls and signals"
- Verwendet ptrace(2)

Syscall-Name (NR)

Syscall-ARGS

```
→ ~ strace sleep 5
execve("/usr/bin/sleep", ["sleep", "5"], 0xffffffff1e5f048 /* 25 vars */) = 0
brk(NULL)                                = 0xaaaaad0698000
faccessat(AT_FDCWD, "/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=29405, ...}) = 0
mmap(NULL, 29405, PROT_READ, MAP_PRIVATE, 3, 0) = 0xfffff88deb000
close(3)                                = 0
...
clock_nanosleep(CLOCK_REALTIME, 0, {tv_sec=5, tv_nsec=0},
```

strace(1)

- man: "trace system calls and signals"
- Verwendet ptrace(2)

Syscall-Name (NR)

Syscall-ARGS

Syscall-RTN-VAL

```
→ ~ strace sleep 5
execve("/usr/bin/sleep", ["sleep", "5"], 0xffffffff1e5f048 /* 25 vars */) = 0
brk(NULL)                                = 0xaaaaad0698000
faccessat(AT_FDCWD, "/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=29405, ...}) = 0
mmap(NULL, 29405, PROT_READ, MAP_PRIVATE, 3, 0) = 0xfffff88deb000
close(3)                                = 0
...
clock_nanosleep(CLOCK_REALTIME, 0, {tv_sec=5, tv_nsec=0},
```

strace(1)

- man: "trace system calls and signals"
- Verwendet ptrace(2)

Syscall-Name (NR)

Syscall-ARGS

Syscall-RTN-VAL

```
→ ~ strace sleep 5
execve("/usr/bin/sleep", ["sleep", "5"], 0xffffffff1e5f048 /* 25 vars */) = 0
brk(NULL)                                = 0xaaaaad0698000
faccessat(AT_FDCWD, "/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=29405, ...}) = 0
mmap(NULL, 29405, PROT_READ, MAP_PRIVATE, 3, 0) = 0xfffff88deb000
close(3)                                = 0
...
clock_nanosleep(CLOCK_REALTIME, 0, {tv_sec=5, tv_nsec=0},
```

- ministrace = Minimaler strace Nachbau

Syscall Parsing

Syscall Parsing

- Von Kernel Source:

Syscall Parsing

- Von Kernel Source:
- **Nr**, **Name**: tbl file

```
# <number> <abi> <name> ...  
0   common restart_syscall      sys_...  
1   common exit                 sys_exit  
2   common fork                 sys_fork  
3   common read                 sys_read  
...
```

Syscall Parsing

- Von Kernel Source:
- **Nr**, **Name**: tbl file

```
# <number> <abi> <name> ...  
0    common restart_syscall      sys_...  
1    common exit                  sys_exit  
2    common fork                  sys_fork  
3    common read                  sys_read  
...
```

Syscall Parsing

- Von Kernel Source:
- **Nr**, **Name**: tbl file

```
# <number> <abi> <name> ...  
0    common restart_syscall      sys_...  
1    common exit                  sys_exit  
2    common fork                  sys_fork  
3    common read                  sys_read  
...
```

Syscall Parsing

- Von Kernel Source:
- **Nr**, **Name**: tbl file
- **Args**: Macros

```
# <number> <abi> <name> ...  
0    common restart_syscall      sys_...  
1    common exit                  sys_exit  
2    common fork                  sys_fork  
3    common read                  sys_read  
...  
  
SYSCALL_DEFINE1(exit, int, error_code)
```

Syscall Parsing

- Von Kernel Source:
- **Nr**, **Name**: tbl file
- **Args**: Macros

```
# <number> <abi> <name> ...  
0    common restart_syscall      sys_...  
1    common exit                  sys_exit  
2    common fork                  sys_fork  
3    common read                  sys_read  
...  
  
SYSCALL_DEFINE1(exit, int, error_code)
```

Syscall Parsing

- Von Kernel Source:
- **Nr**, **Name**: tbl file
- **Args**: Macros
- Parsing script:

```
# <number> <abi> <name> ...  
0    common restart_syscall      sys_...  
1    common exit                  sys_exit  
2    common fork                  sys_fork  
3    common read                  sys_read  
...  
  
SYSCALL_DEFINE1(exit, int, error_code)
```


Syscall Parsing

- Von Kernel Source:
 - **Nr**, **Name**: tbl file
 - **Args**: Macros
- Parsing script:
 - Via RegEx

```
# <number> <abi> <name> ...  
0    common restart_syscall      sys_...  
1    common exit                  sys_exit  
2    common fork                  sys_fork  
3    common read                  sys_read  
...  
  
SYSCALL_DEFINE1(exit, int, error_code)
```

Syscall Parsing

- Von Kernel Source:
 - **Nr**, **Name**: tbl file
 - **Args**: Macros
- Parsing script:
 - Via RegEx
 - Output: Syscall table

```
# <number> <abi> <name> ...
0   common restart_syscall      sys_...
1   common exit                  sys_exit
2   common fork                  sys_fork
3   common read                  sys_read
...

SYSCALL_DEFINE1(exit, int, error_code)
```

Syscall Parsing

- Von Kernel Source:
 - **Nr**, **Name**: tbl file
 - **Args**: Macros
- Parsing script:
 - Via RegEx
 - Output: Syscall table
 - Platform specific

```
# <number> <abi> <name> ...
0    common restart_syscall      sys_...
1    common exit                  sys_exit
2    common fork                  sys_fork
3    common read                  sys_read
...

SYSCALL_DEFINE1(exit, int, error_code)
```

Syscall Parsing

- Von Kernel Source:
 - **Nr**, **Name**: tbl file
 - **Args**: Macros
- Parsing script:
 - Via RegEx
 - Output: Syscall table
 - Platform specific

```
# <number> <abi> <name> ...
0   common restart_syscall      sys_...
1   common exit                 sys_exit
2   common fork                 sys_fork
3   common read                 sys_read
...

SYSCALL_DEFINE1(exit, int, error_code)
```

```
// kernel version 5.4.0-97-generic on aarch64

const syscall_entry syscalls[] = {
    [__SNR_restart_syscall] = {
        .name   = "restart_syscall",
        .nargs  = 0,
        .args   = {-1, -1, -1, -1, -1, -1}},
    [__SNR_exit] = {
        .name   = "exit",
        .nargs  = 1,
        .args   = {ARG_INT, -1, -1, -1, -1, -1}},

```

Syscall Parsing

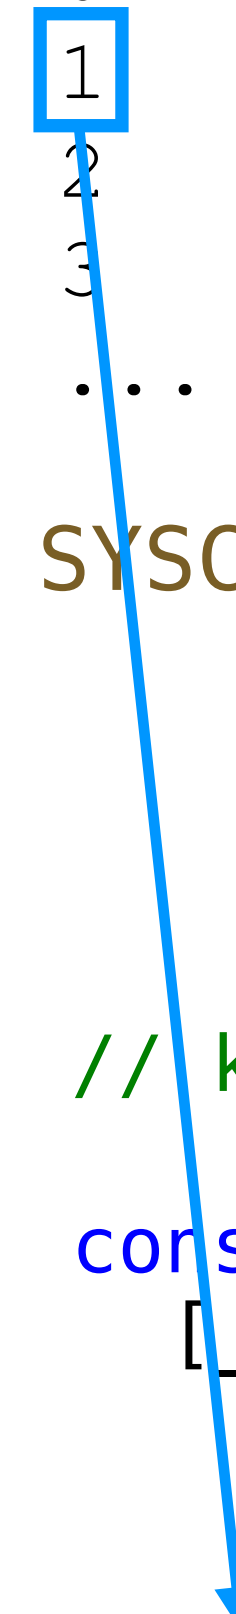
- Von Kernel Source:
- **Nr**, **Name**: tbl file
- **Args**: Macros
- Parsing script:
 - Via RegEx
 - Output: Syscall table
 - Platform specific

```
# <number> <abi> <name> ...
0   common restart_syscall      sys_...
1   common exit                 sys_exit
2   common fork                 sys_fork
3   common read                 sys_read
...

SYSCALL_DEFINE1(exit, int, error_code)

// kernel version 5.4.0-97-generic on aarch64
const syscall_entry syscalls[] = {
    [__SNR_restart_syscall] = {
        .name = "restart_syscall",
        .nargs = 0,
        .args = {-1, -1, -1, -1, -1, -1}},
    [__SNR_exit] = {
        .name = "exit",
        .nargs = 1,
        .args = {ARG_INT, -1, -1, -1, -1, -1}},

```



Syscall Parsing

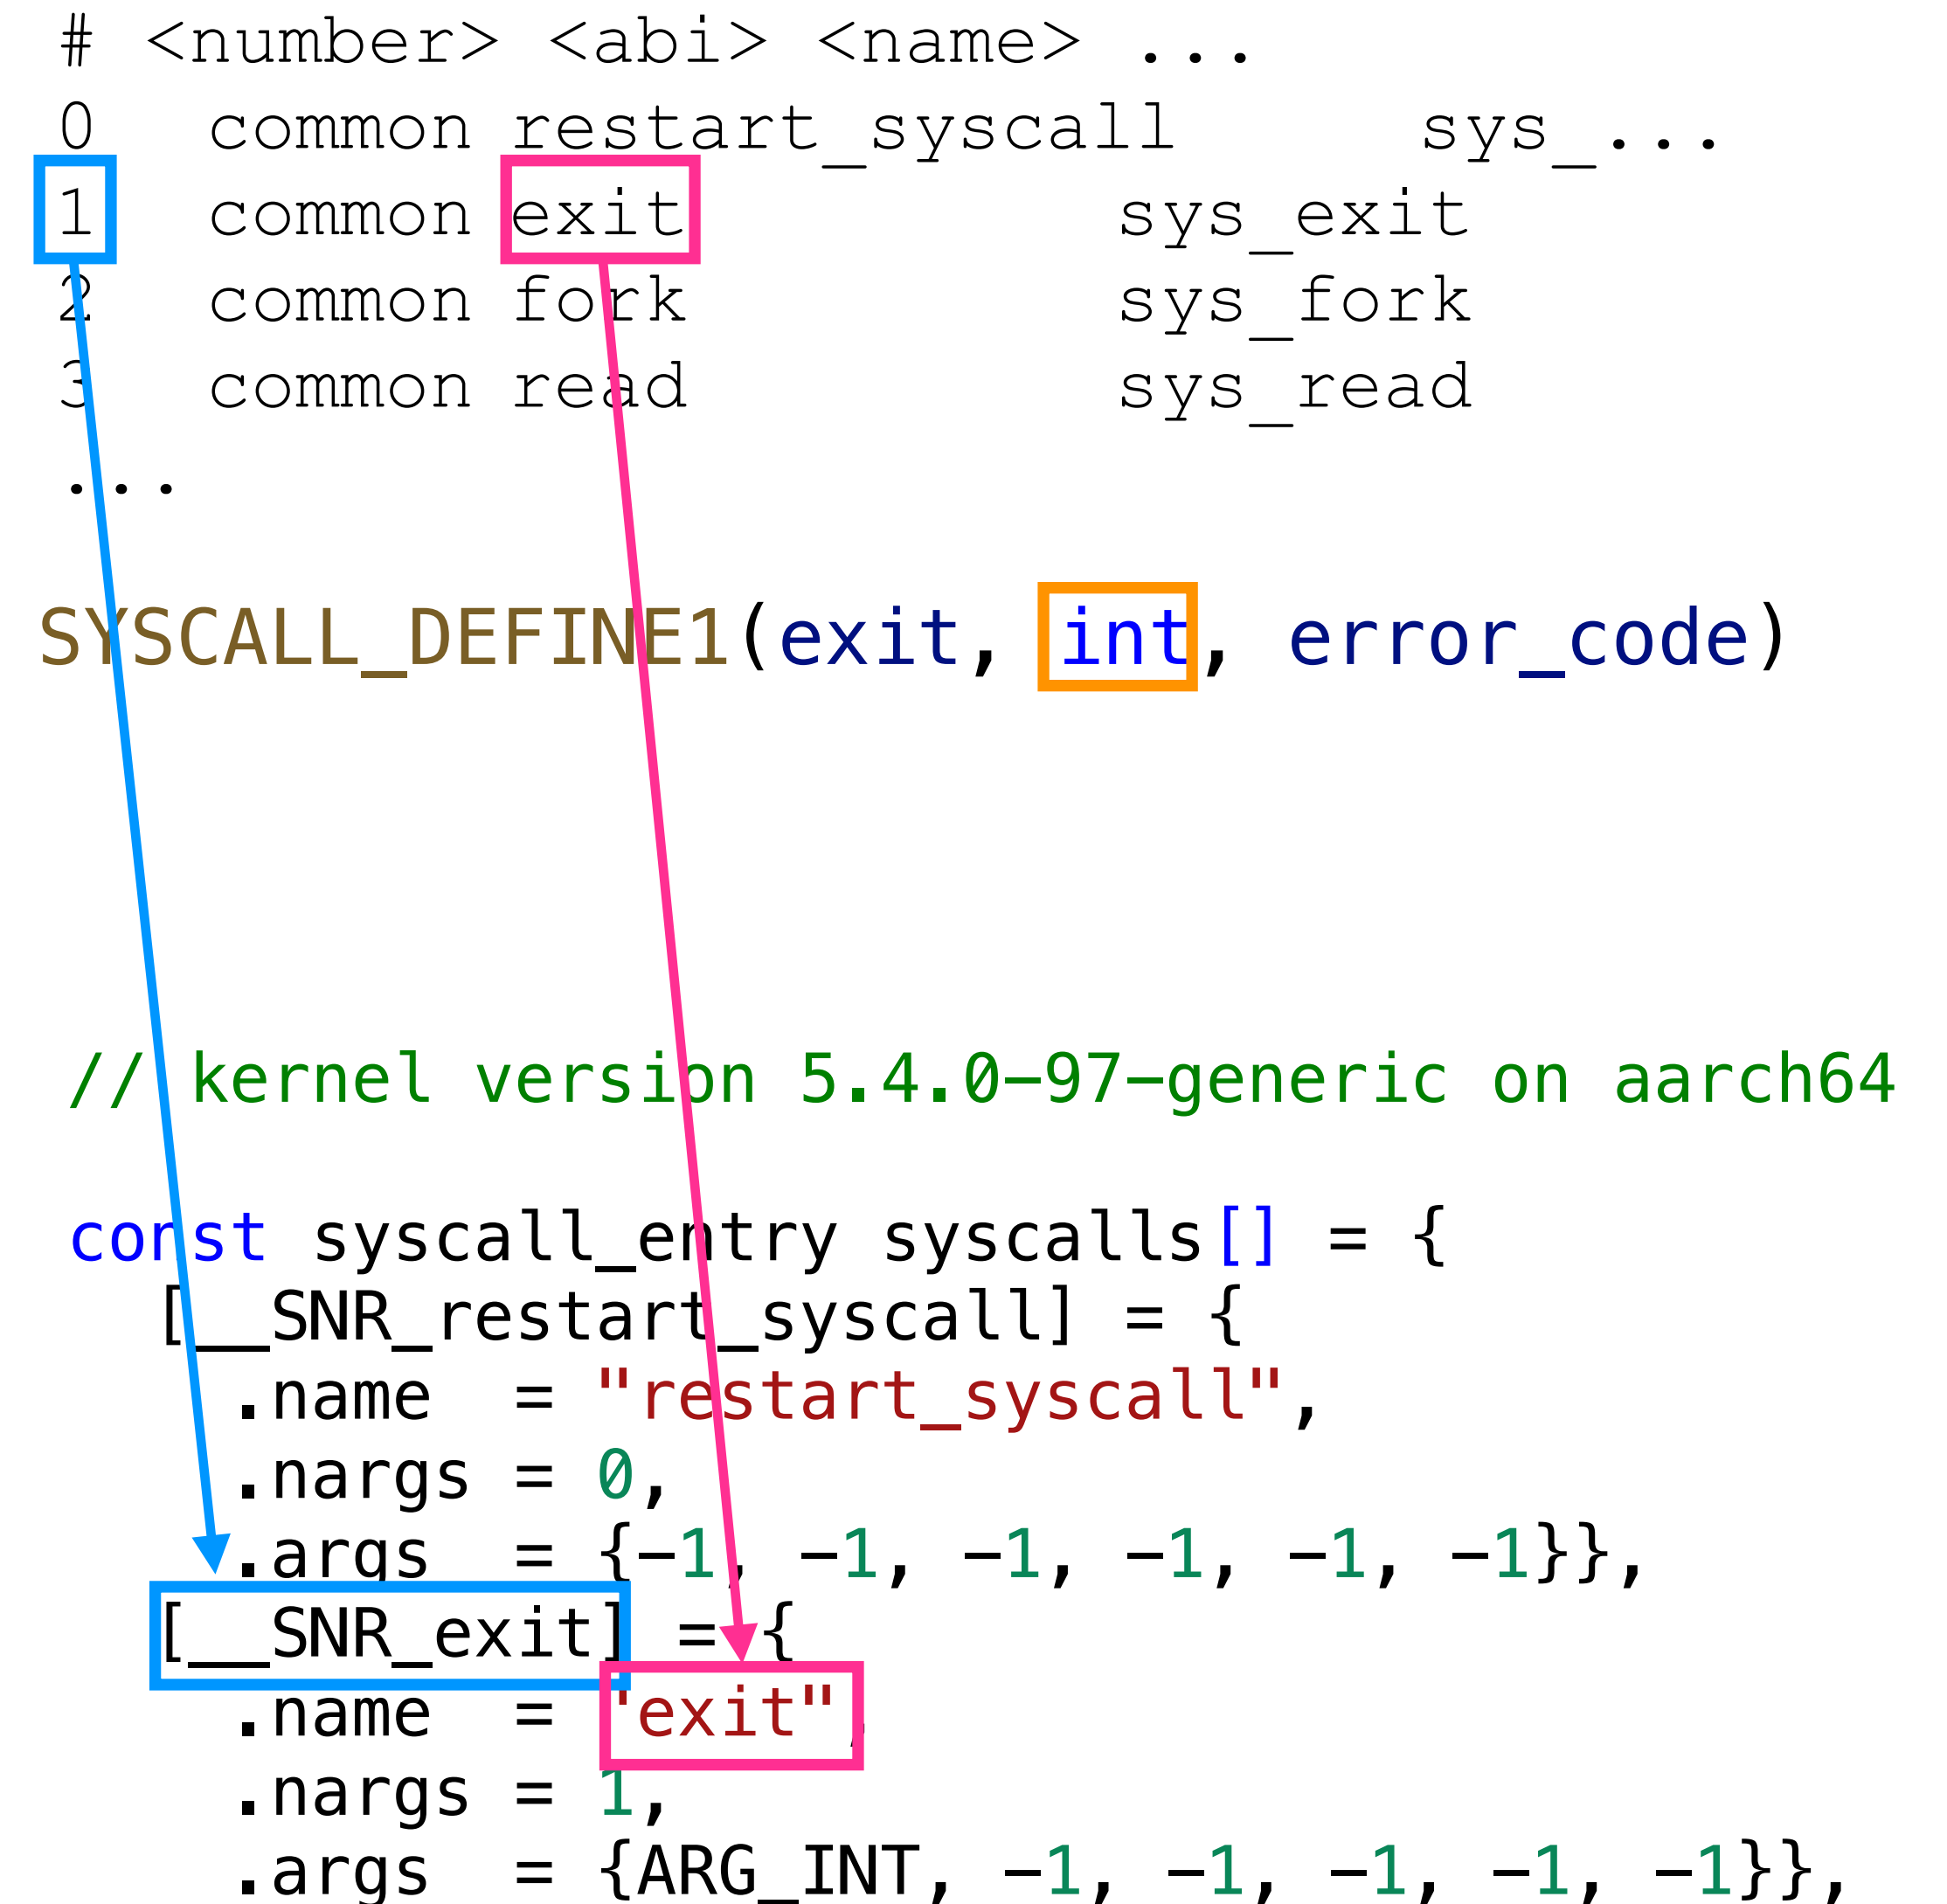
- Von Kernel Source:
- **Nr**, **Name**: tbl file
- **Args**: Macros
- Parsing script:
- Via RegEx
- Output: Syscall table
 - Platform specific

```
# <number> <abi> <name> ...
0   common restart_syscall      sys_...
1   common exit                 sys_exit
2   common fork                 sys_fork
3   common read                 sys_read
...

SYSCALL_DEFINE1(exit, int, error_code)

// kernel version 5.4.0-97-generic on aarch64
const syscall_entry syscalls[] = {
    [__SNR_restart_syscall] = {
        .name = "restart_syscall",
        .nargs = 0,
        .args = {-1, -1, -1, -1, -1, -1}},
    [__SNR_exit] = {
        .name = "exit",
        .nargs = 1,
        .args = {ARG_INT, -1, -1, -1, -1, -1}},

```



Syscall Parsing

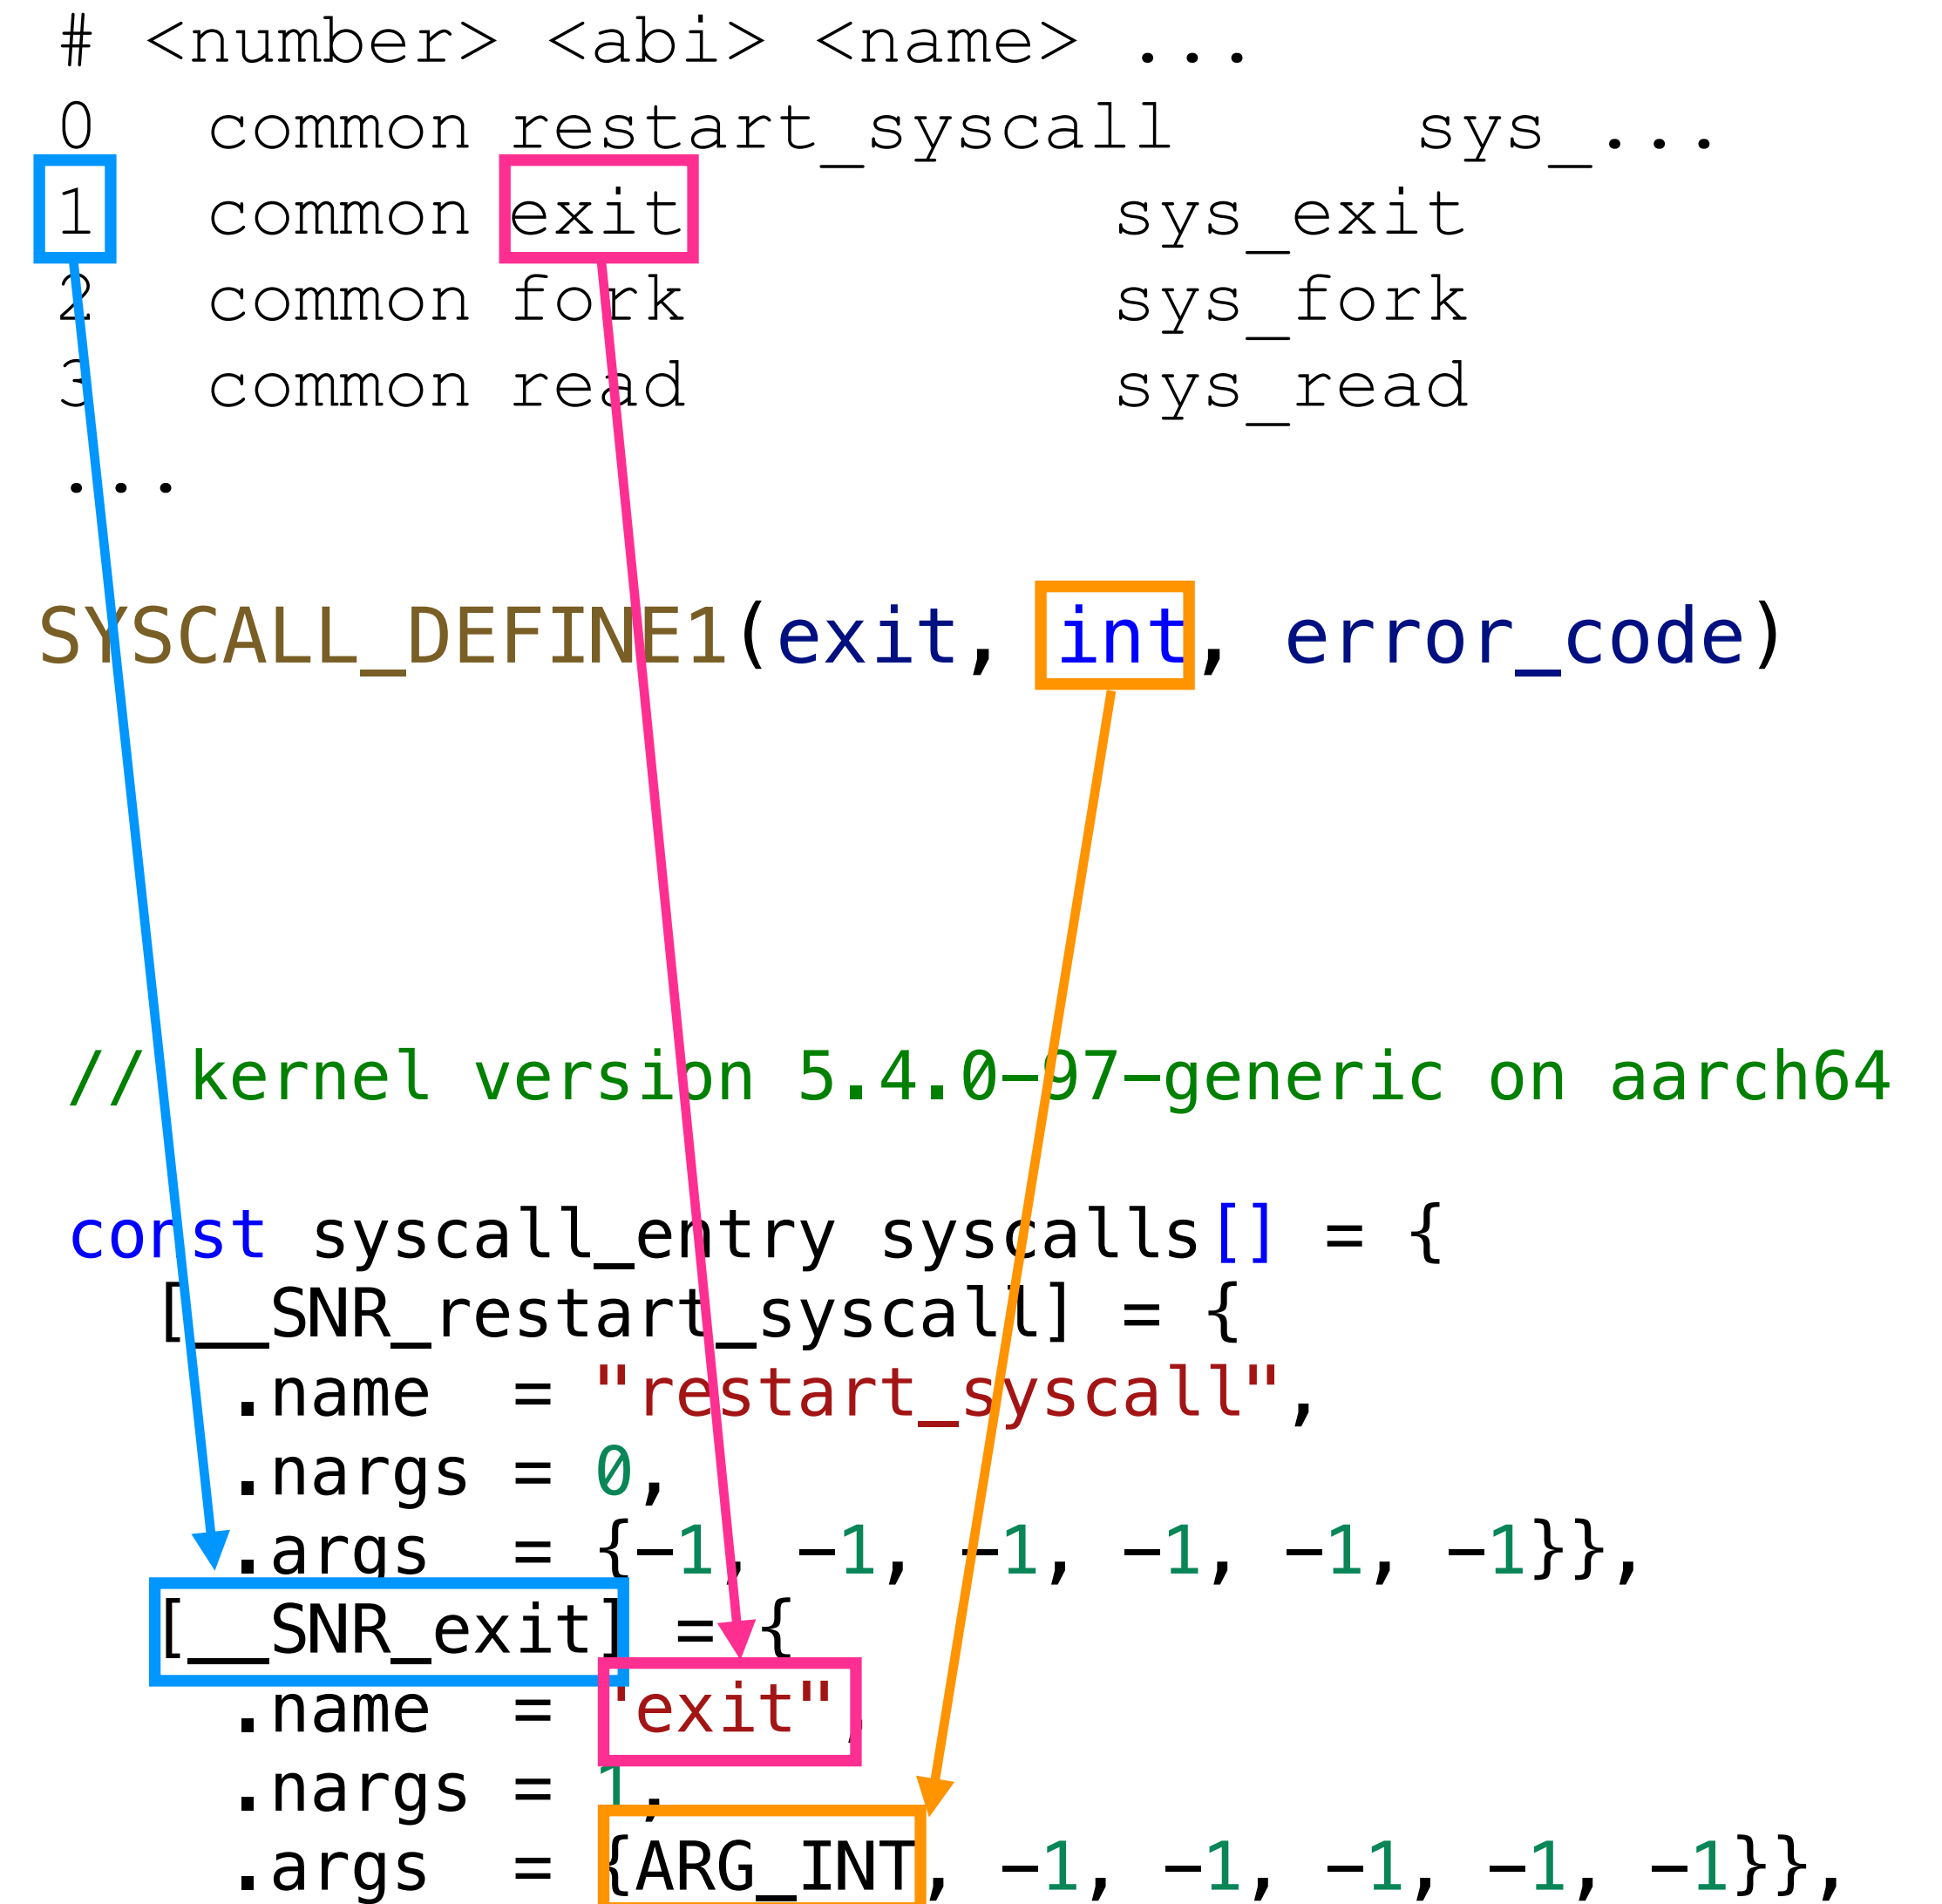
- Von Kernel Source:
- **Nr**, **Name**: tbl file
- **Args**: Macros
- Parsing script:
 - Via RegEx
 - Output: Syscall table
 - Platform specific

```
# <number> <abi> <name> ...
0   common restart_syscall      sys_...
1   common exit                 sys_exit
2   common fork                 sys_fork
3   common read                 sys_read
...

SYSCALL_DEFINE1(exit, int, error_code)

// kernel version 5.4.0-97-generic on aarch64
const syscall_entry syscalls[] = {
    [__SNR_restart_syscall] = {
        .name = "restart_syscall",
        .nargs = 0,
        .args = {-1, -1, -1, -1, -1, -1}},
    [__SNR_exit] = {
        .name = "exit",
        .nargs = 1,
        .args = {ARG_INT, -1, -1, -1, -1, -1}},

```



Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"

Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



Parent


Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



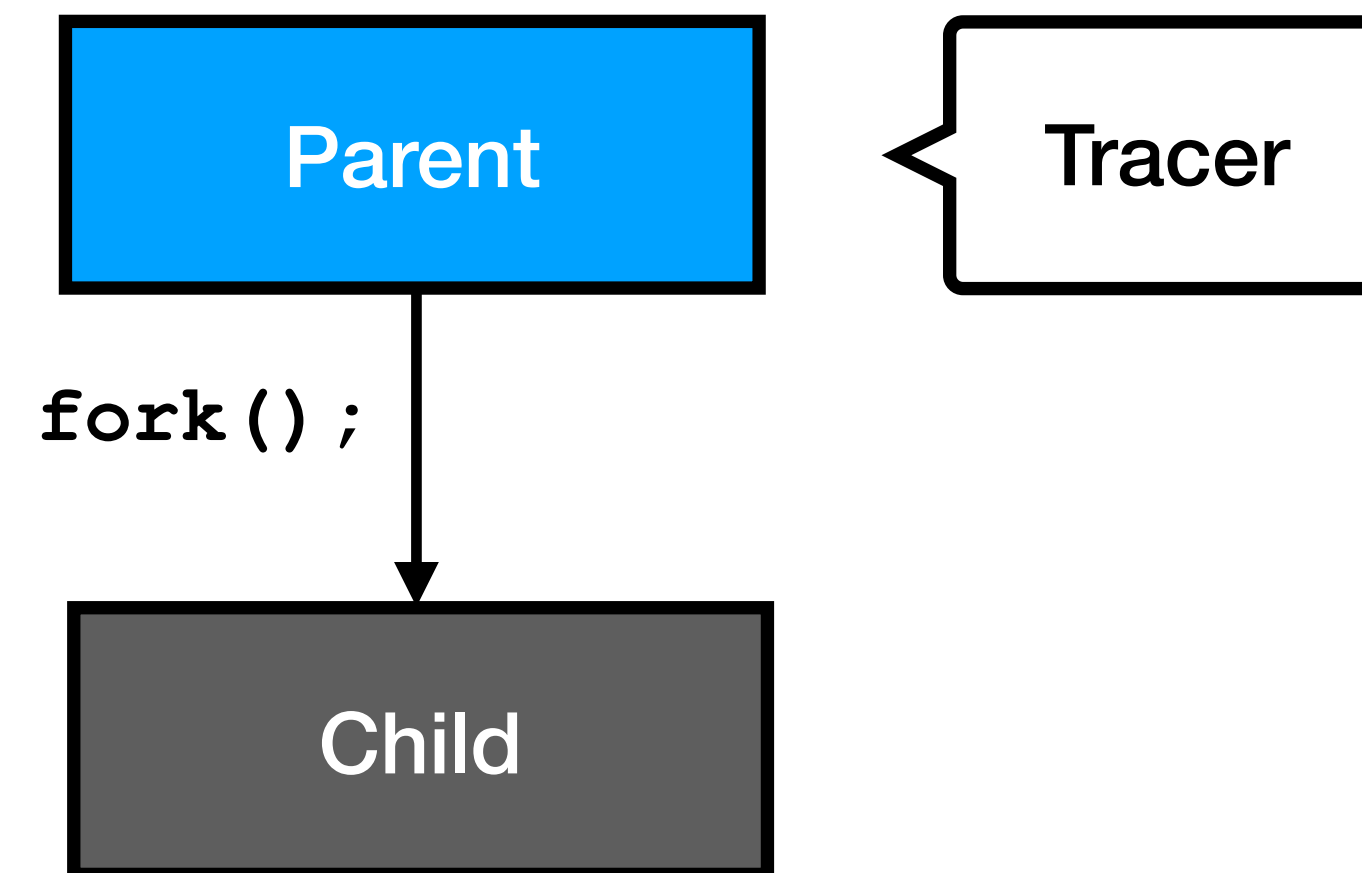
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



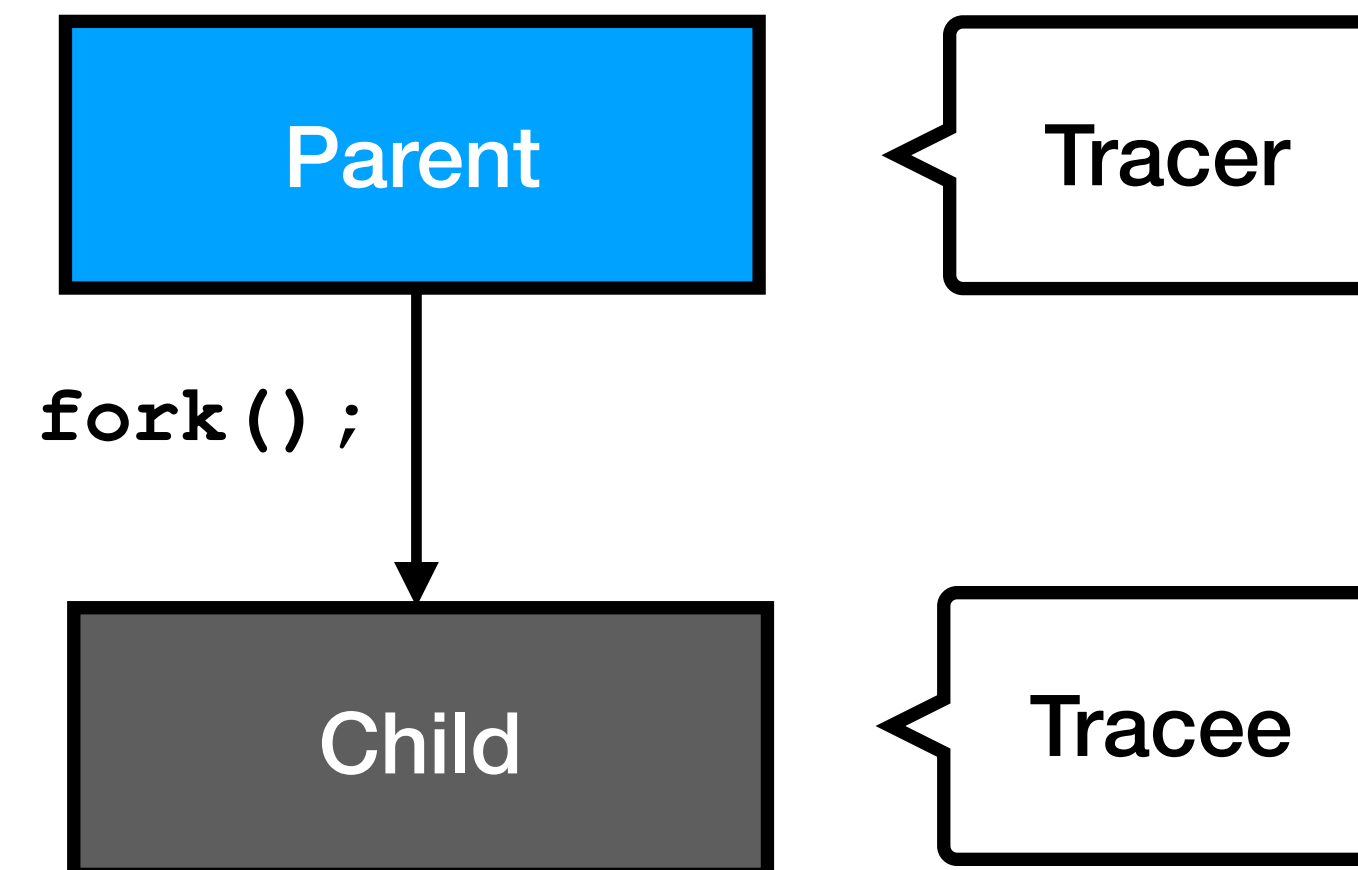
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



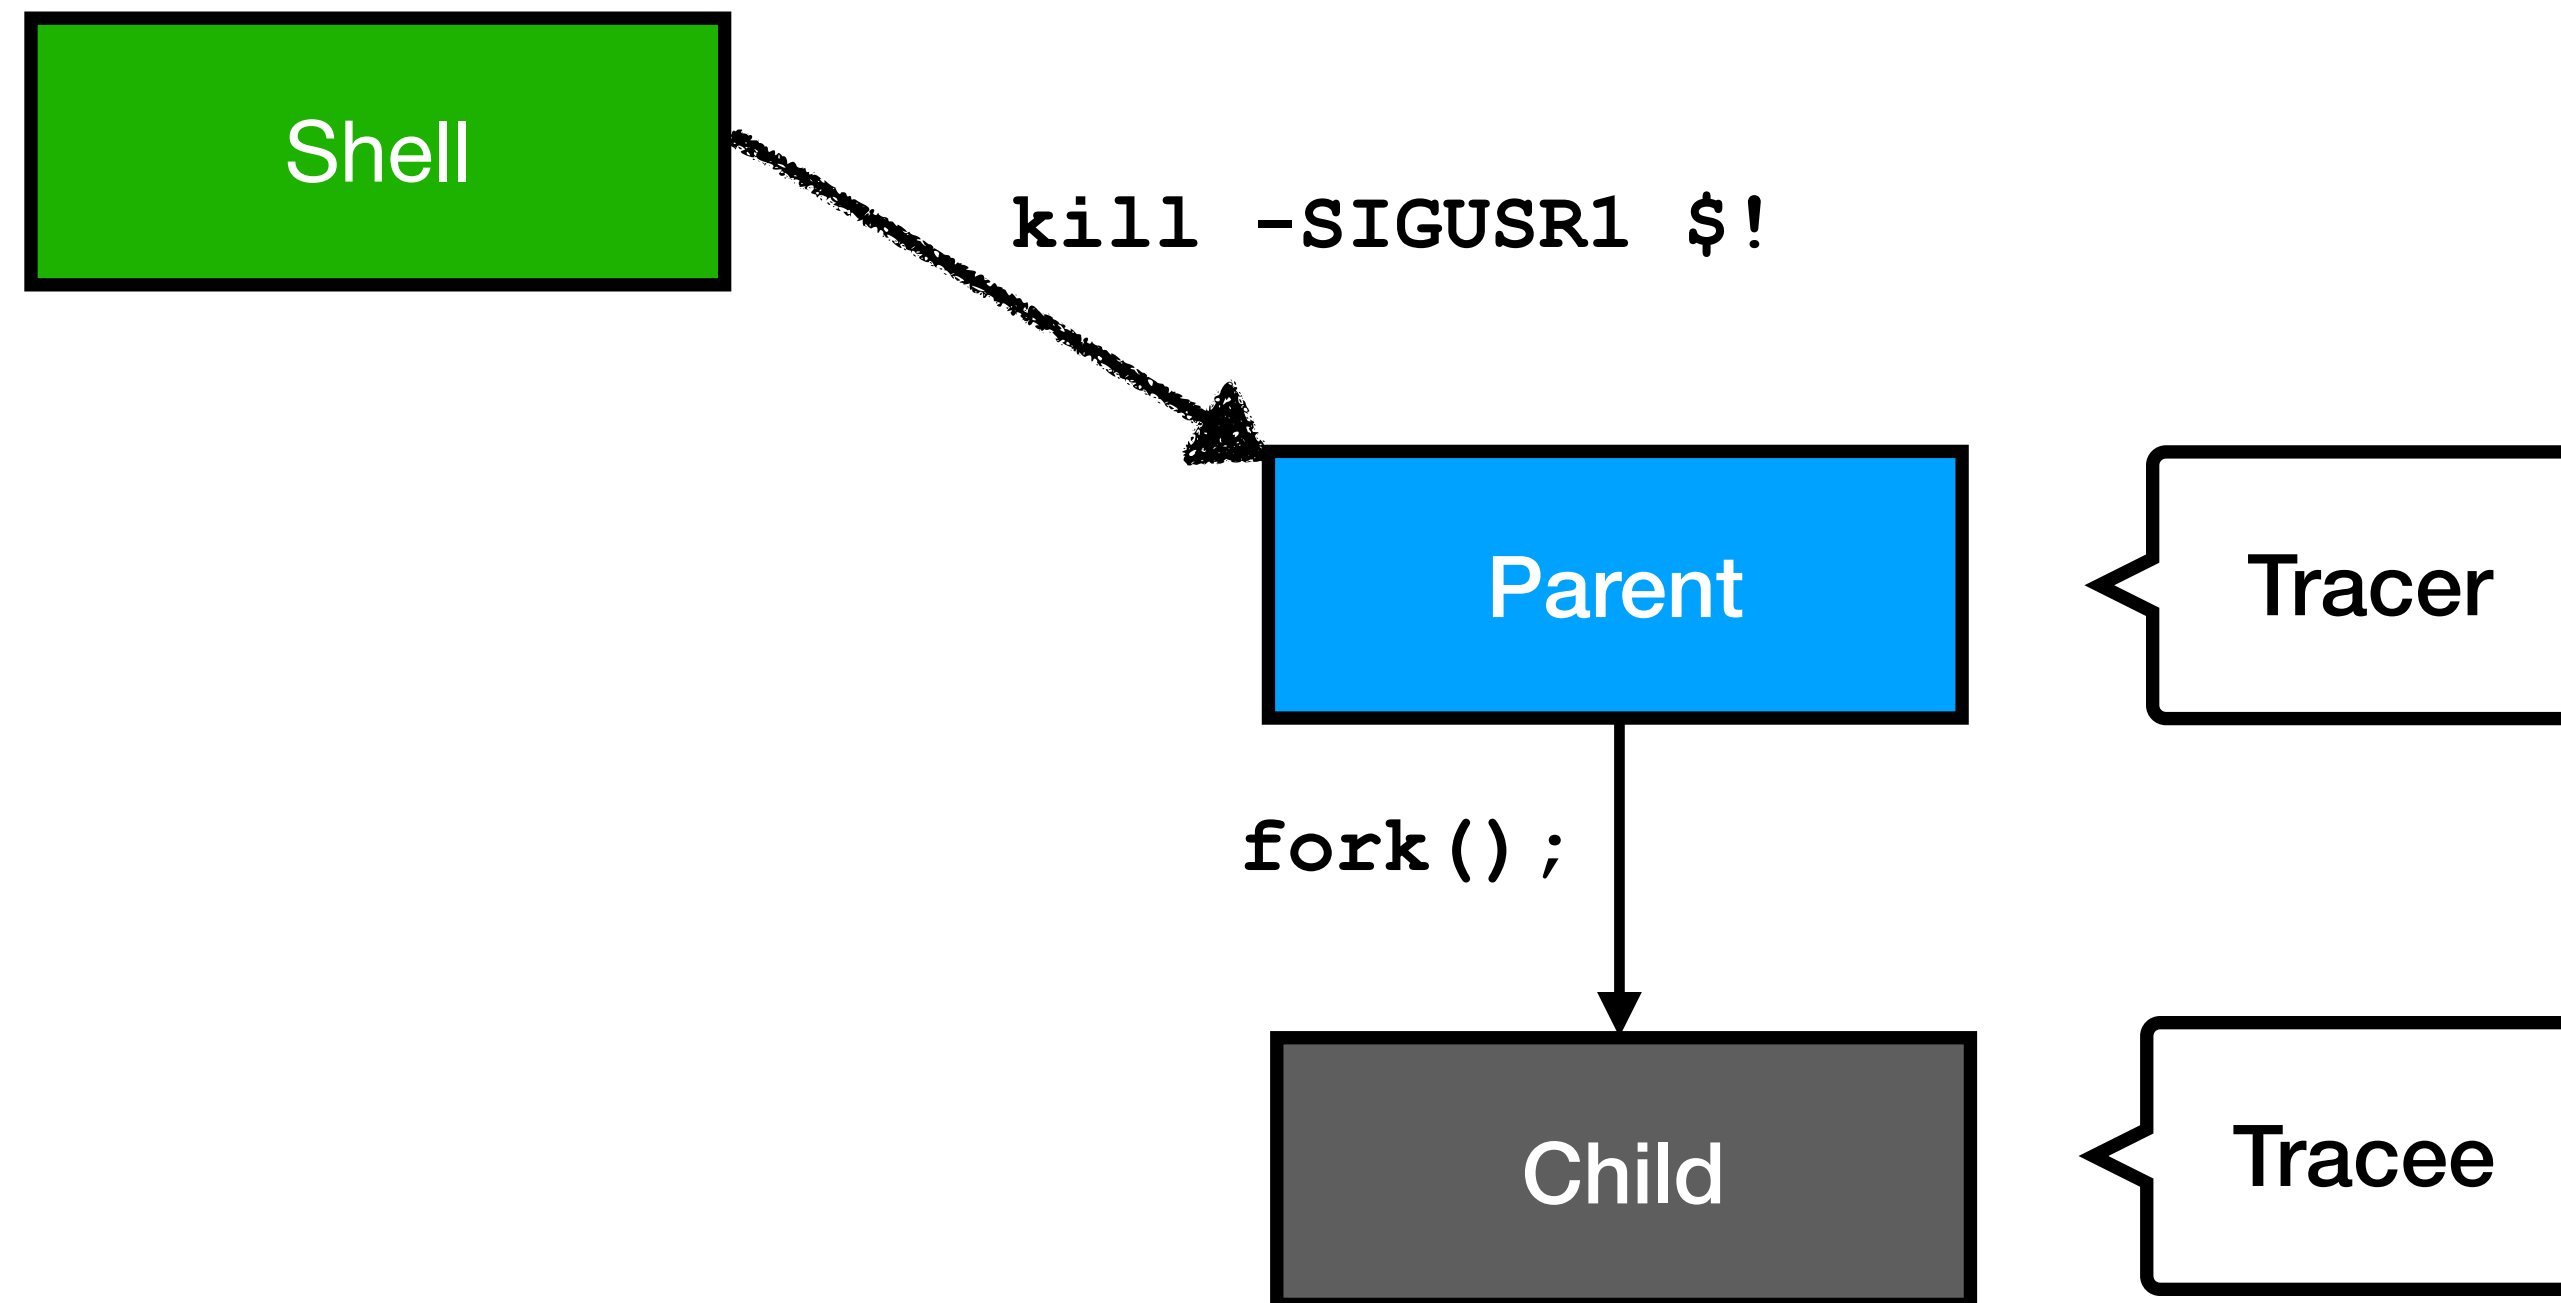
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



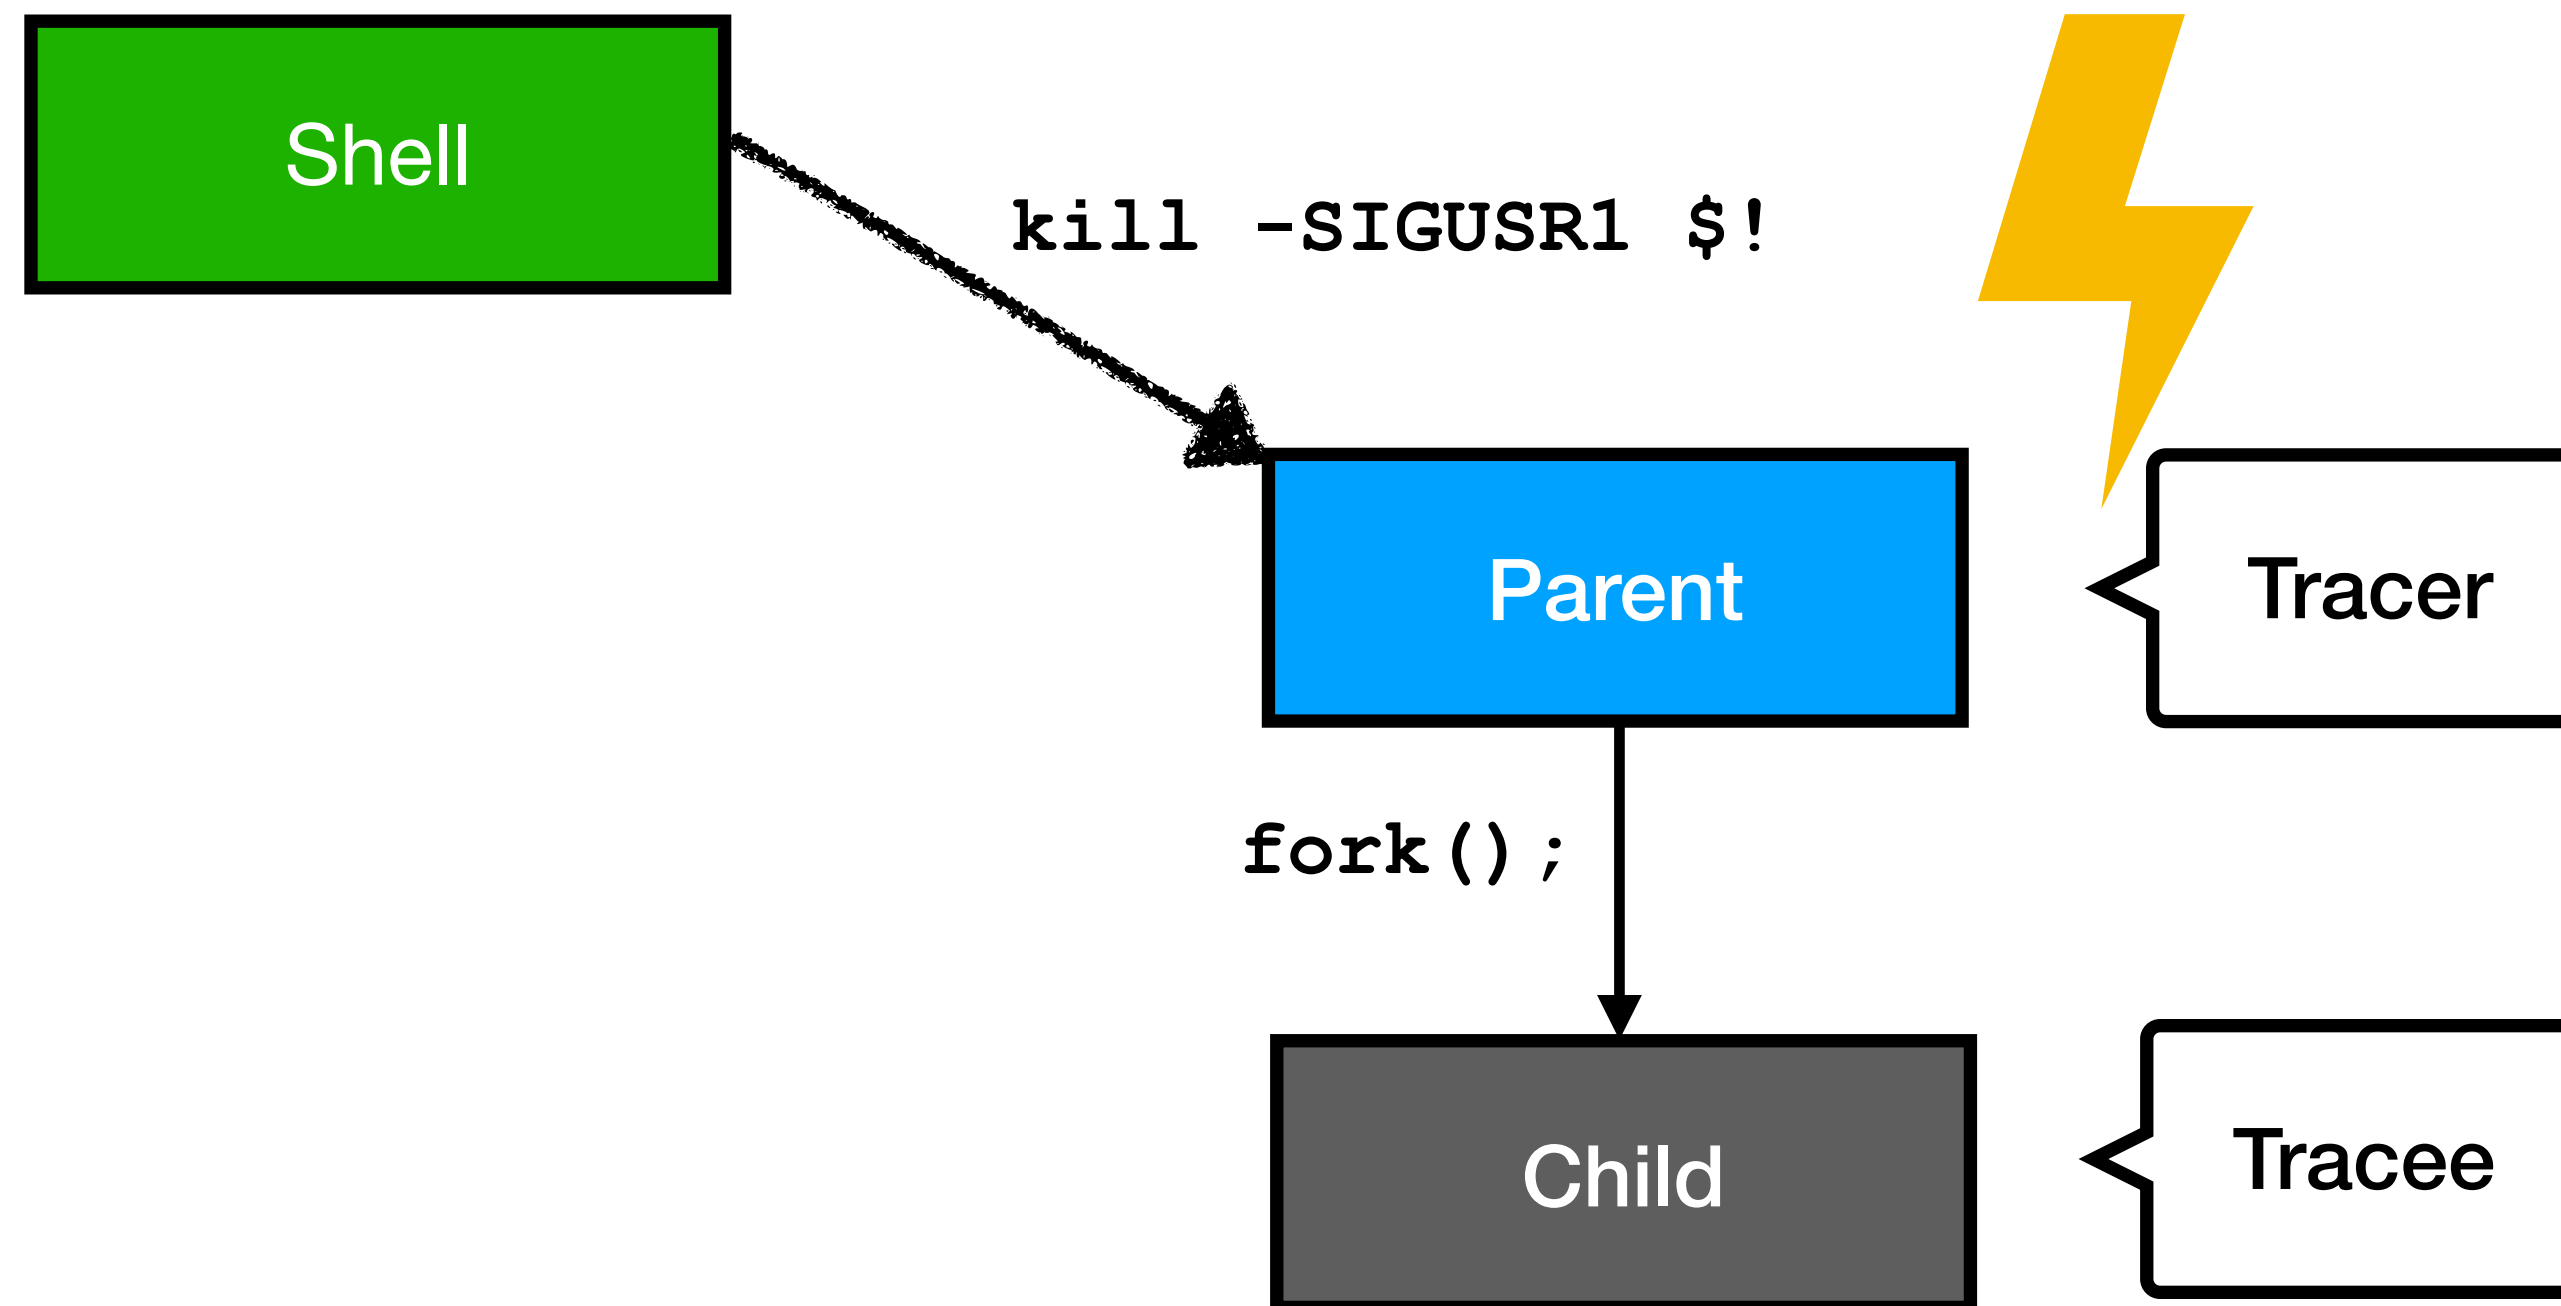
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



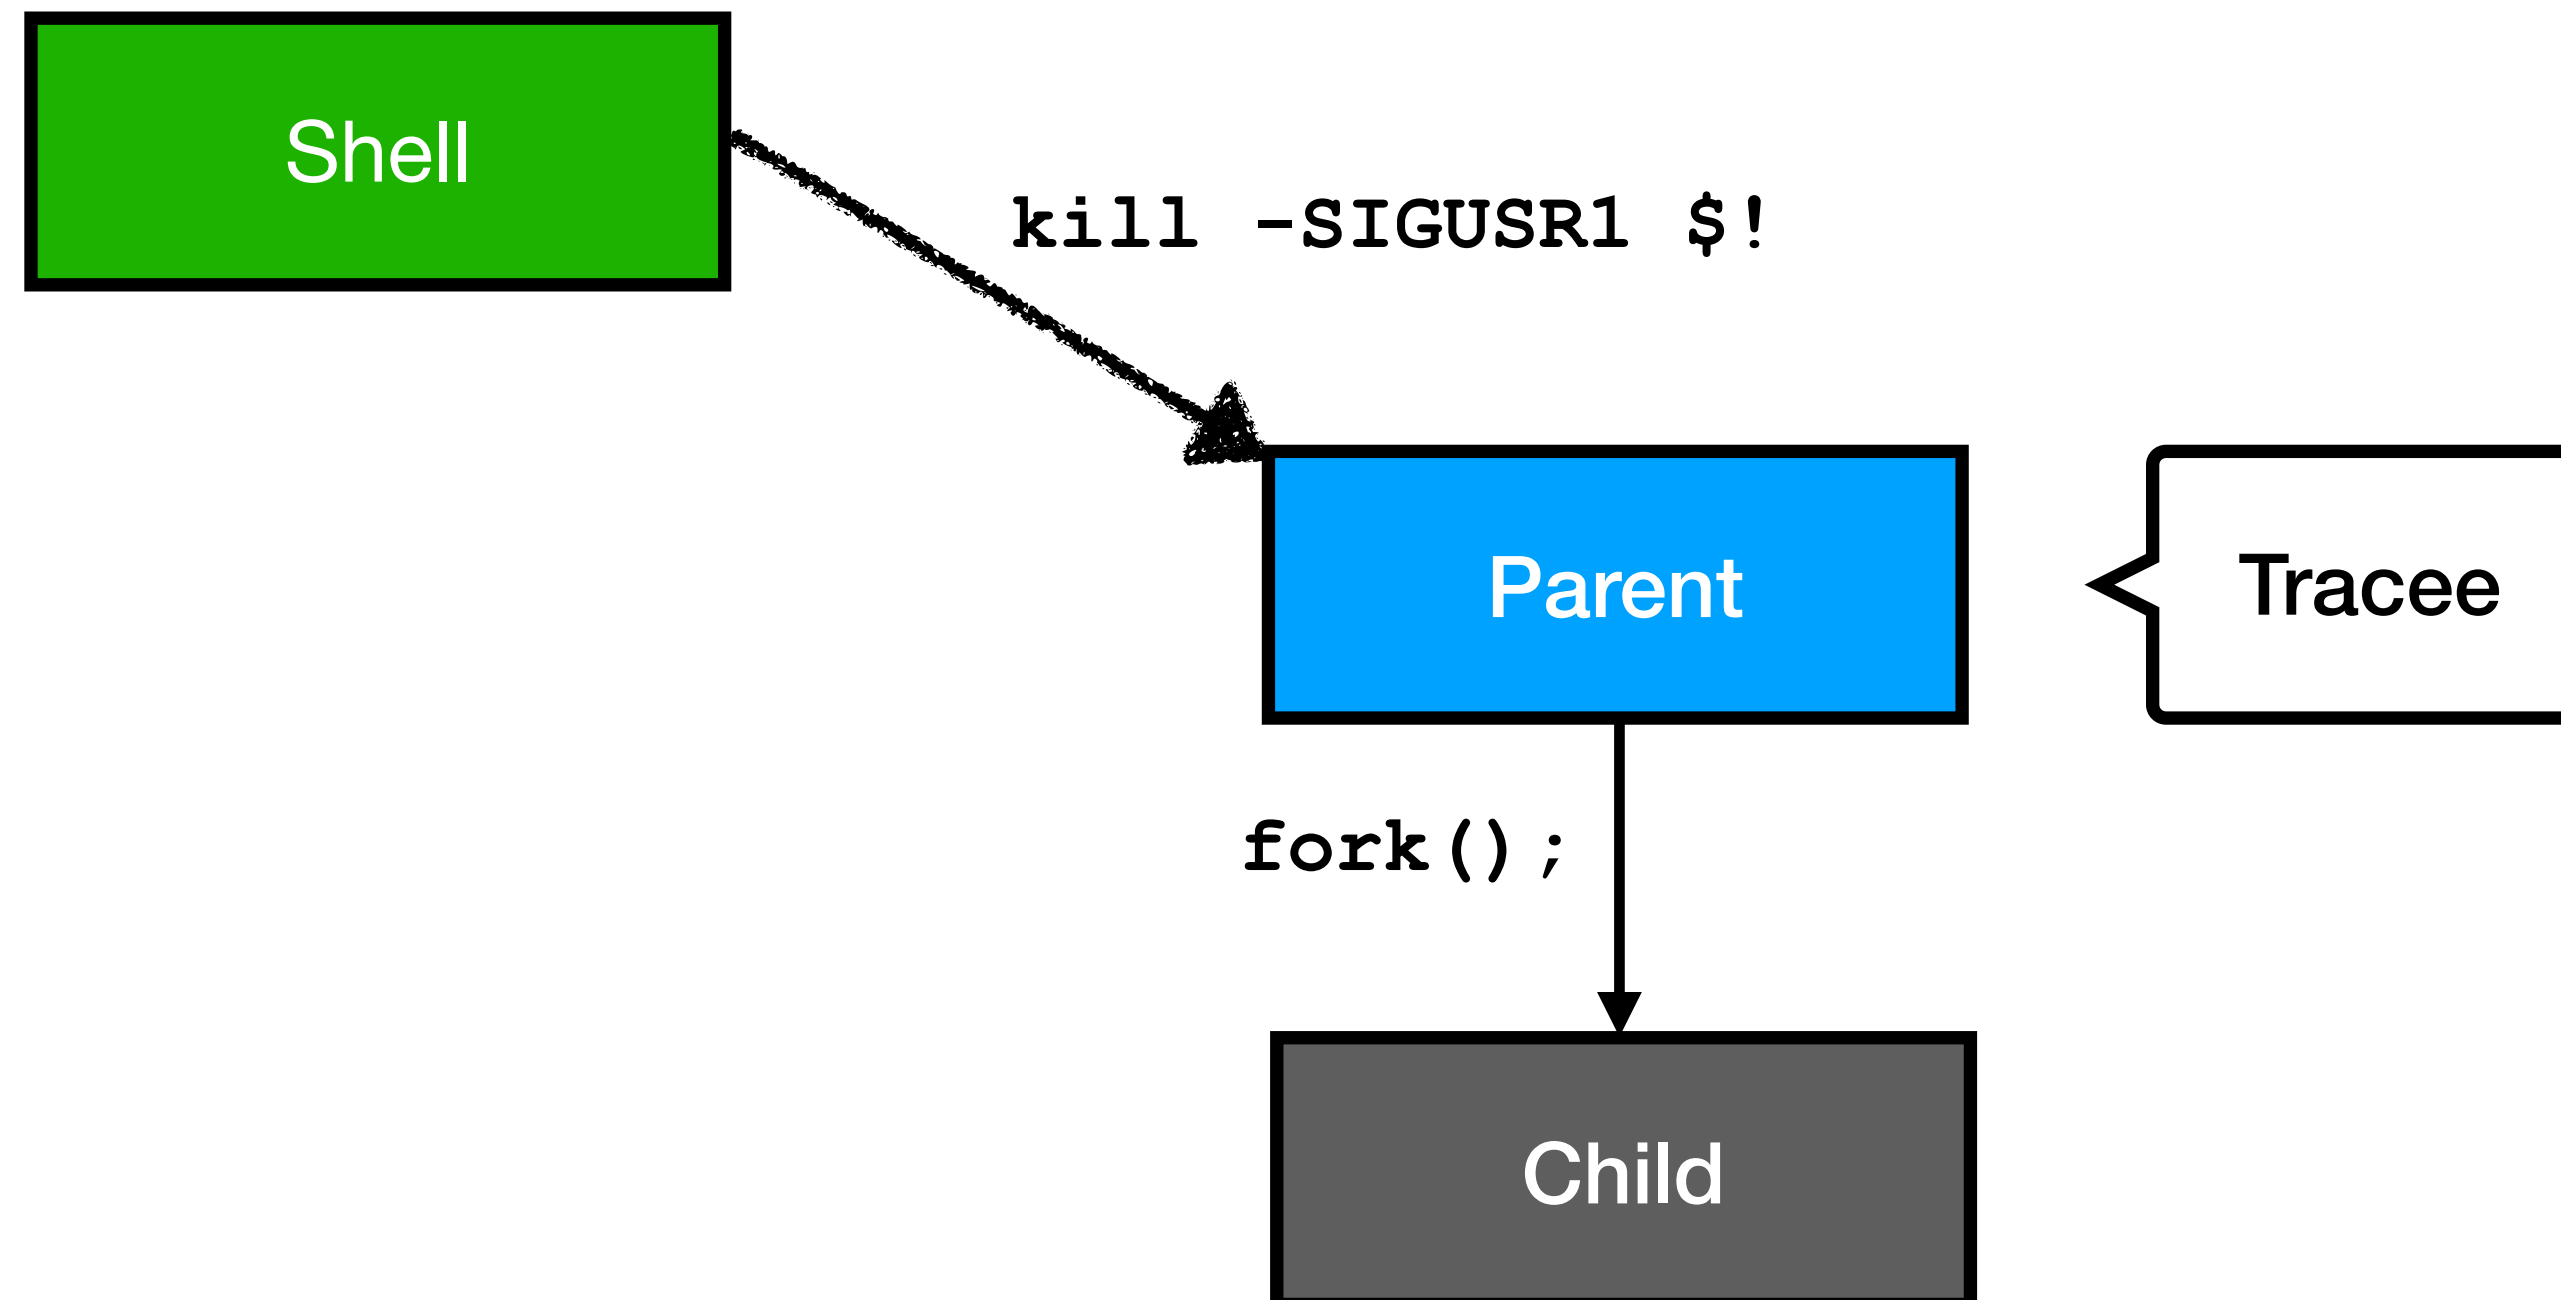
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



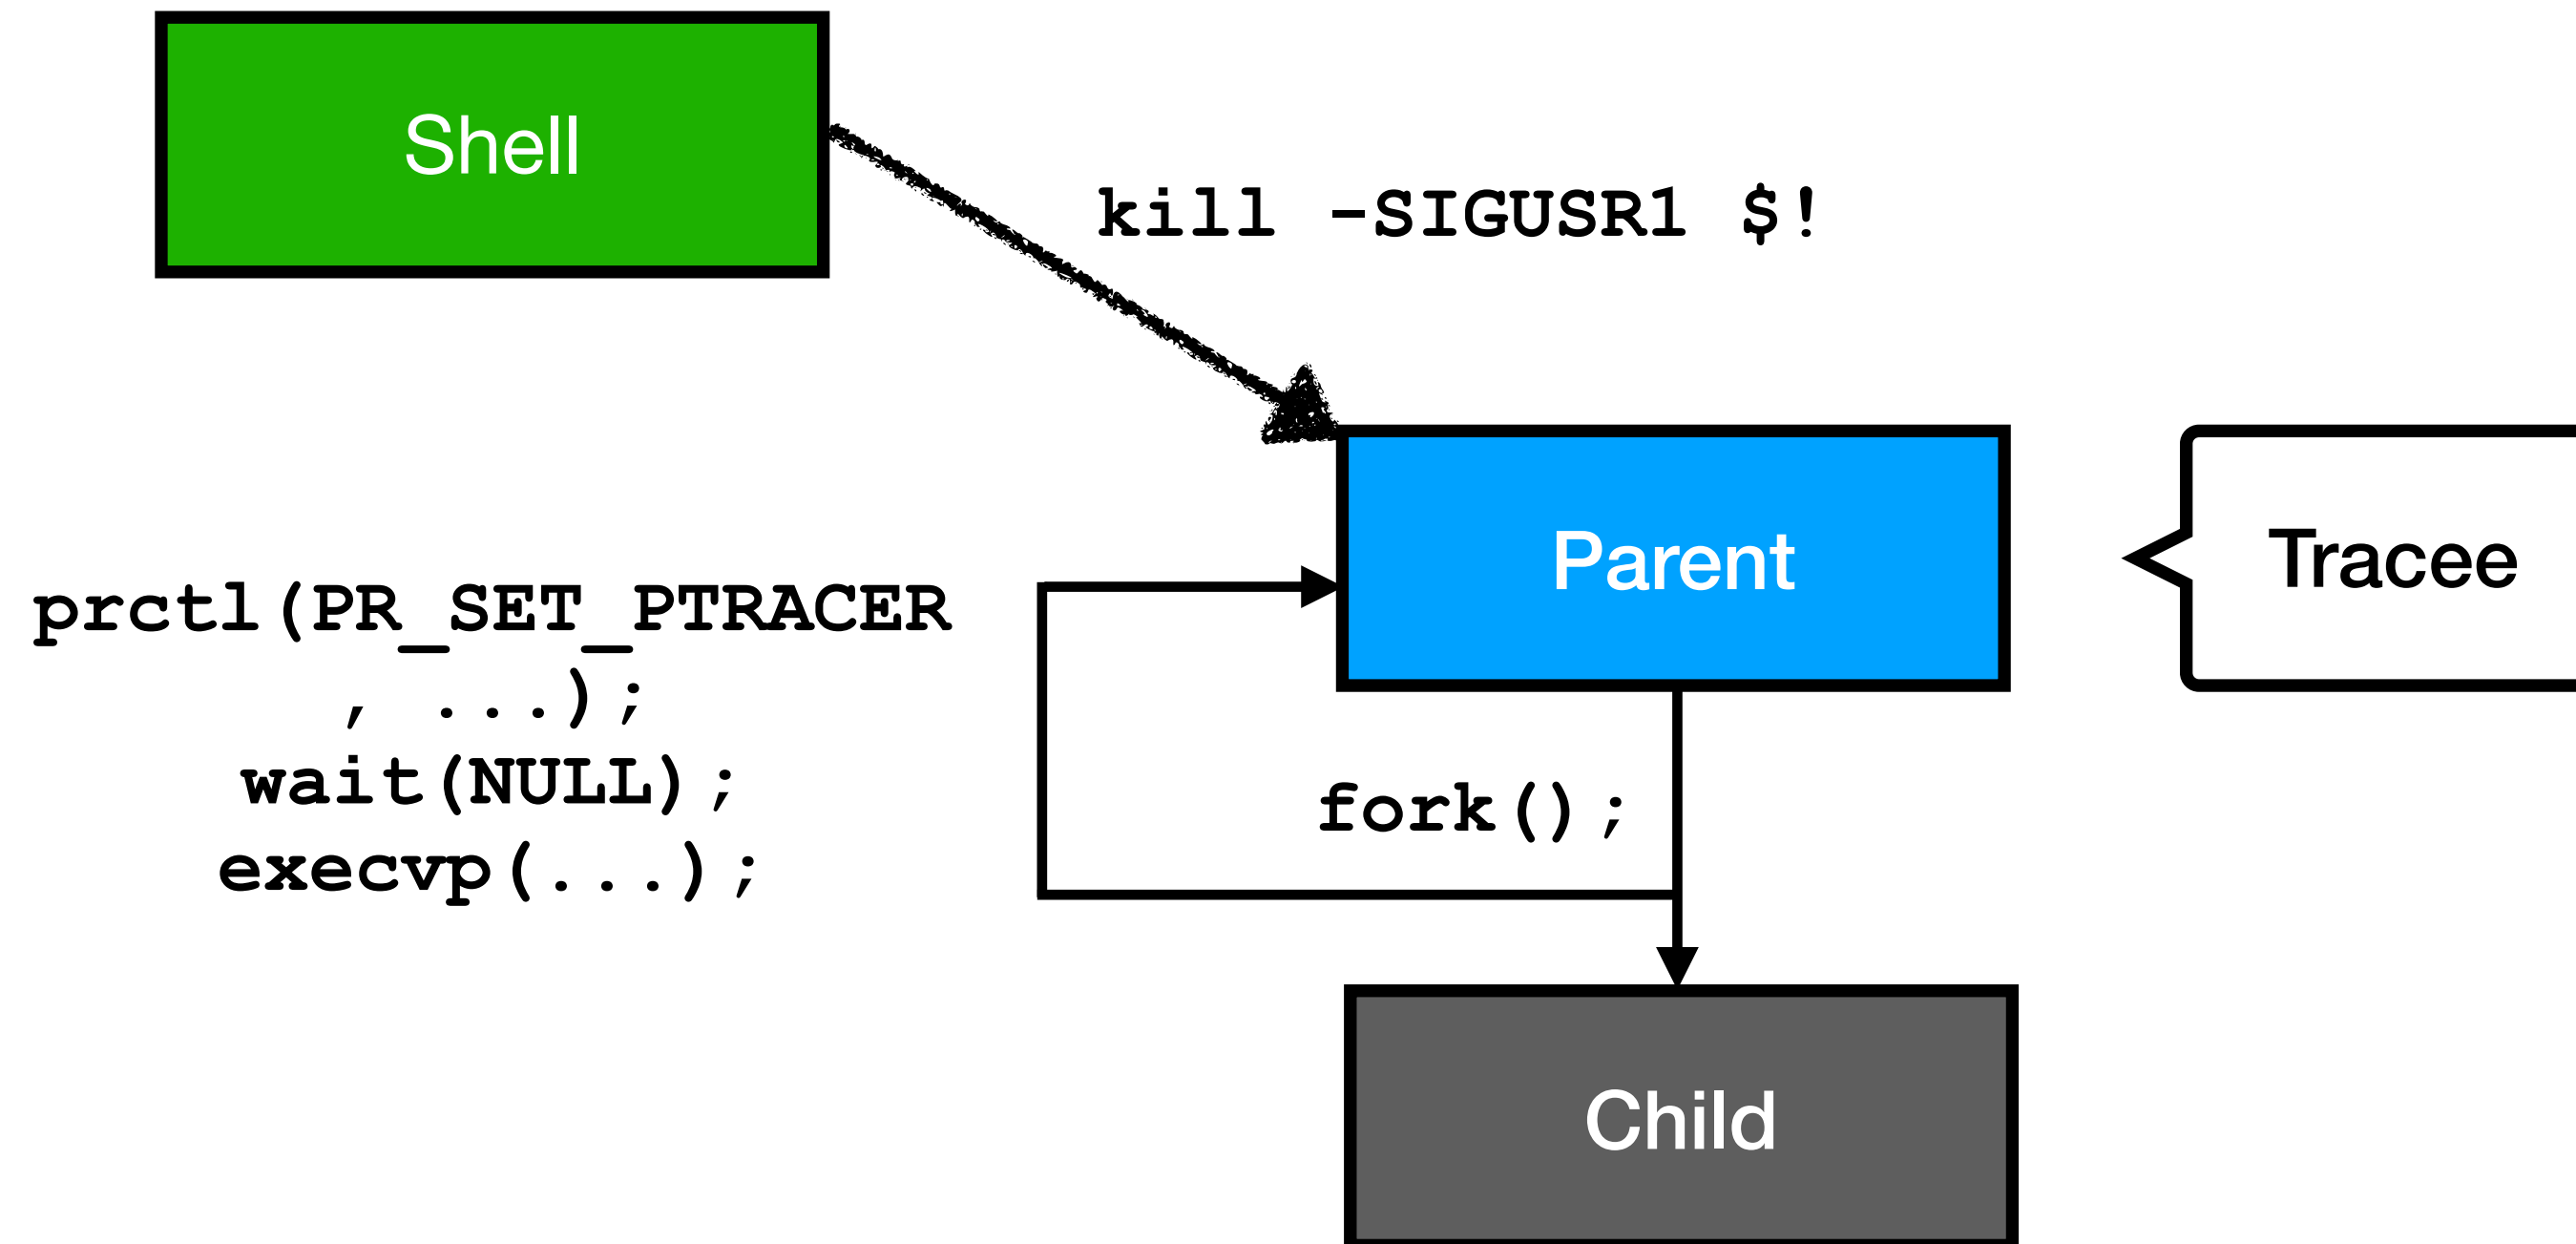
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



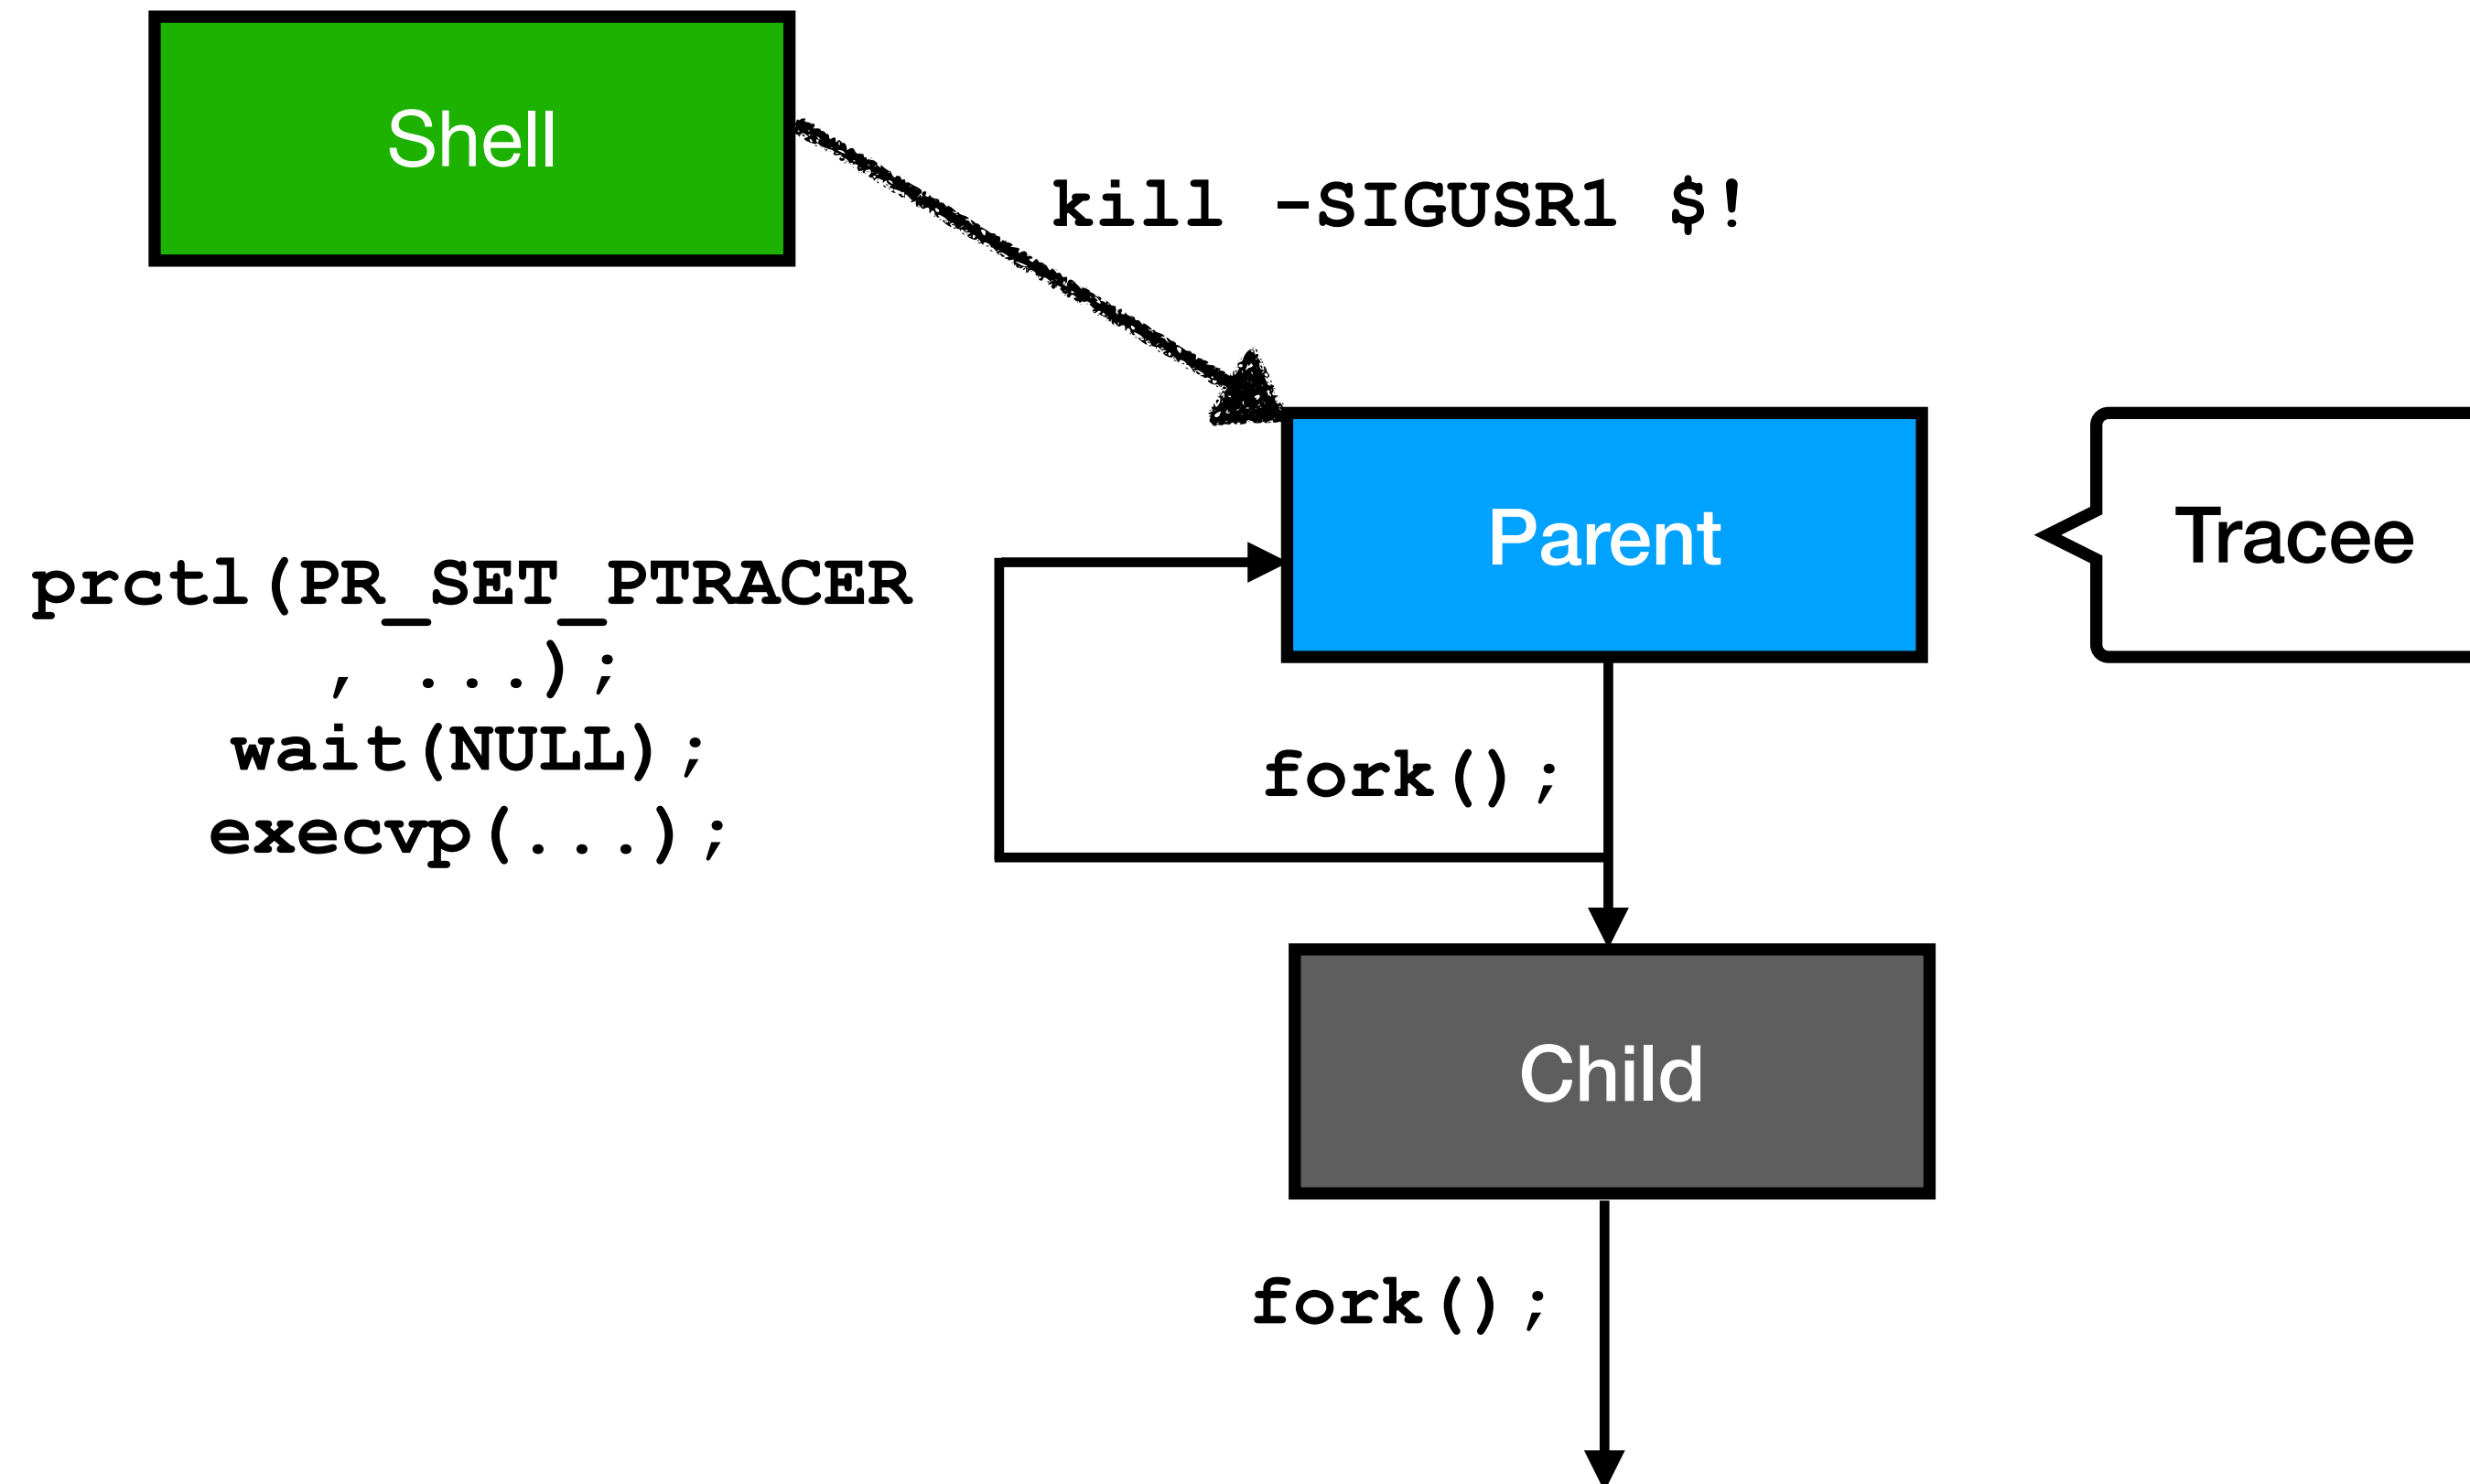
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



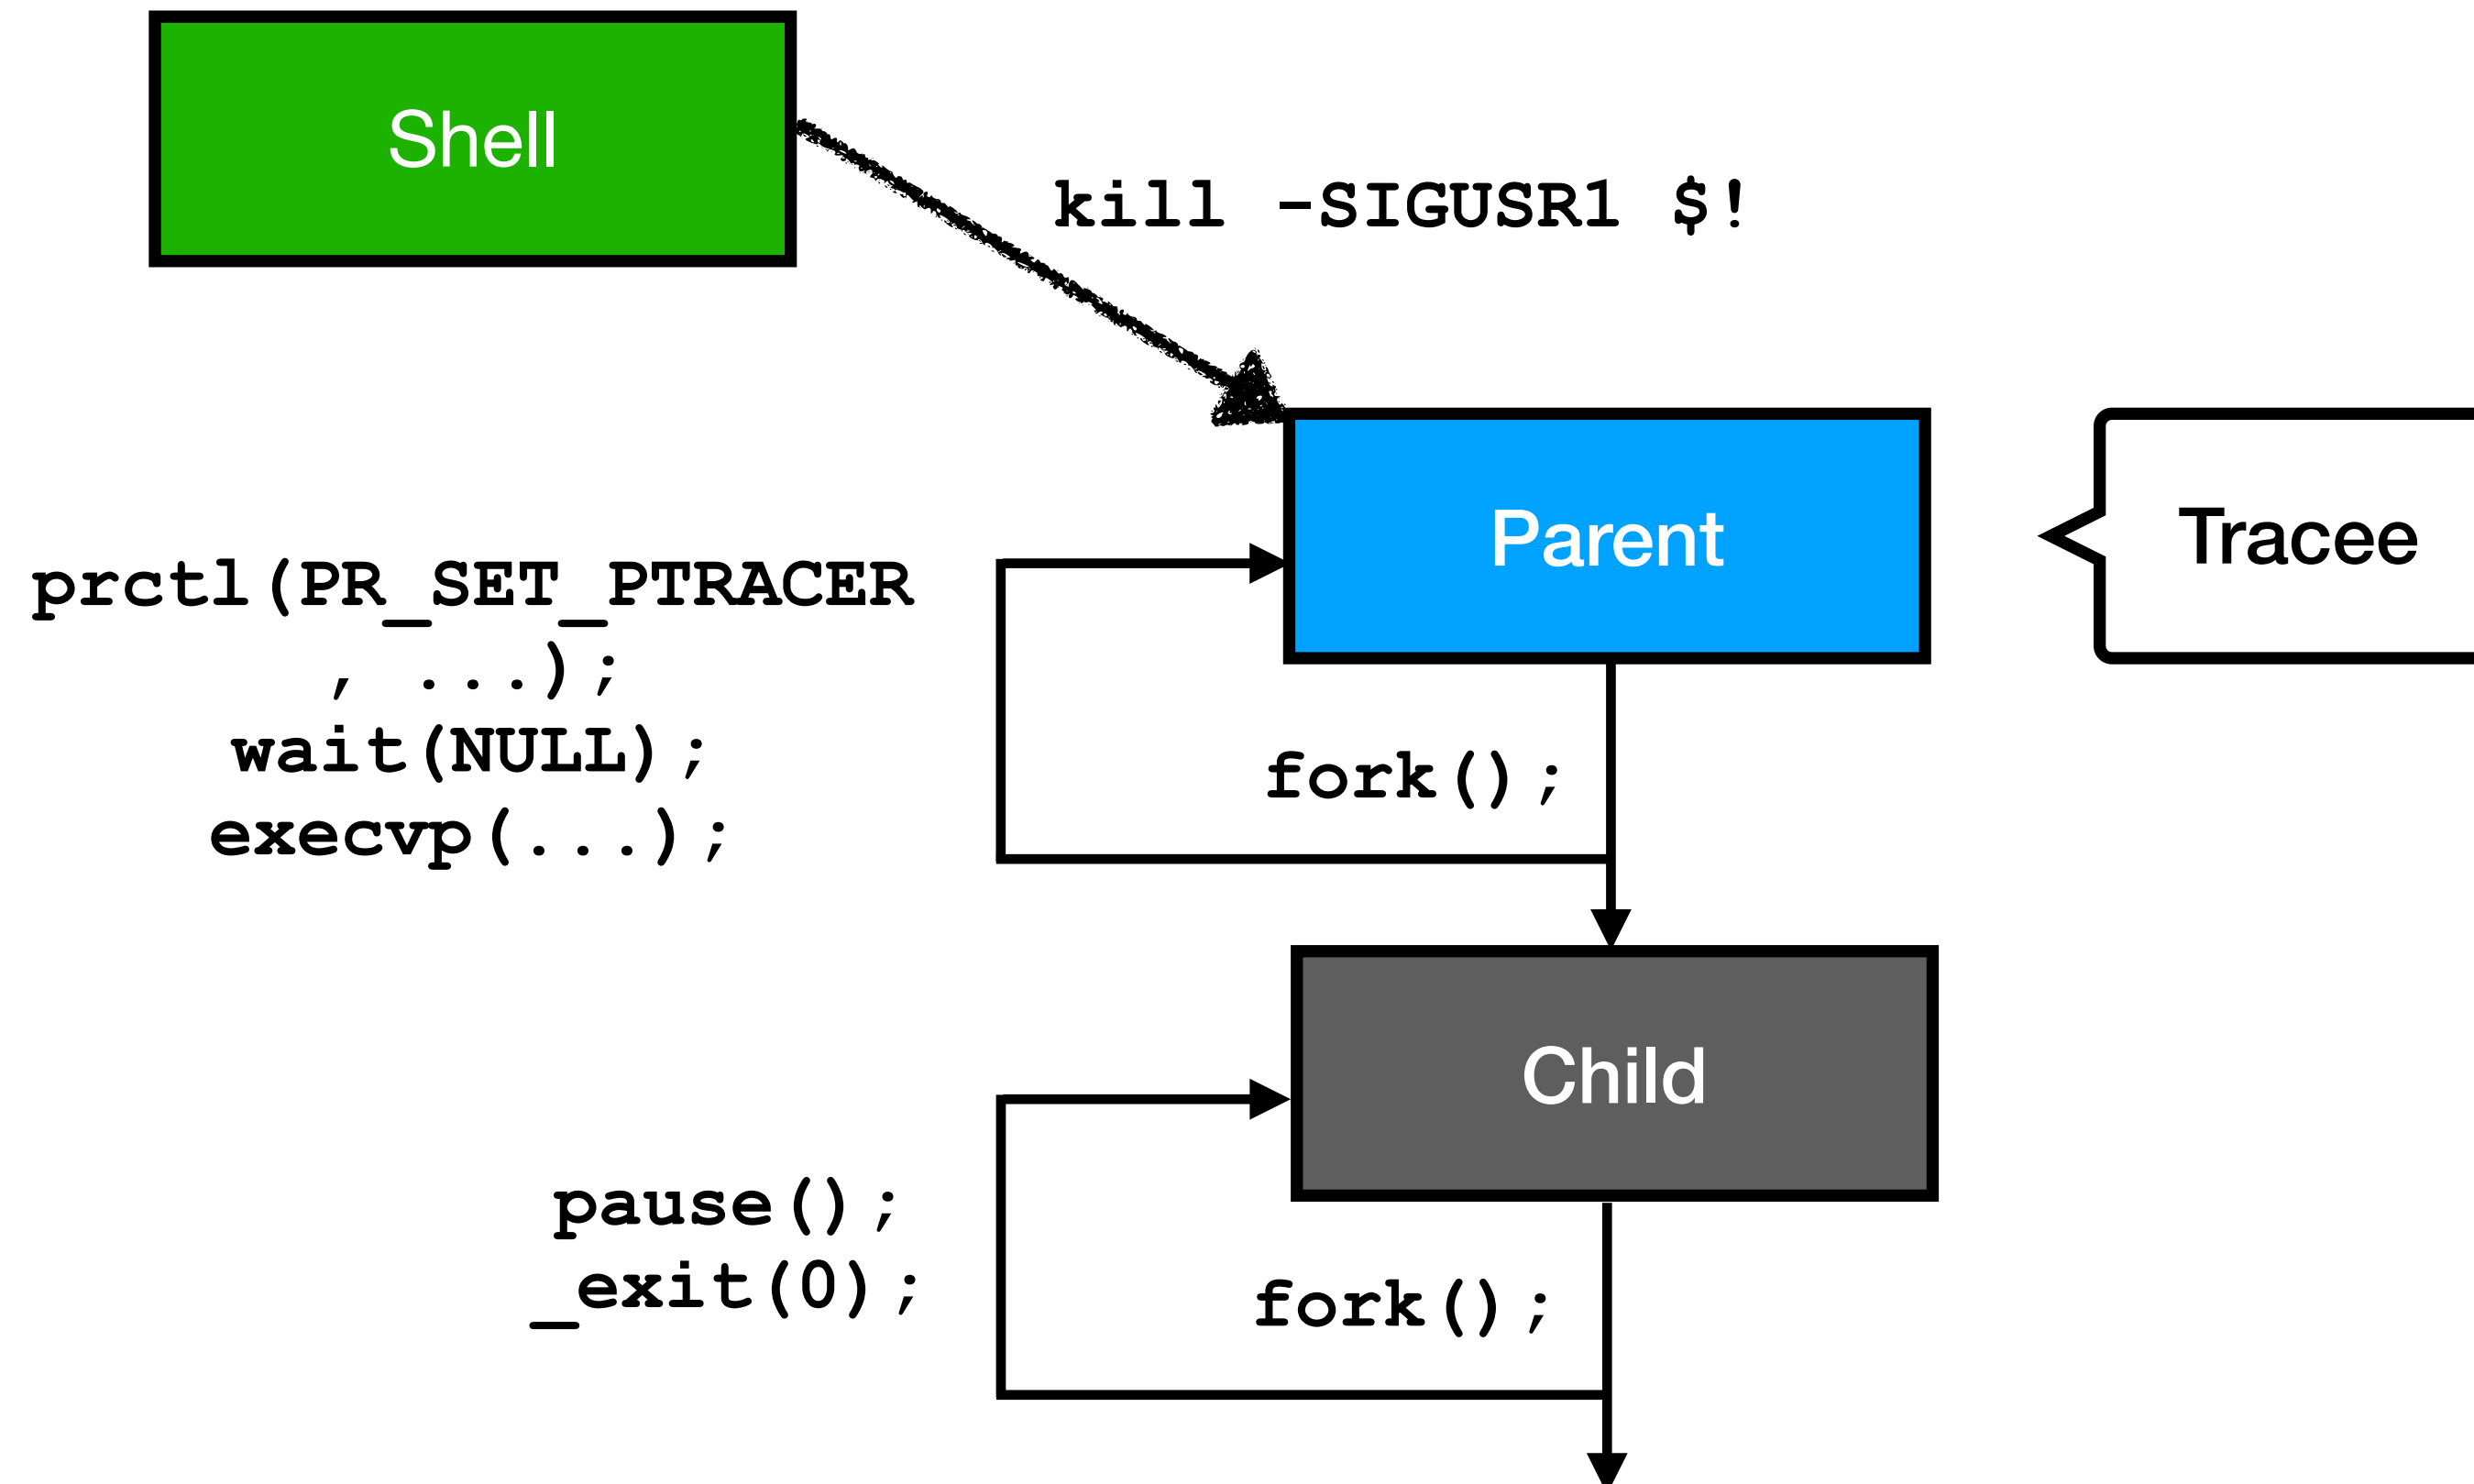
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



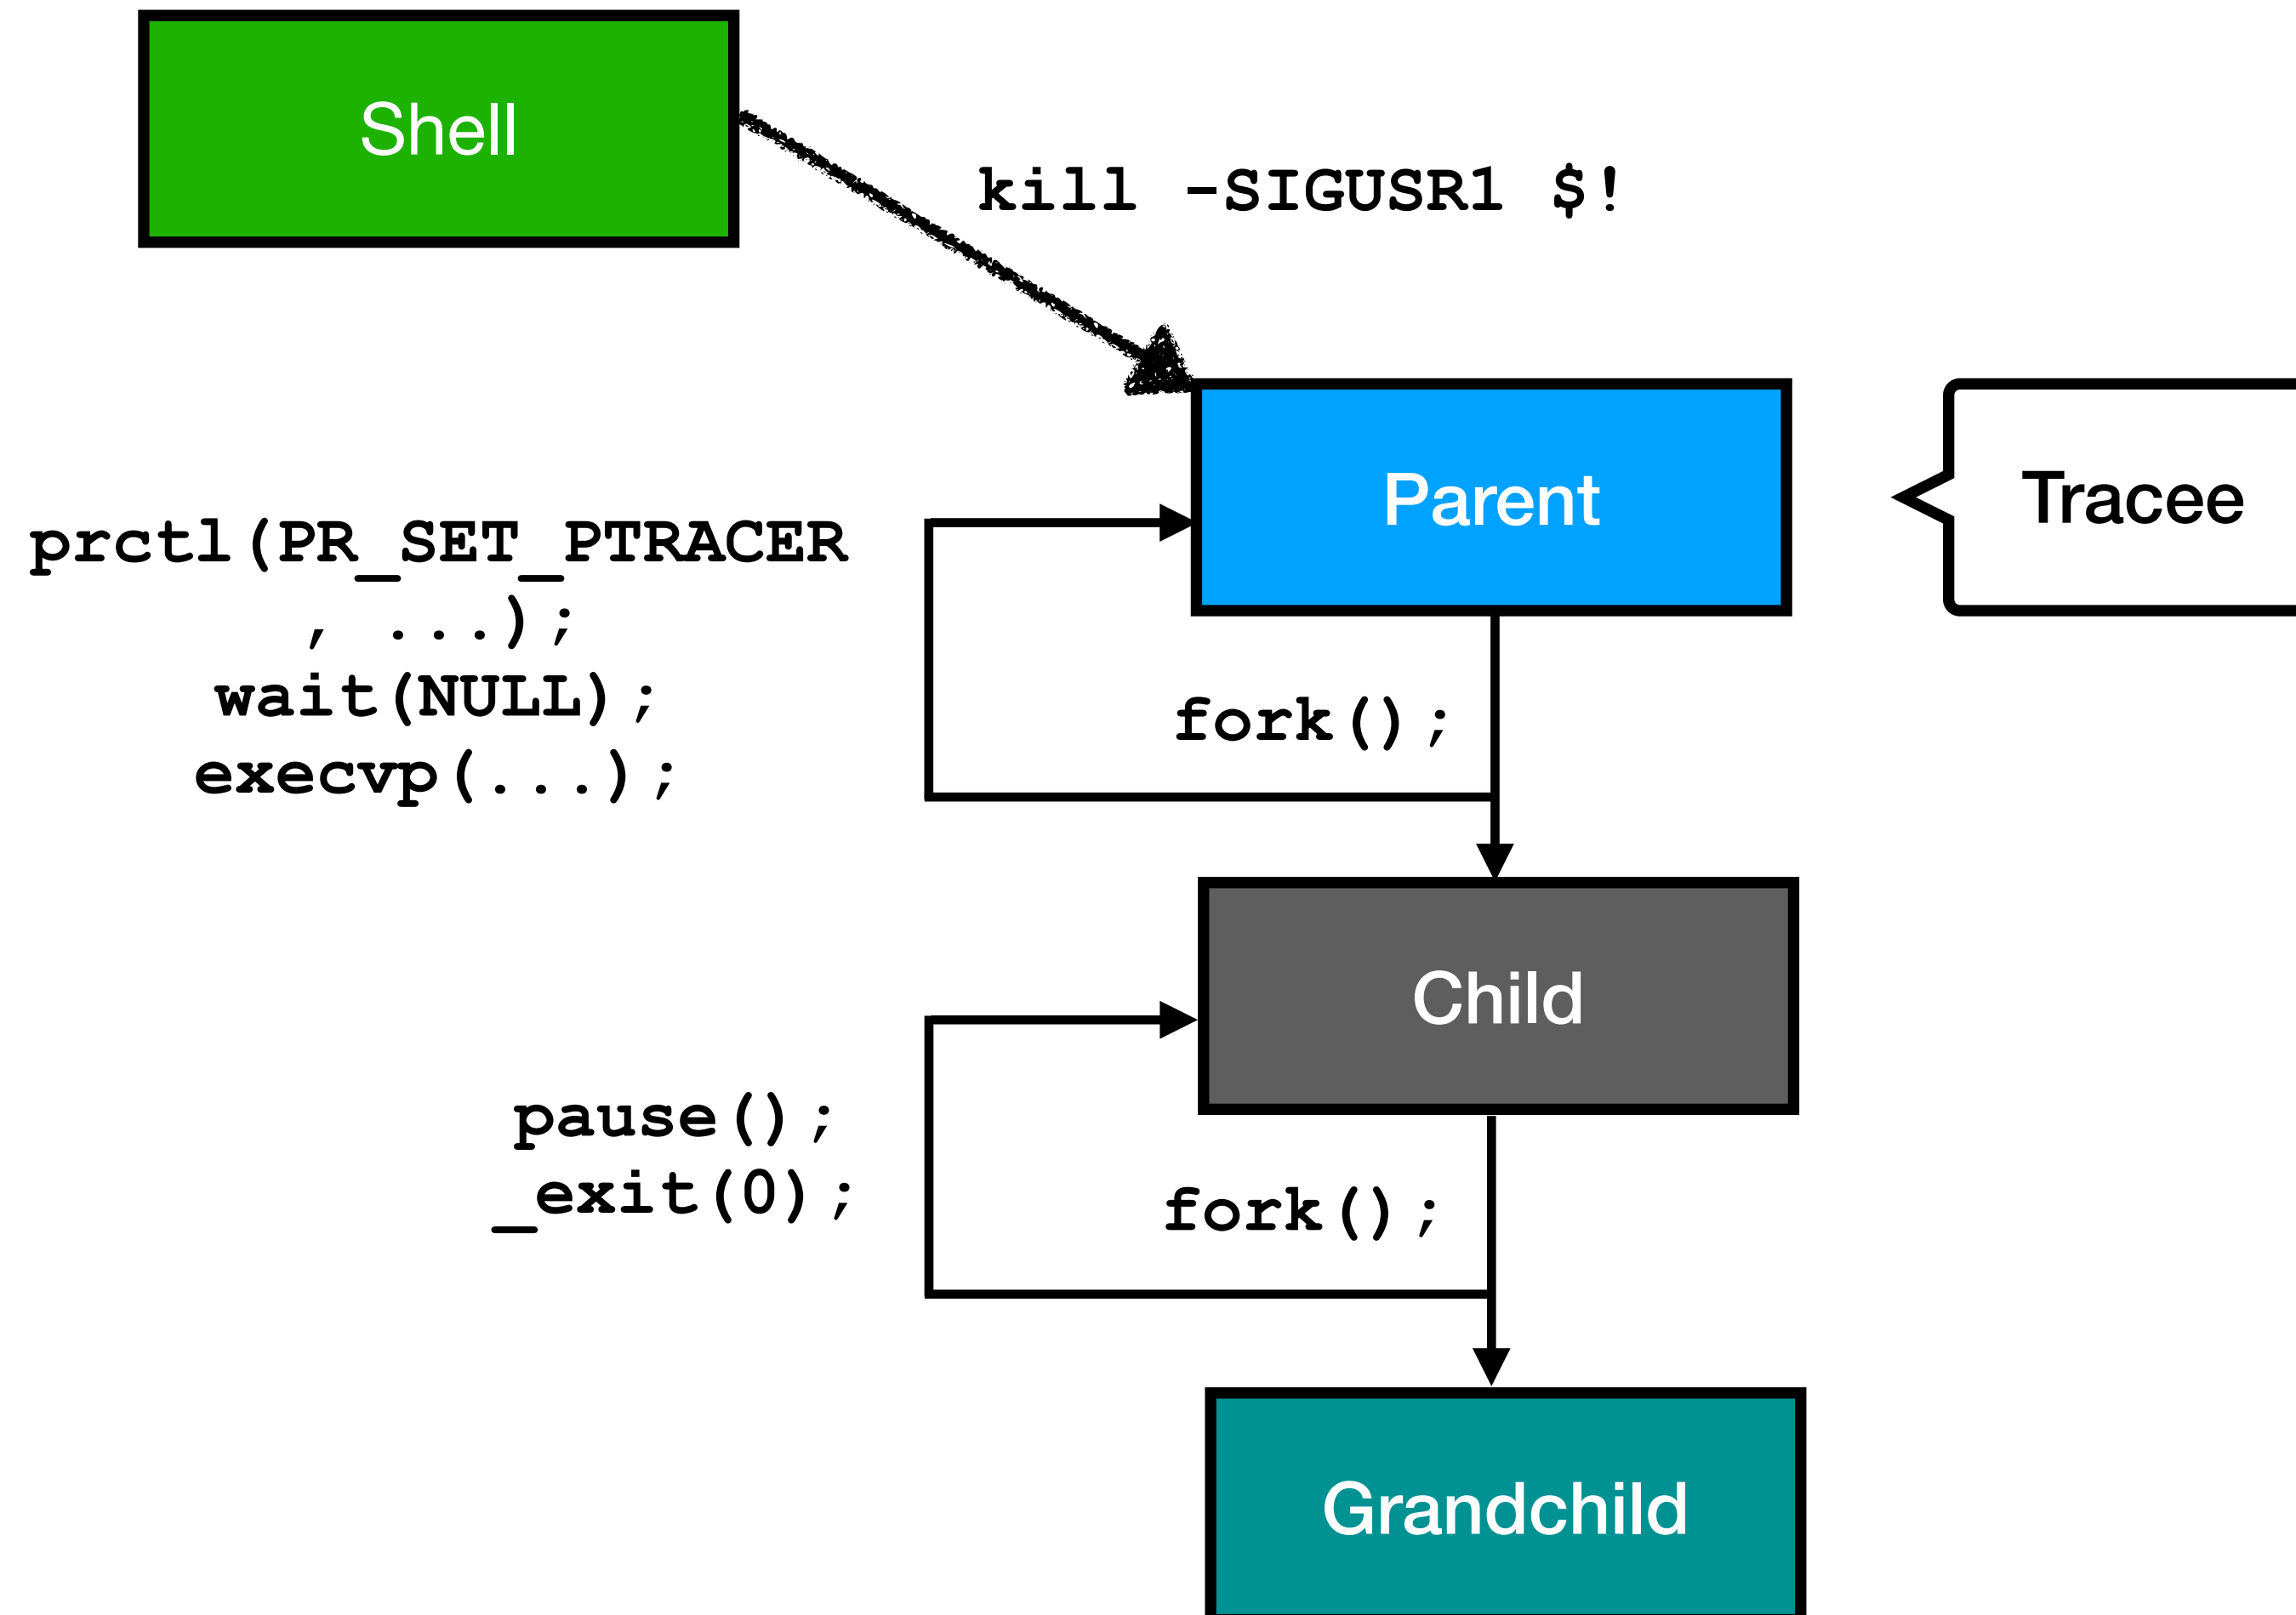
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



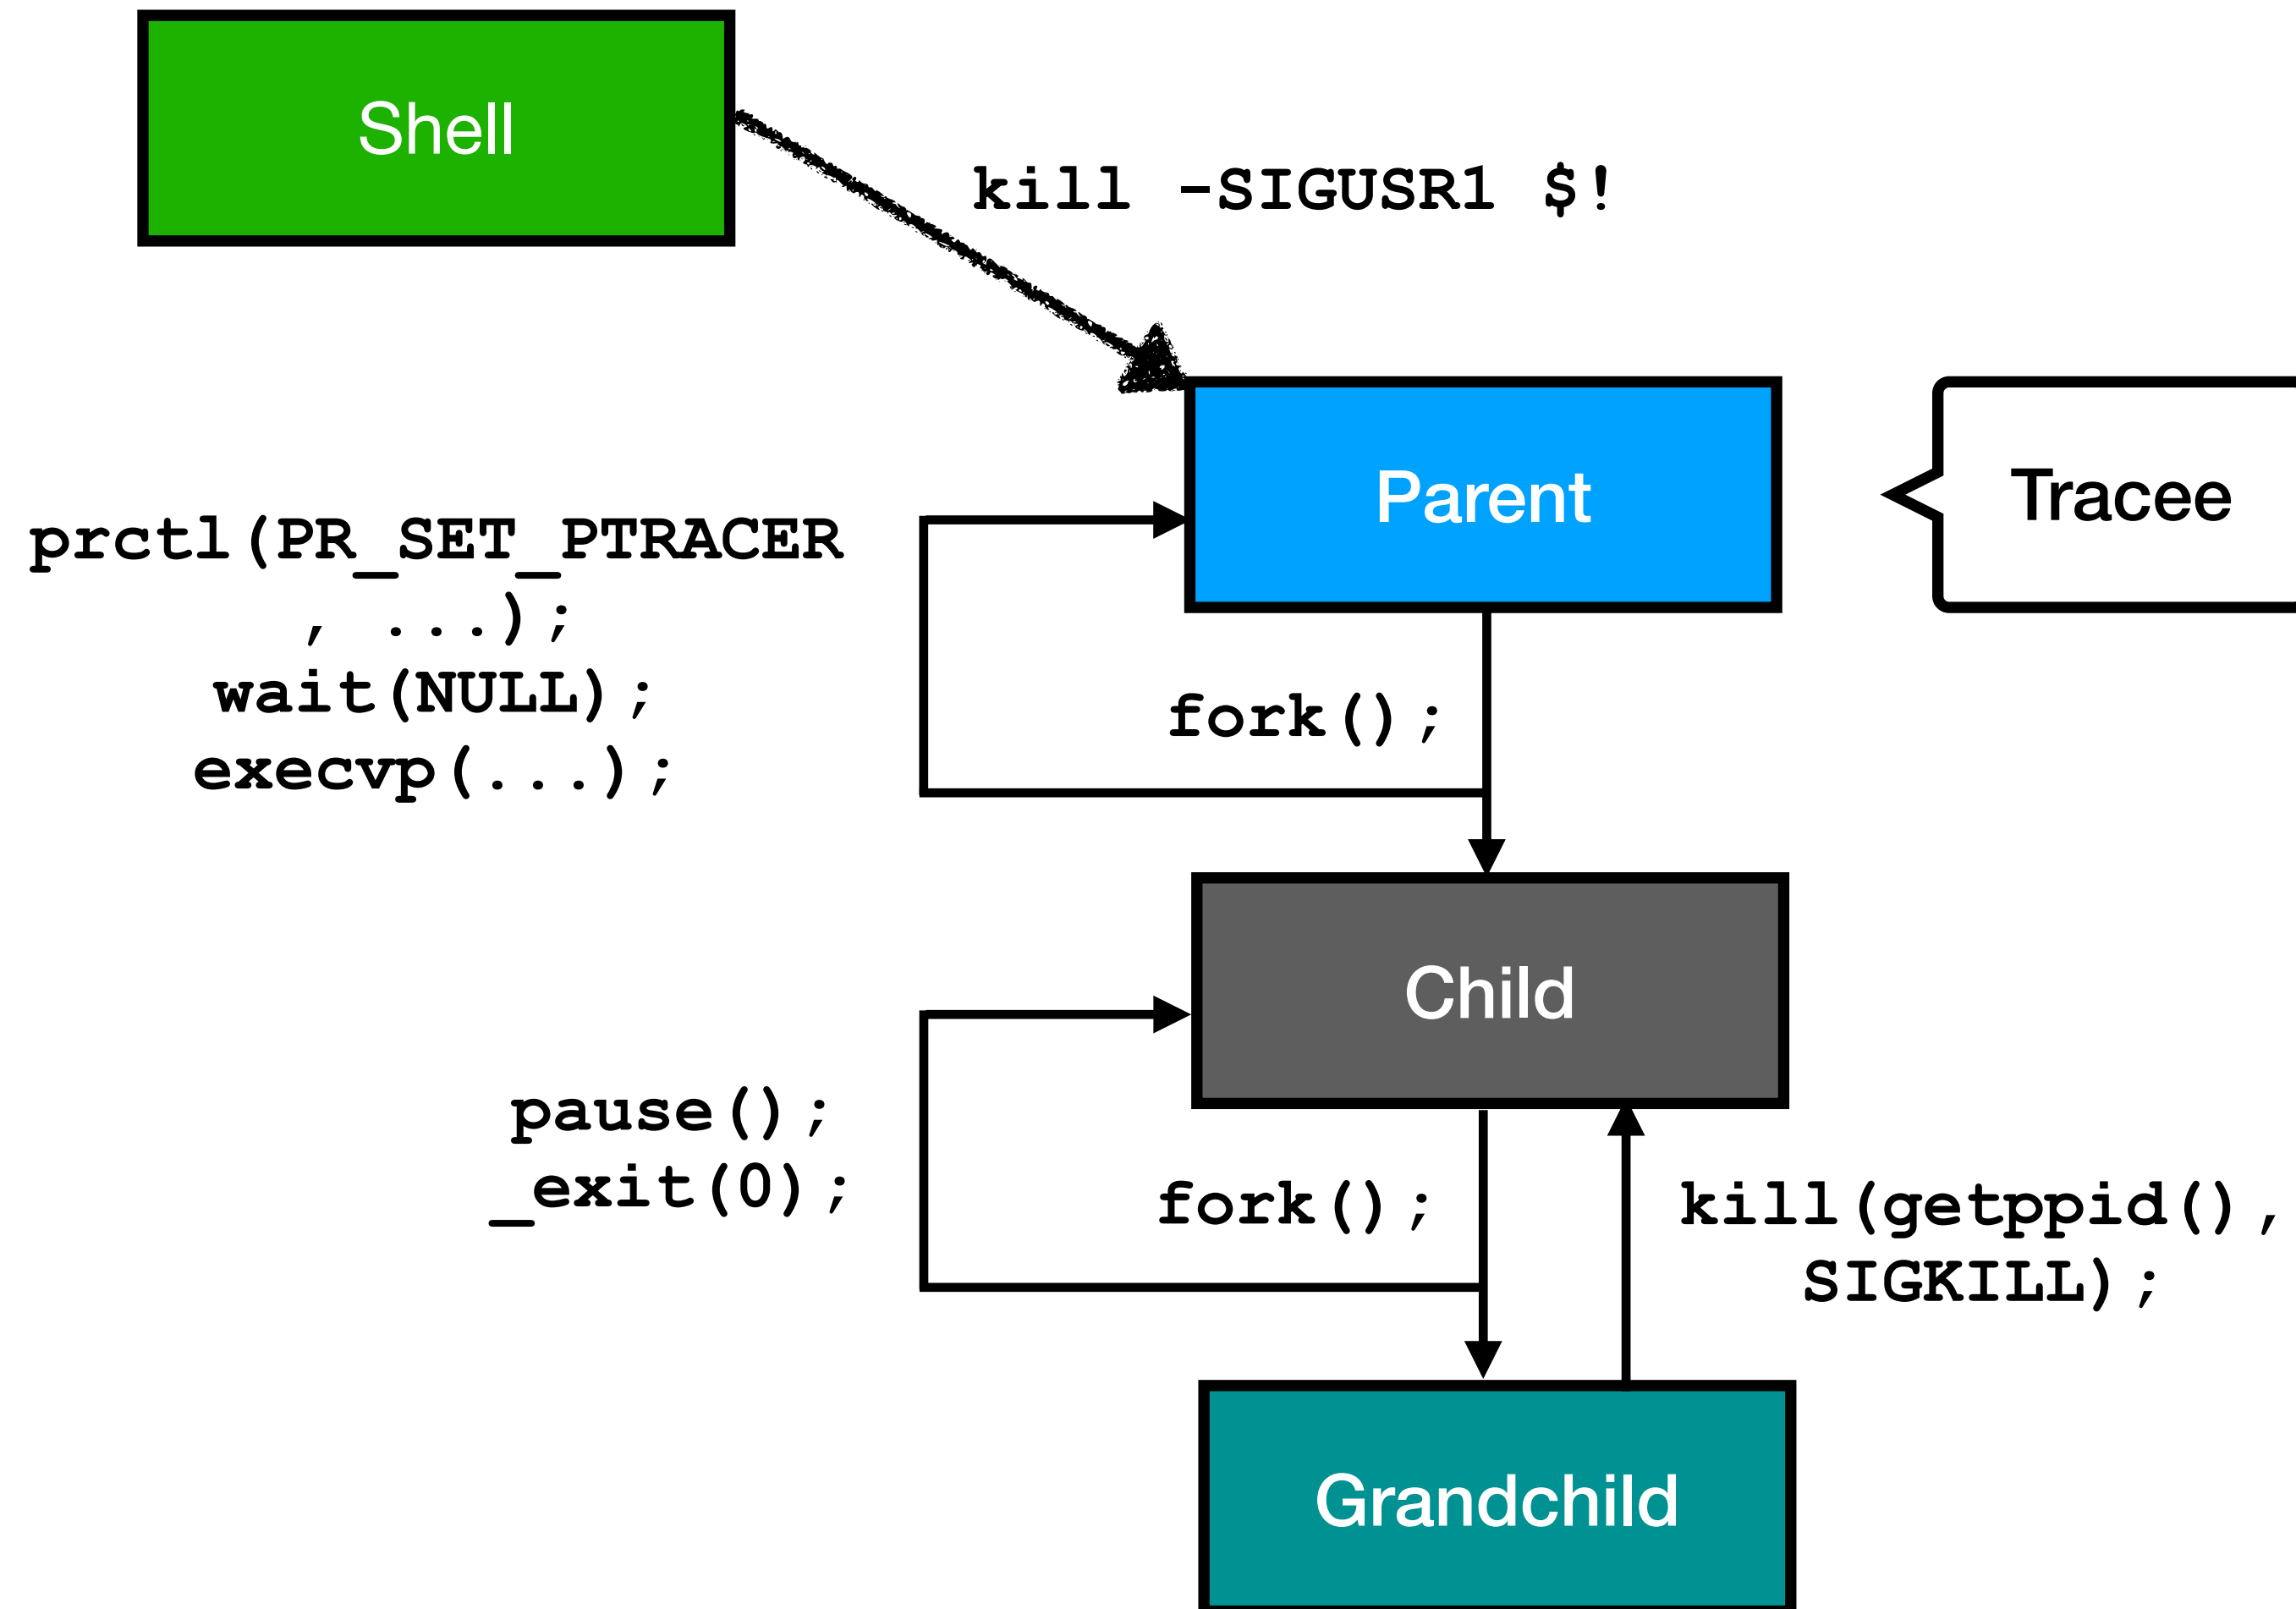
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



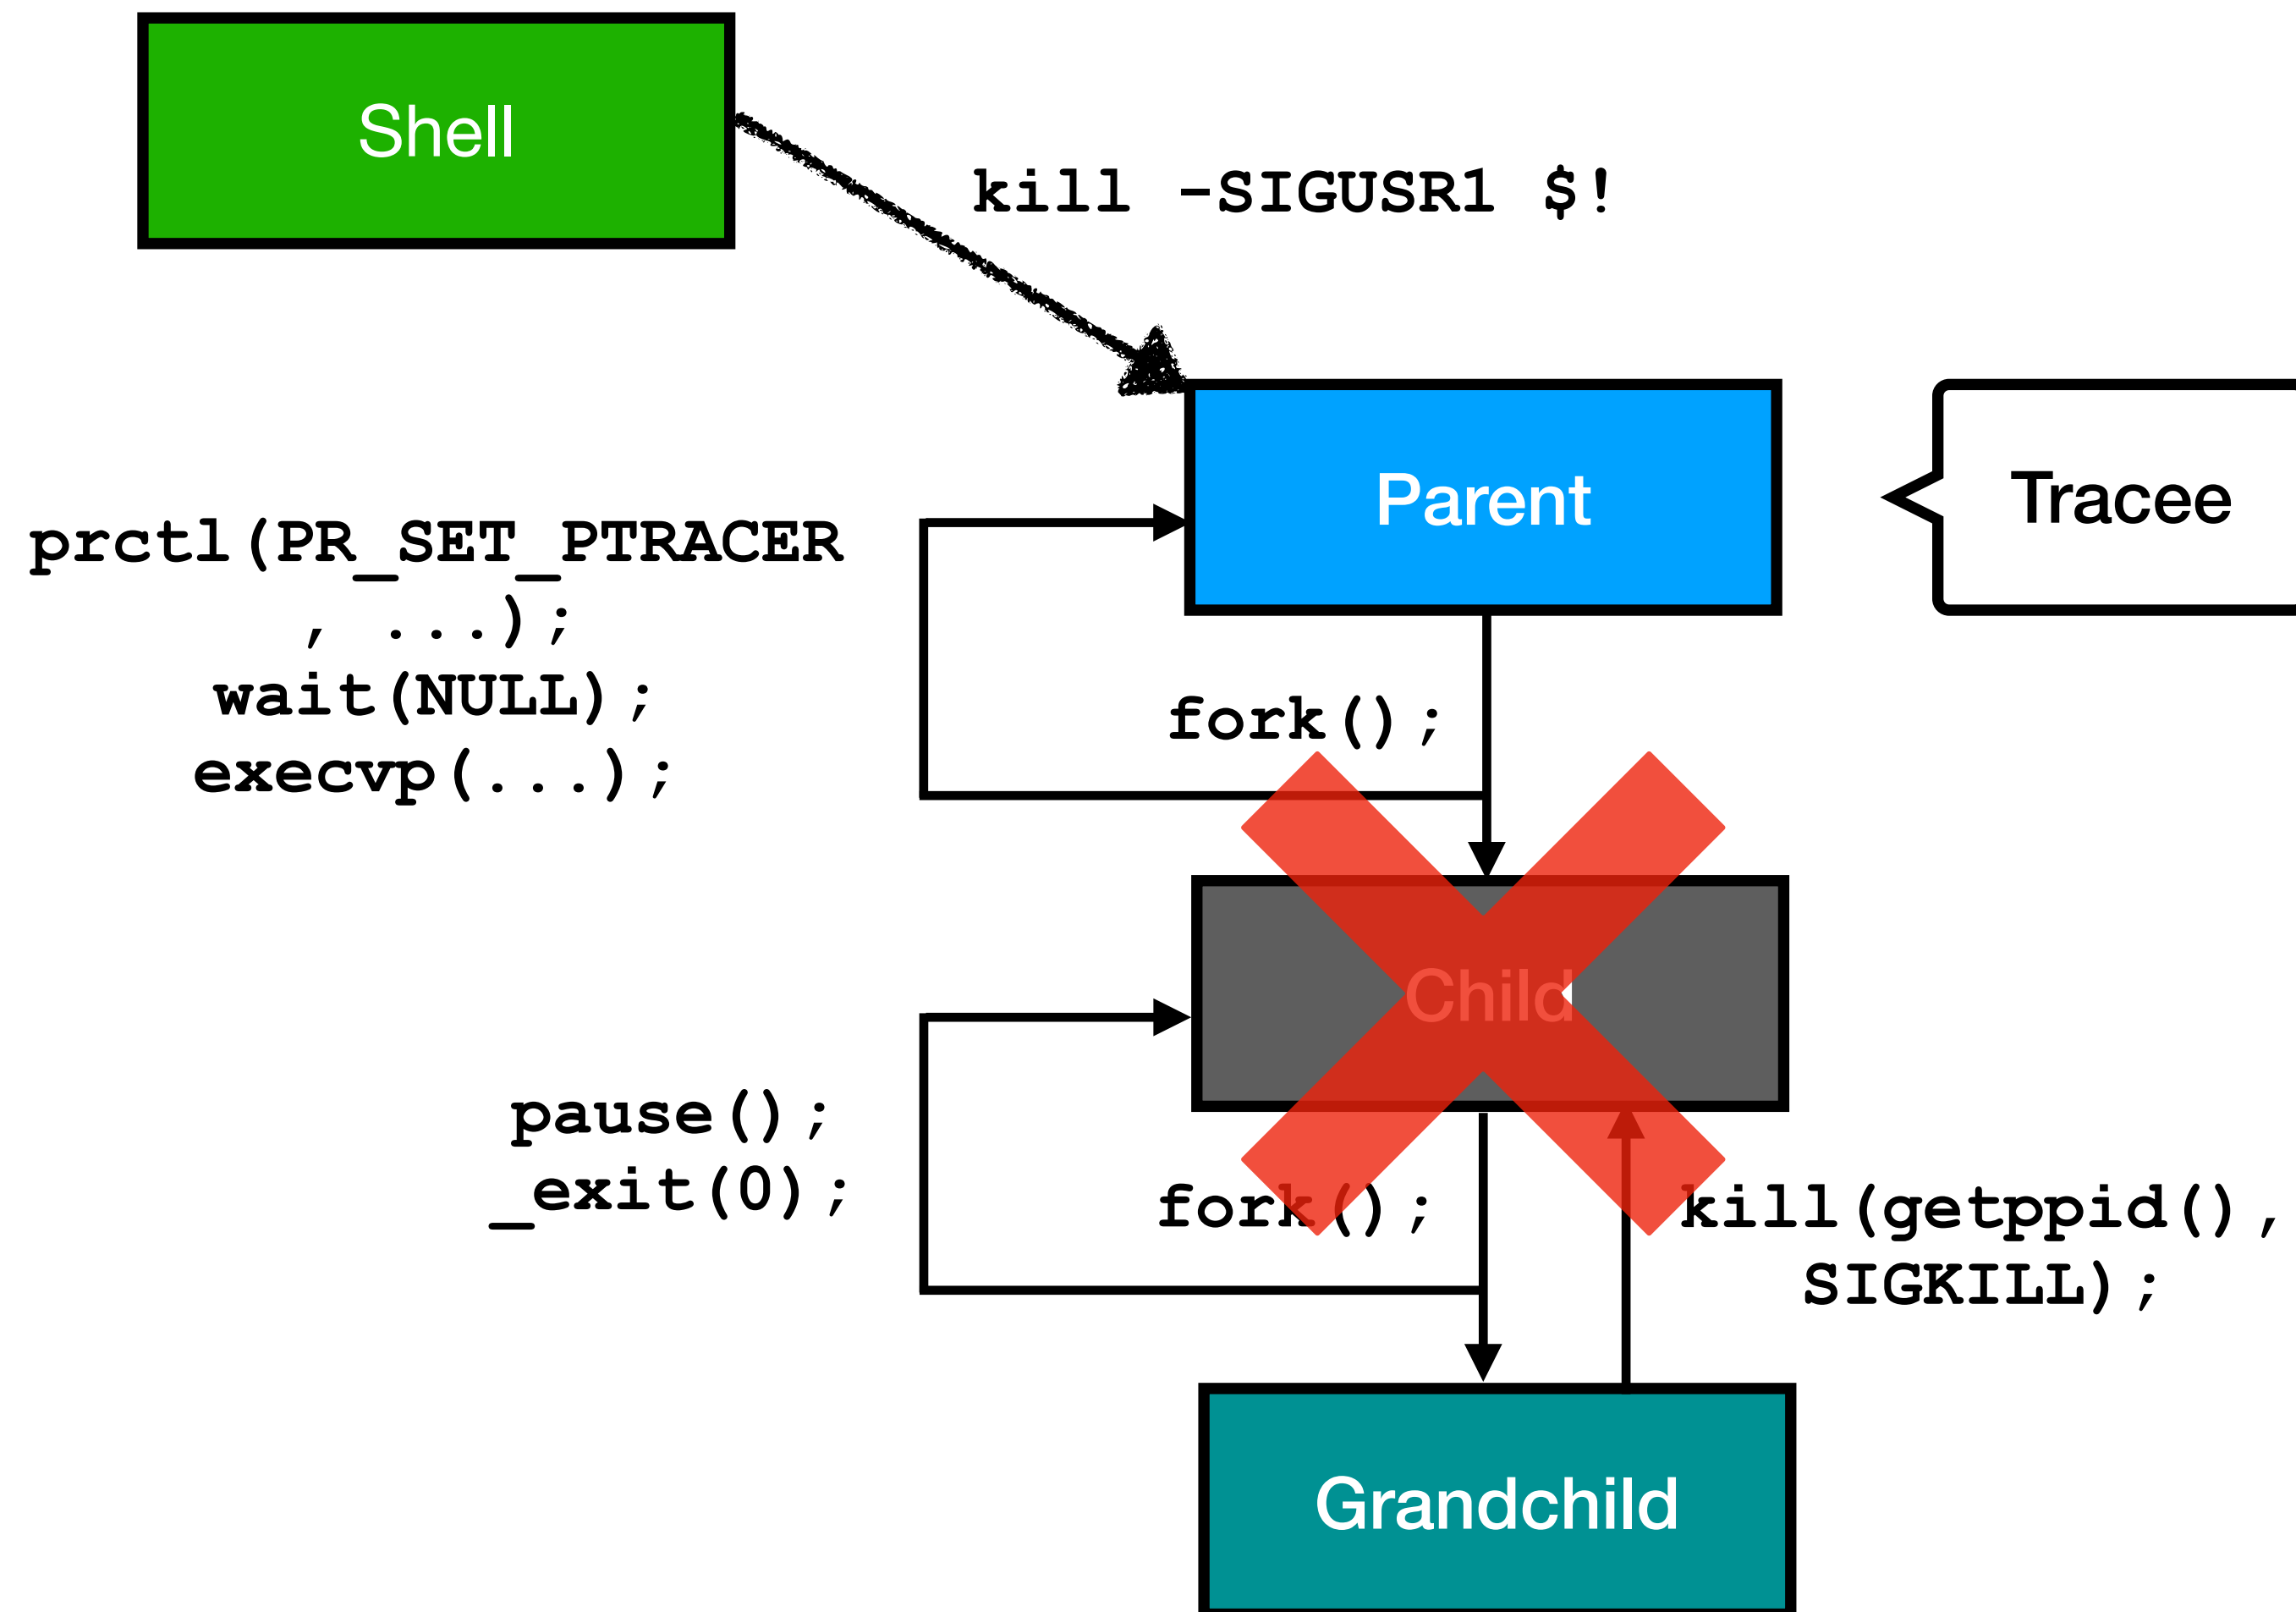
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



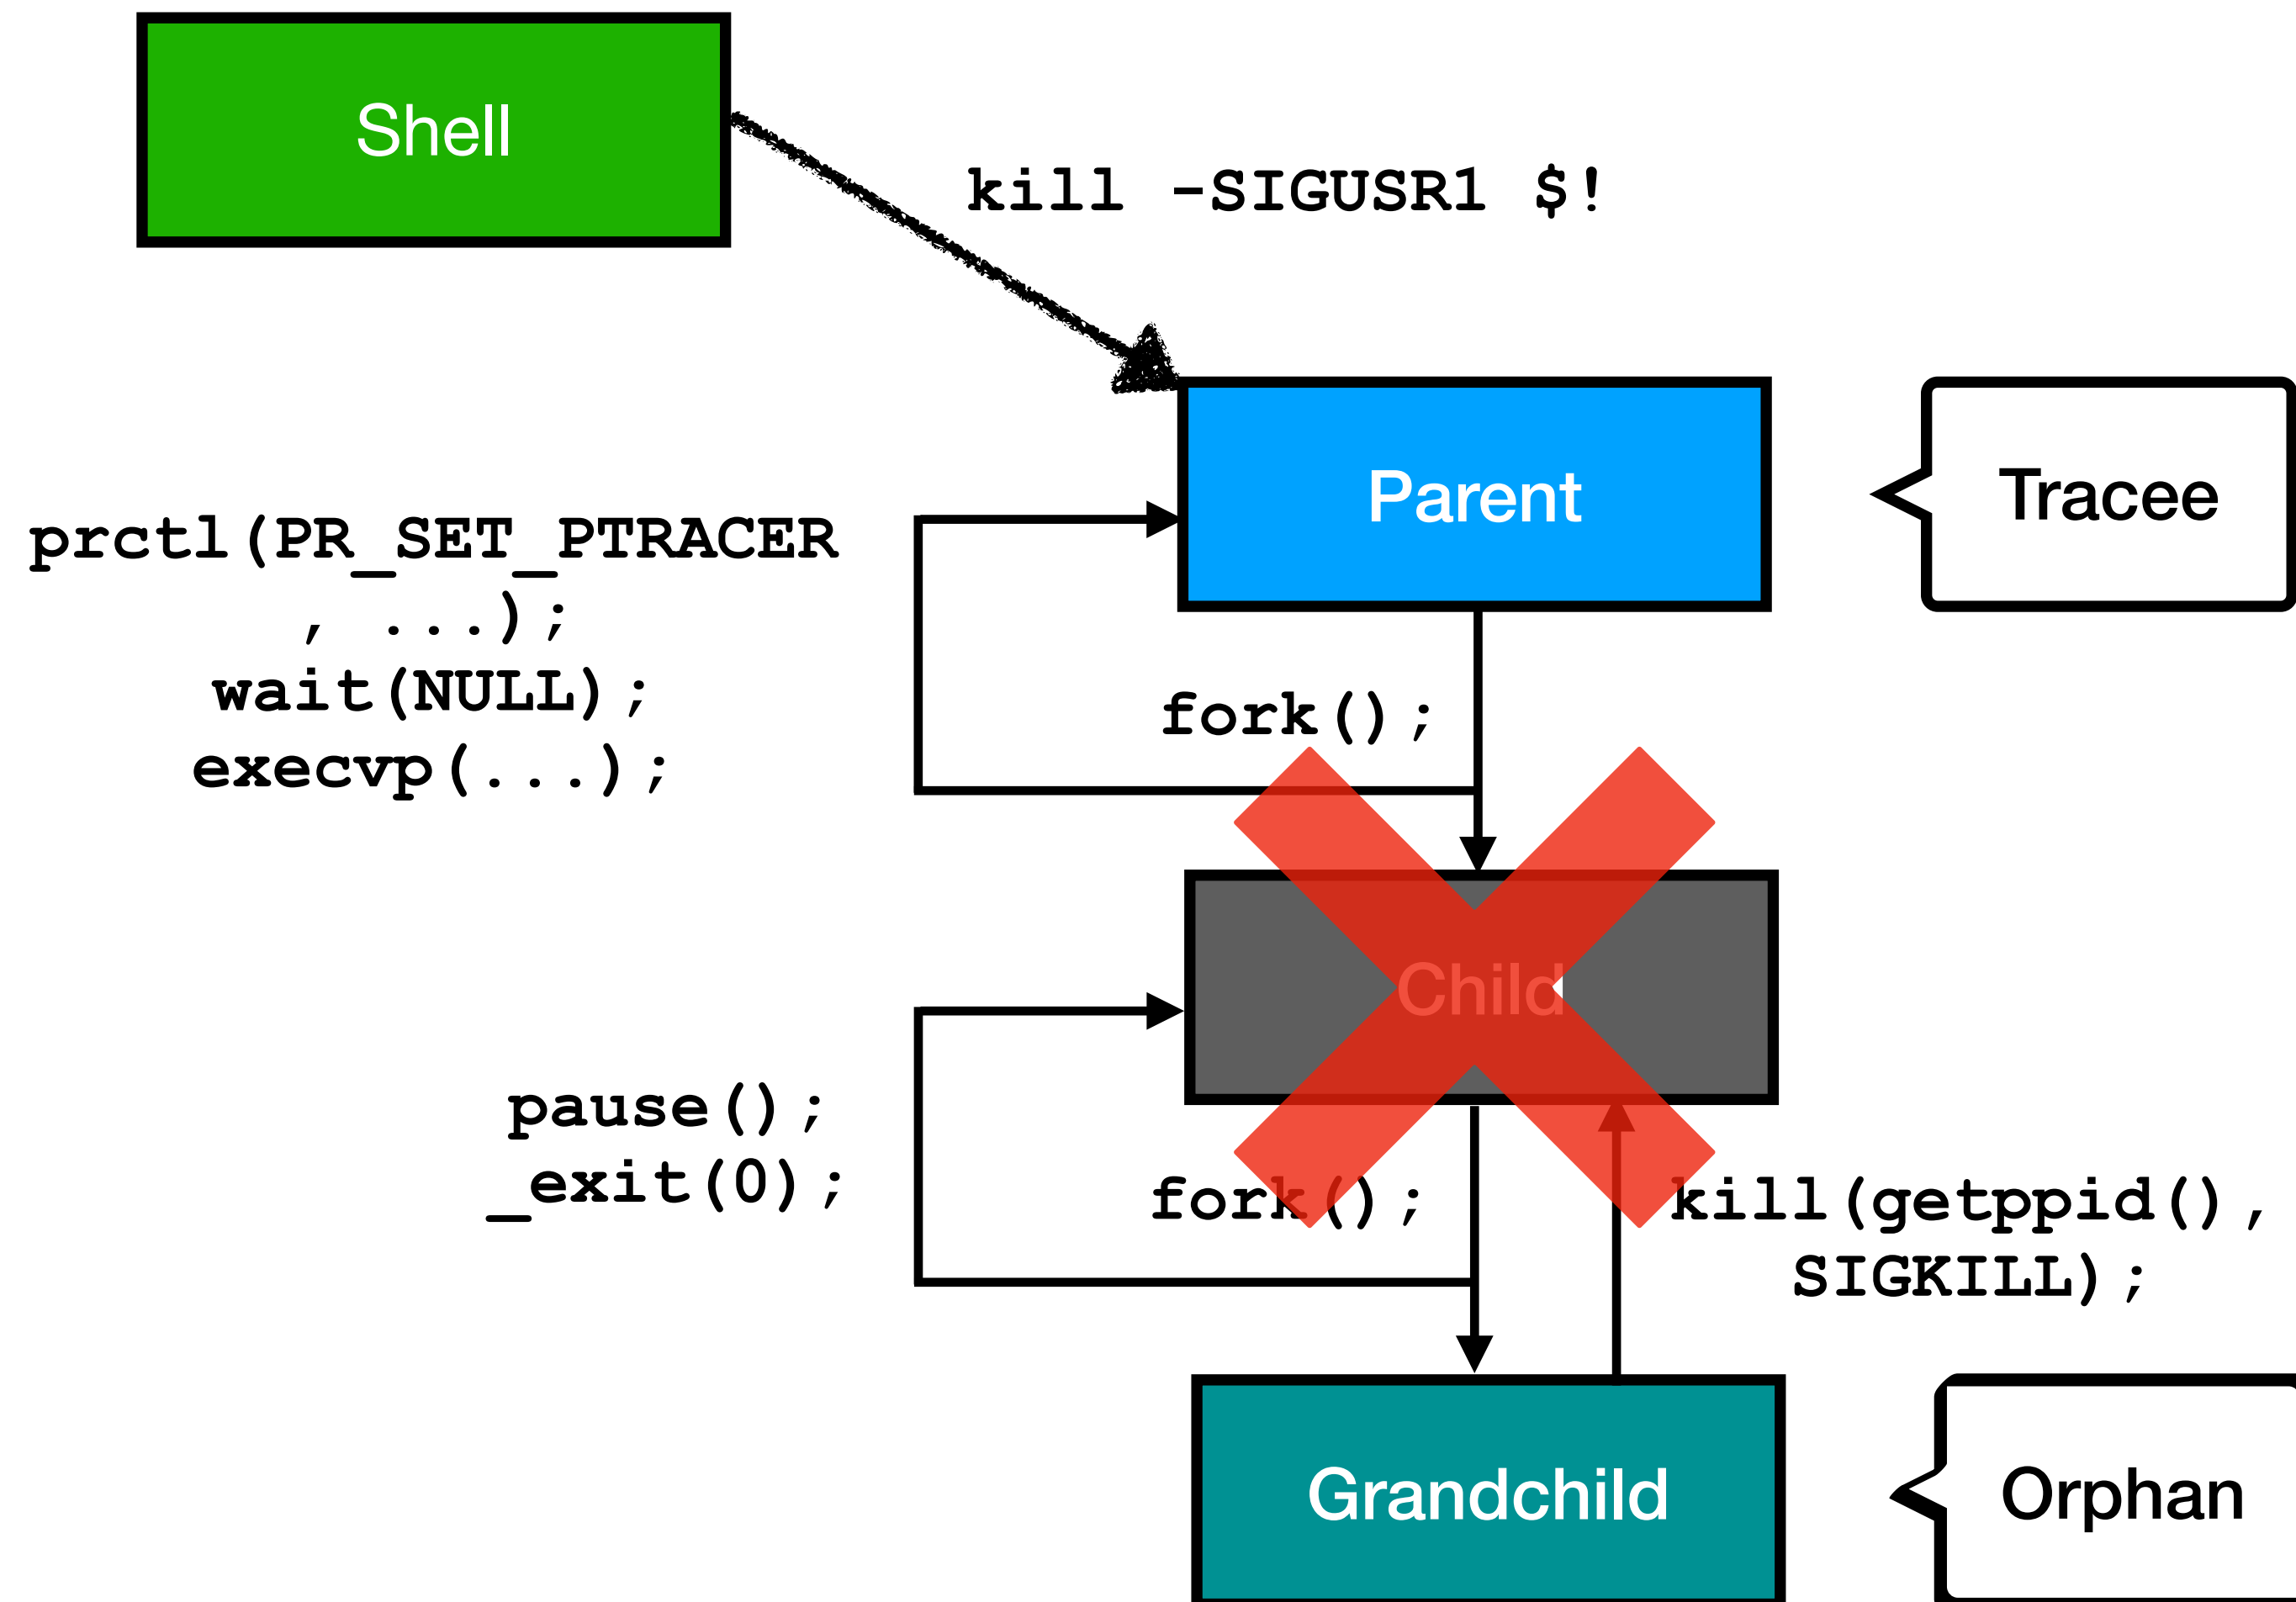
Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"




Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"




Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"



Tracing - Rollen

- "-D Run tracer [...] as a **grandchild**, not as parent"




Tracing - Flow


Tracing - Flow

Start

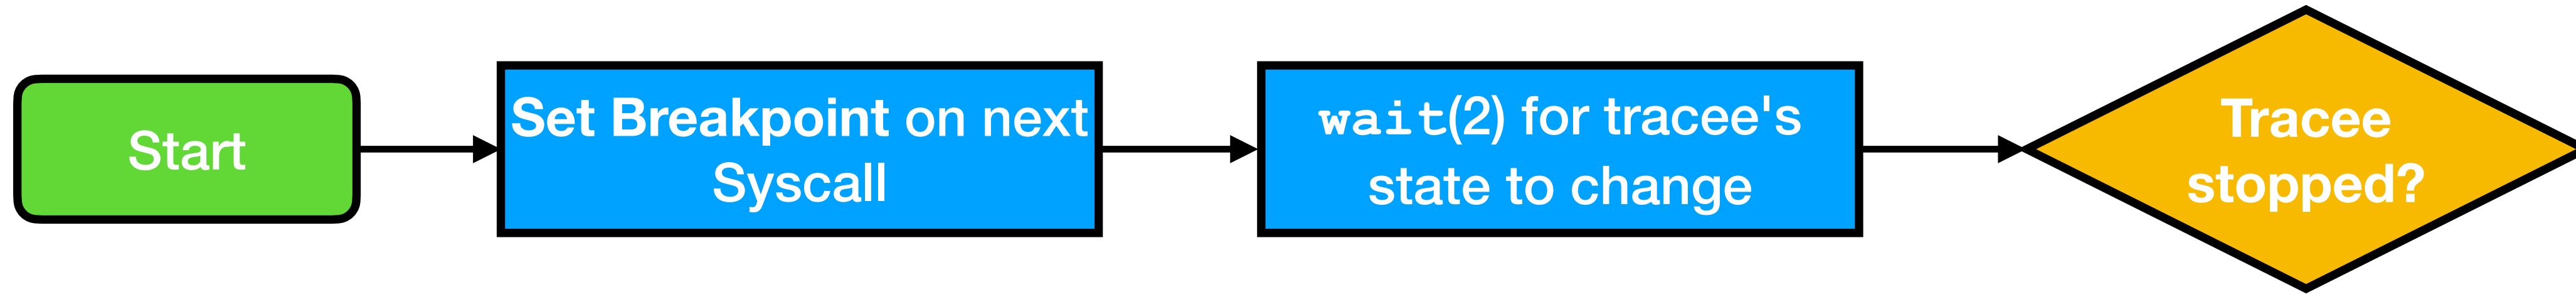
Tracing - Flow



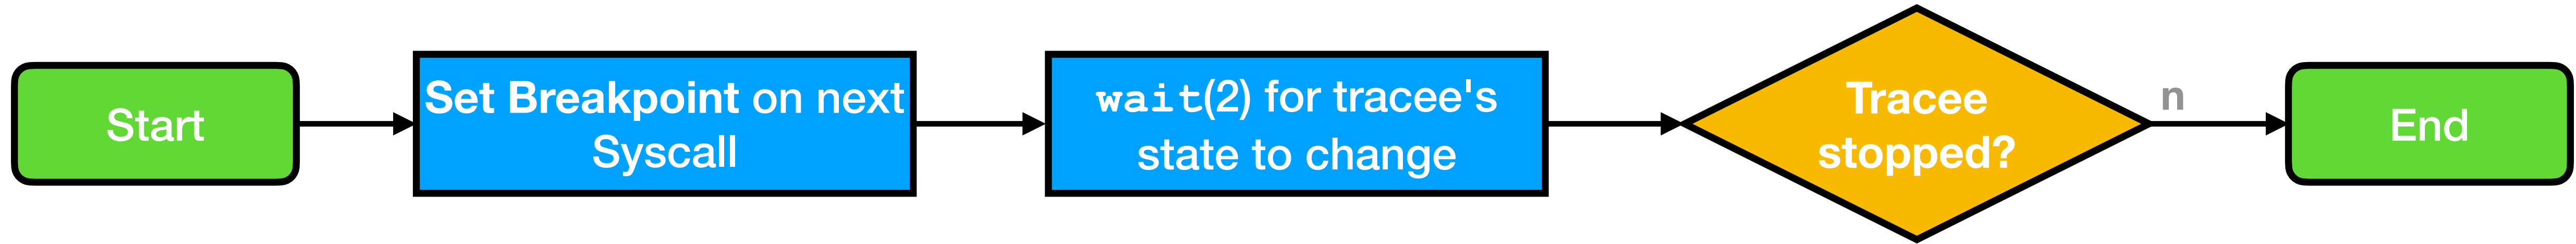
Tracing - Flow



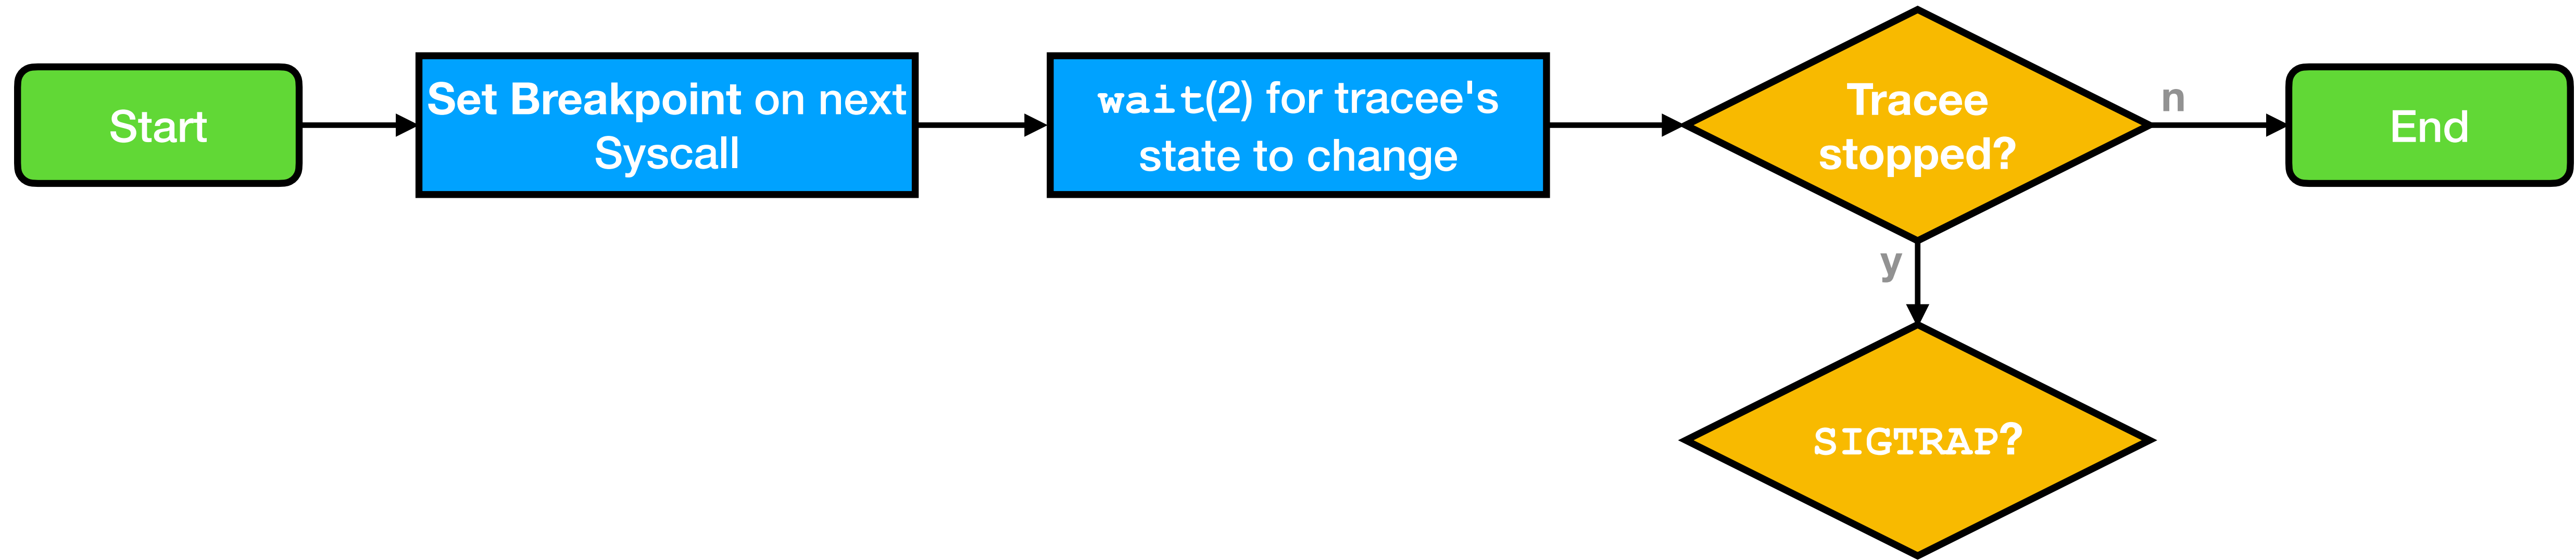
Tracing - Flow



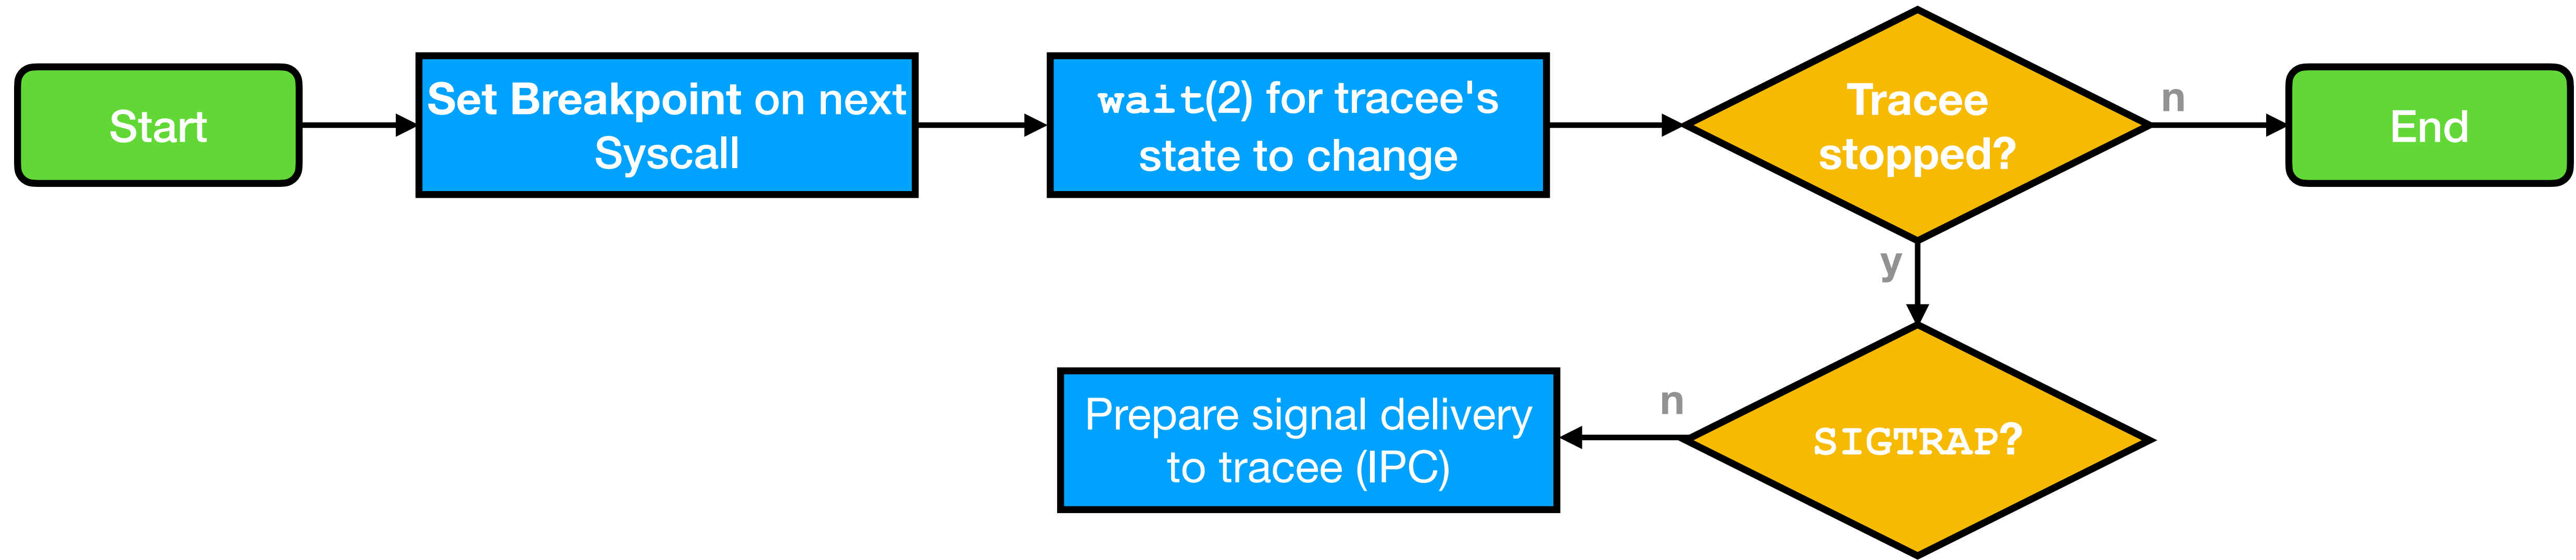
Tracing - Flow



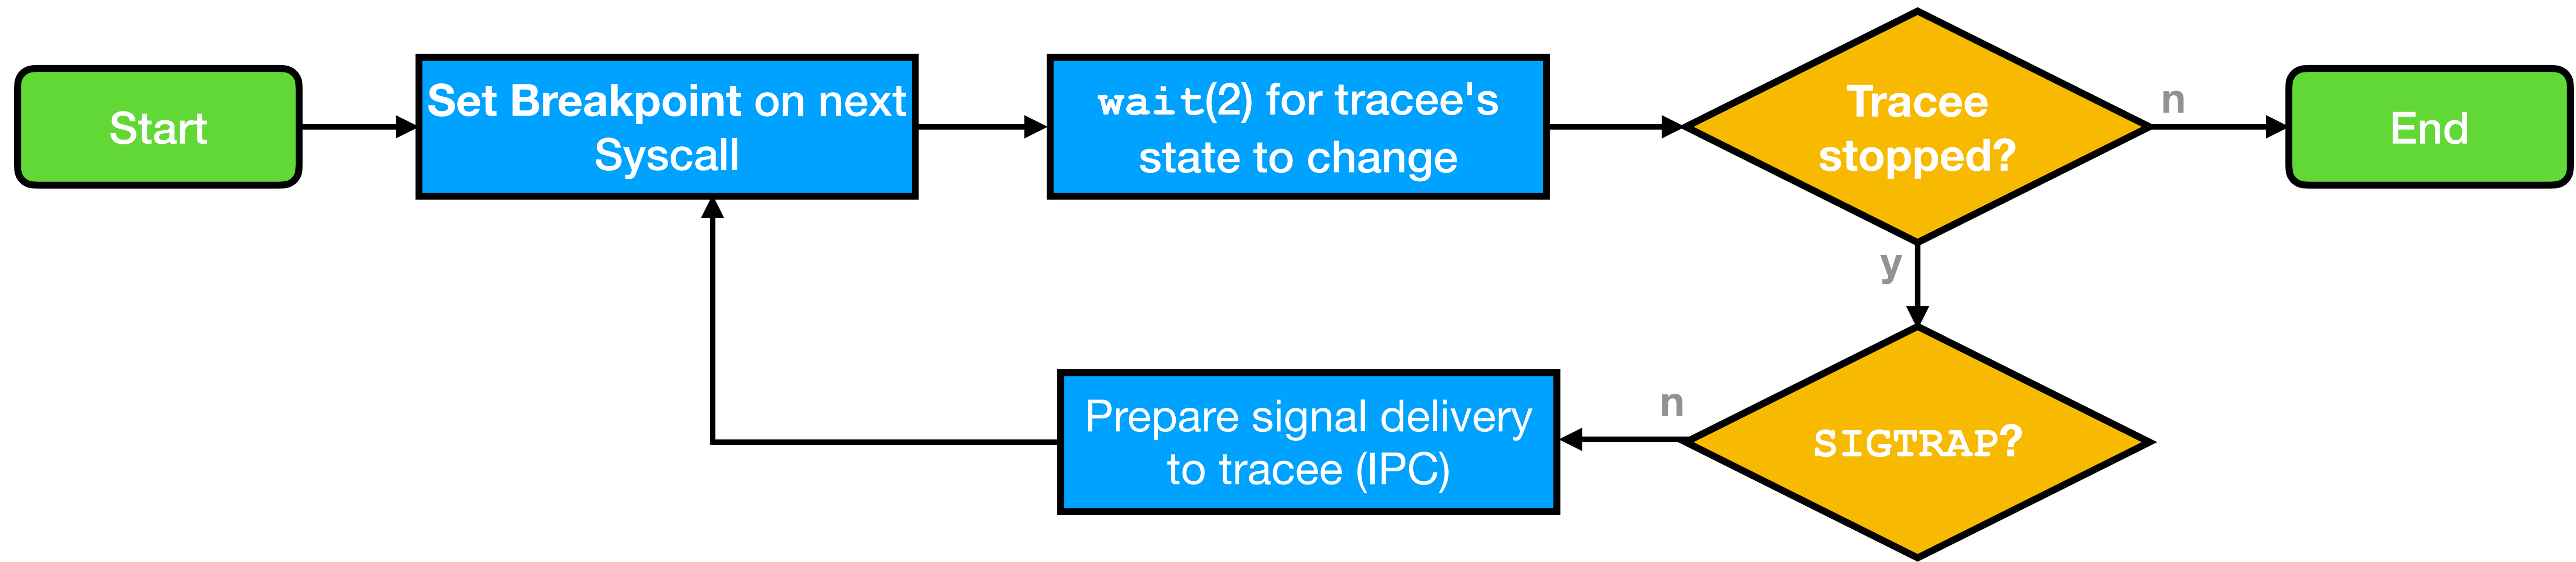
Tracing - Flow



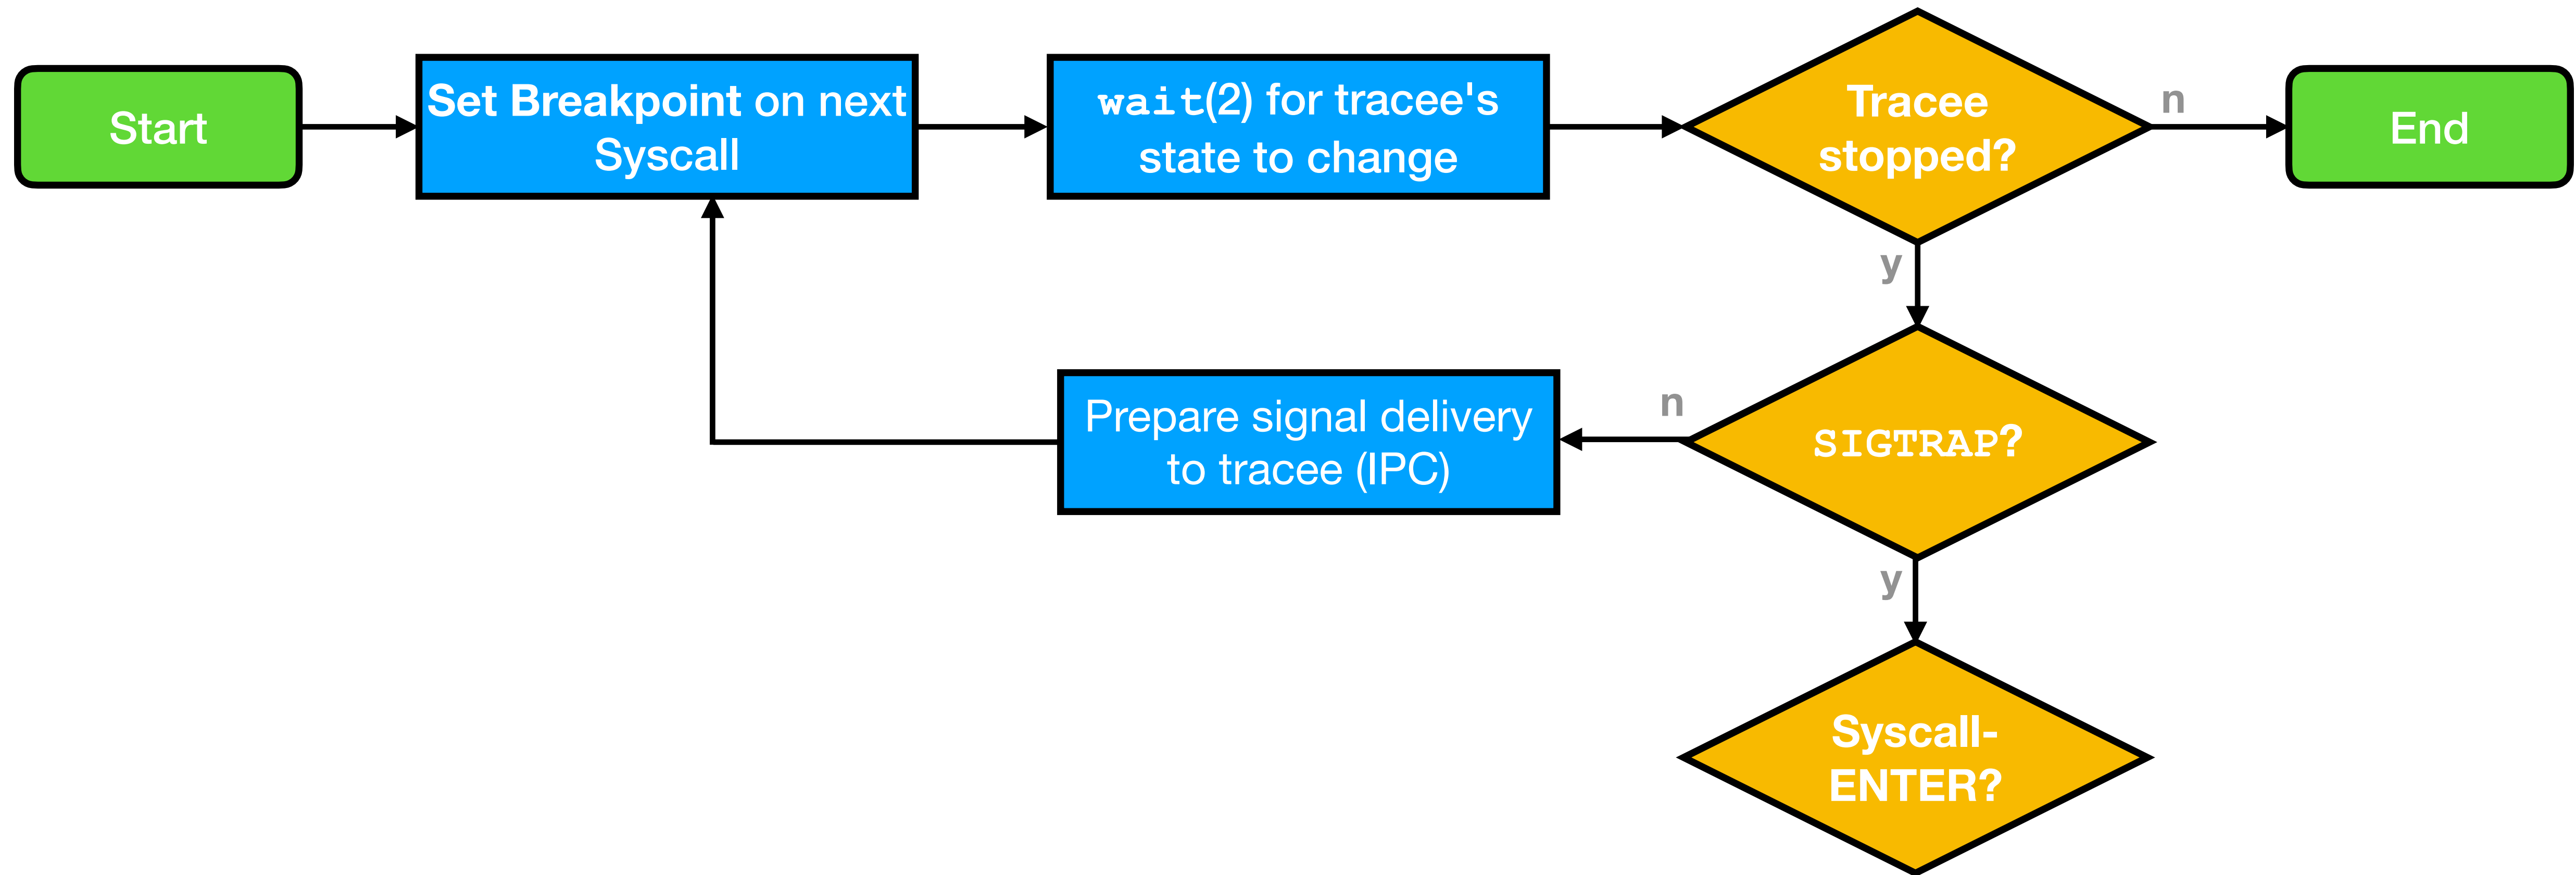
Tracing - Flow



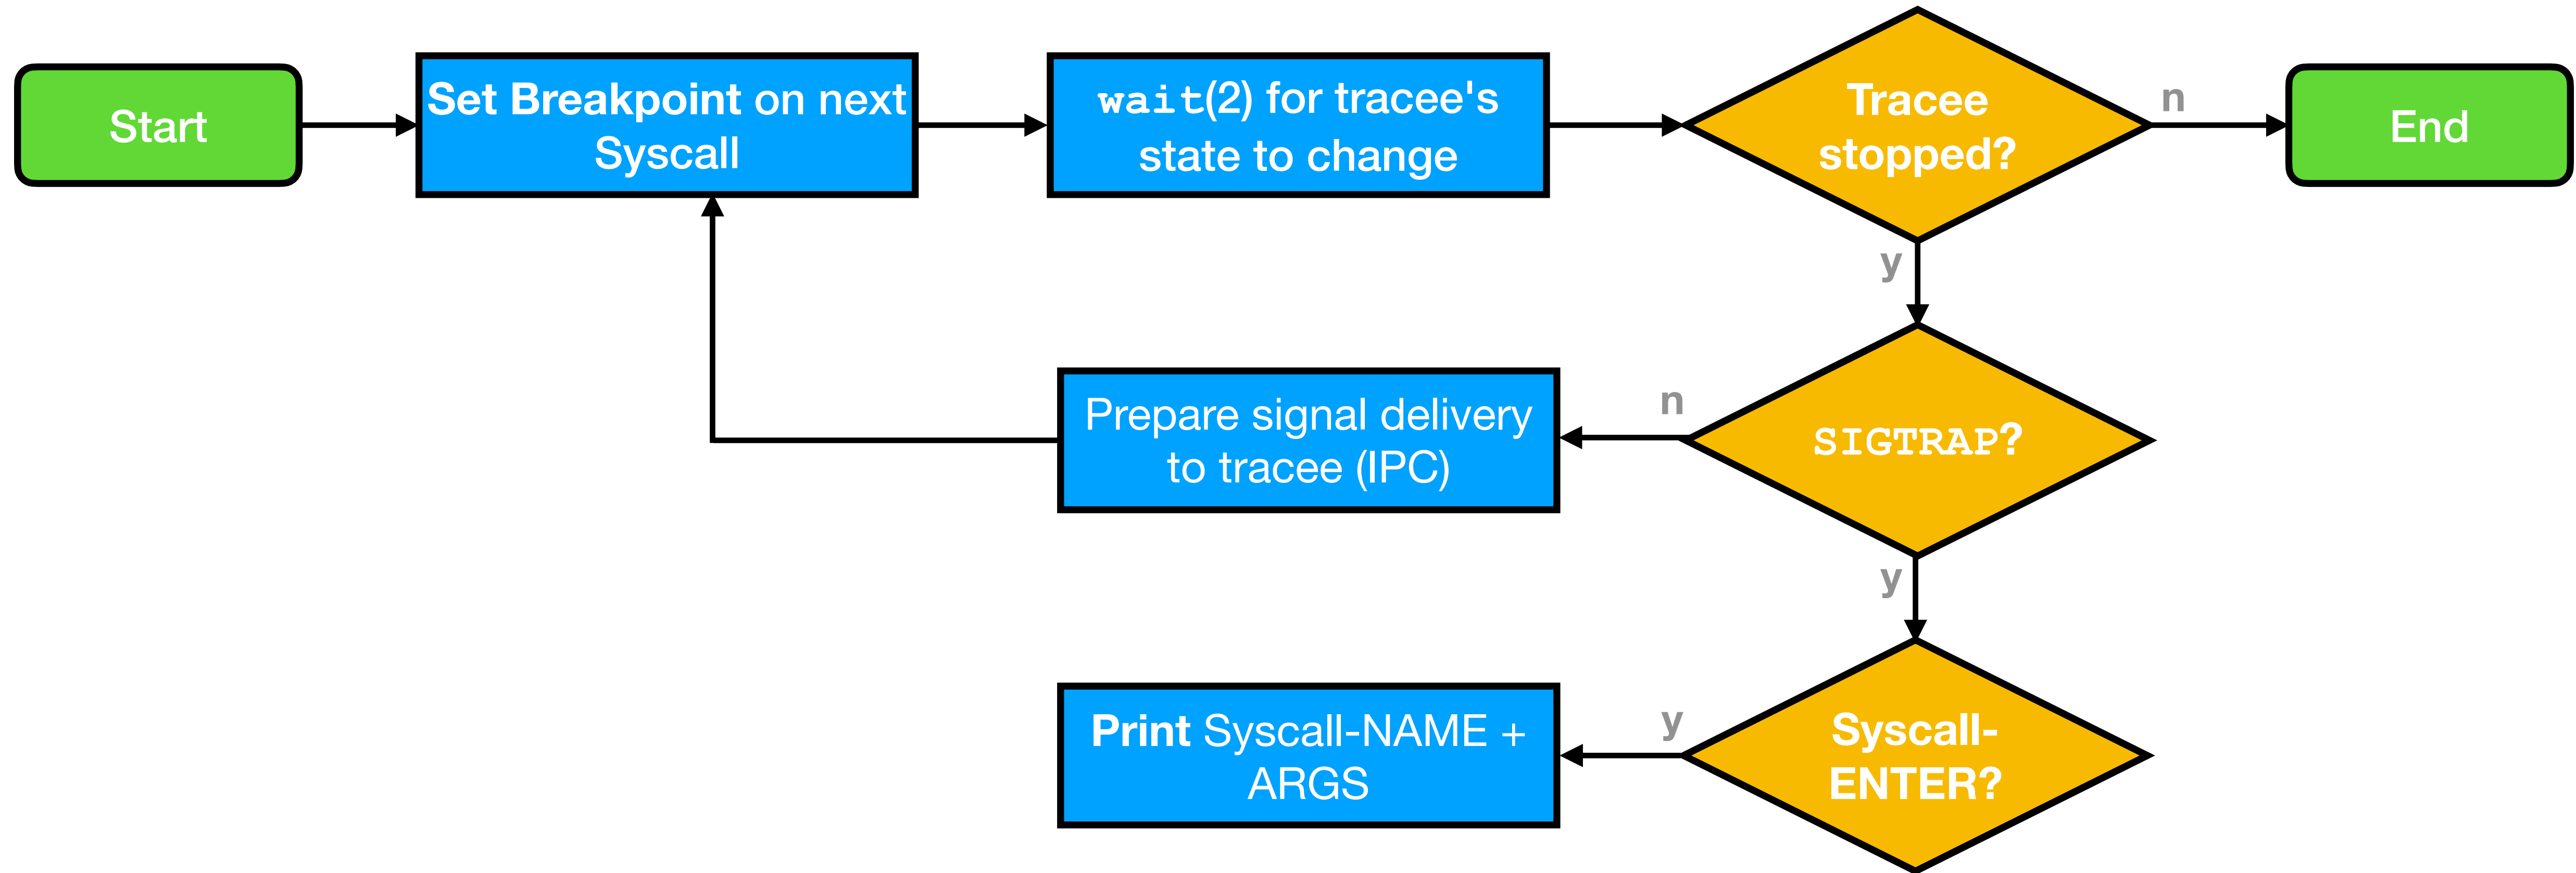
Tracing - Flow



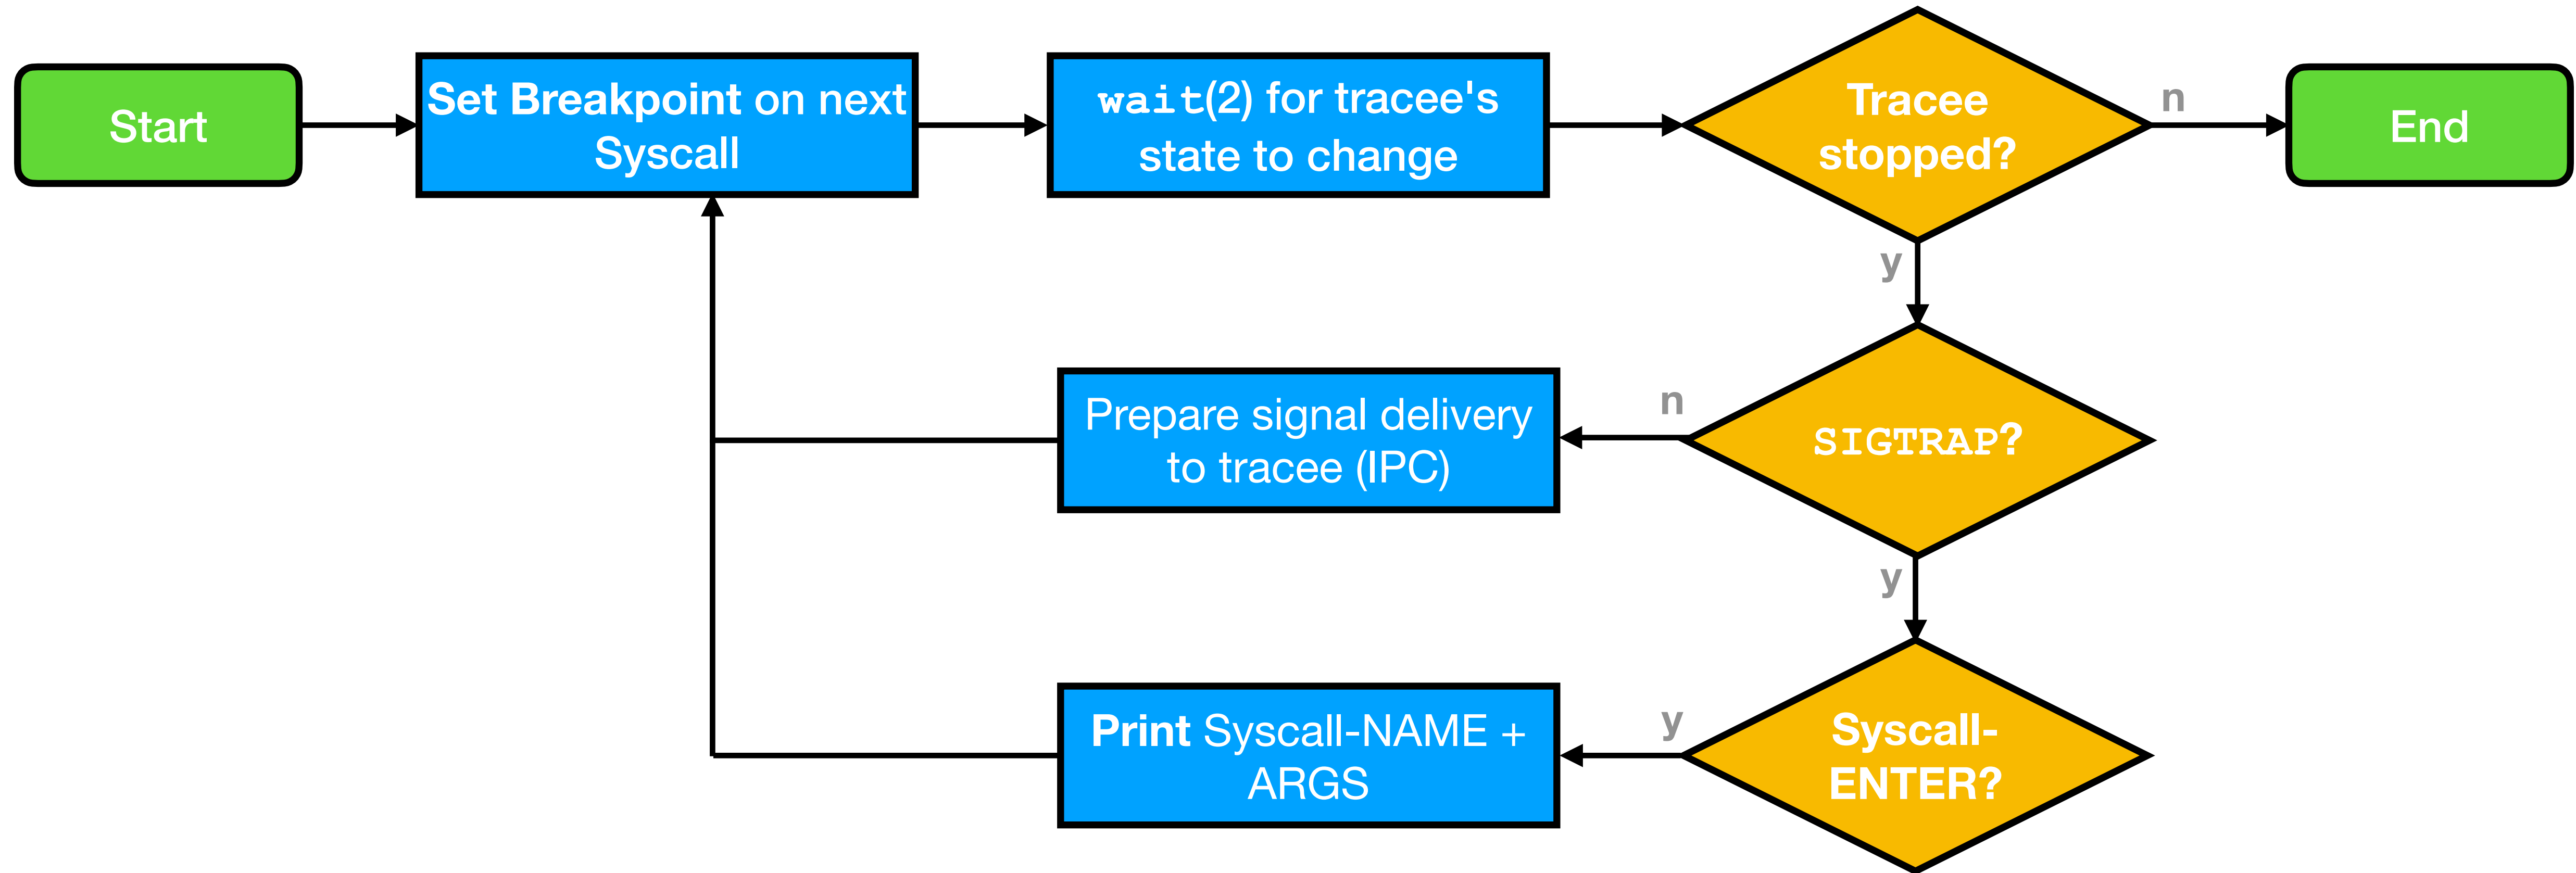
Tracing - Flow



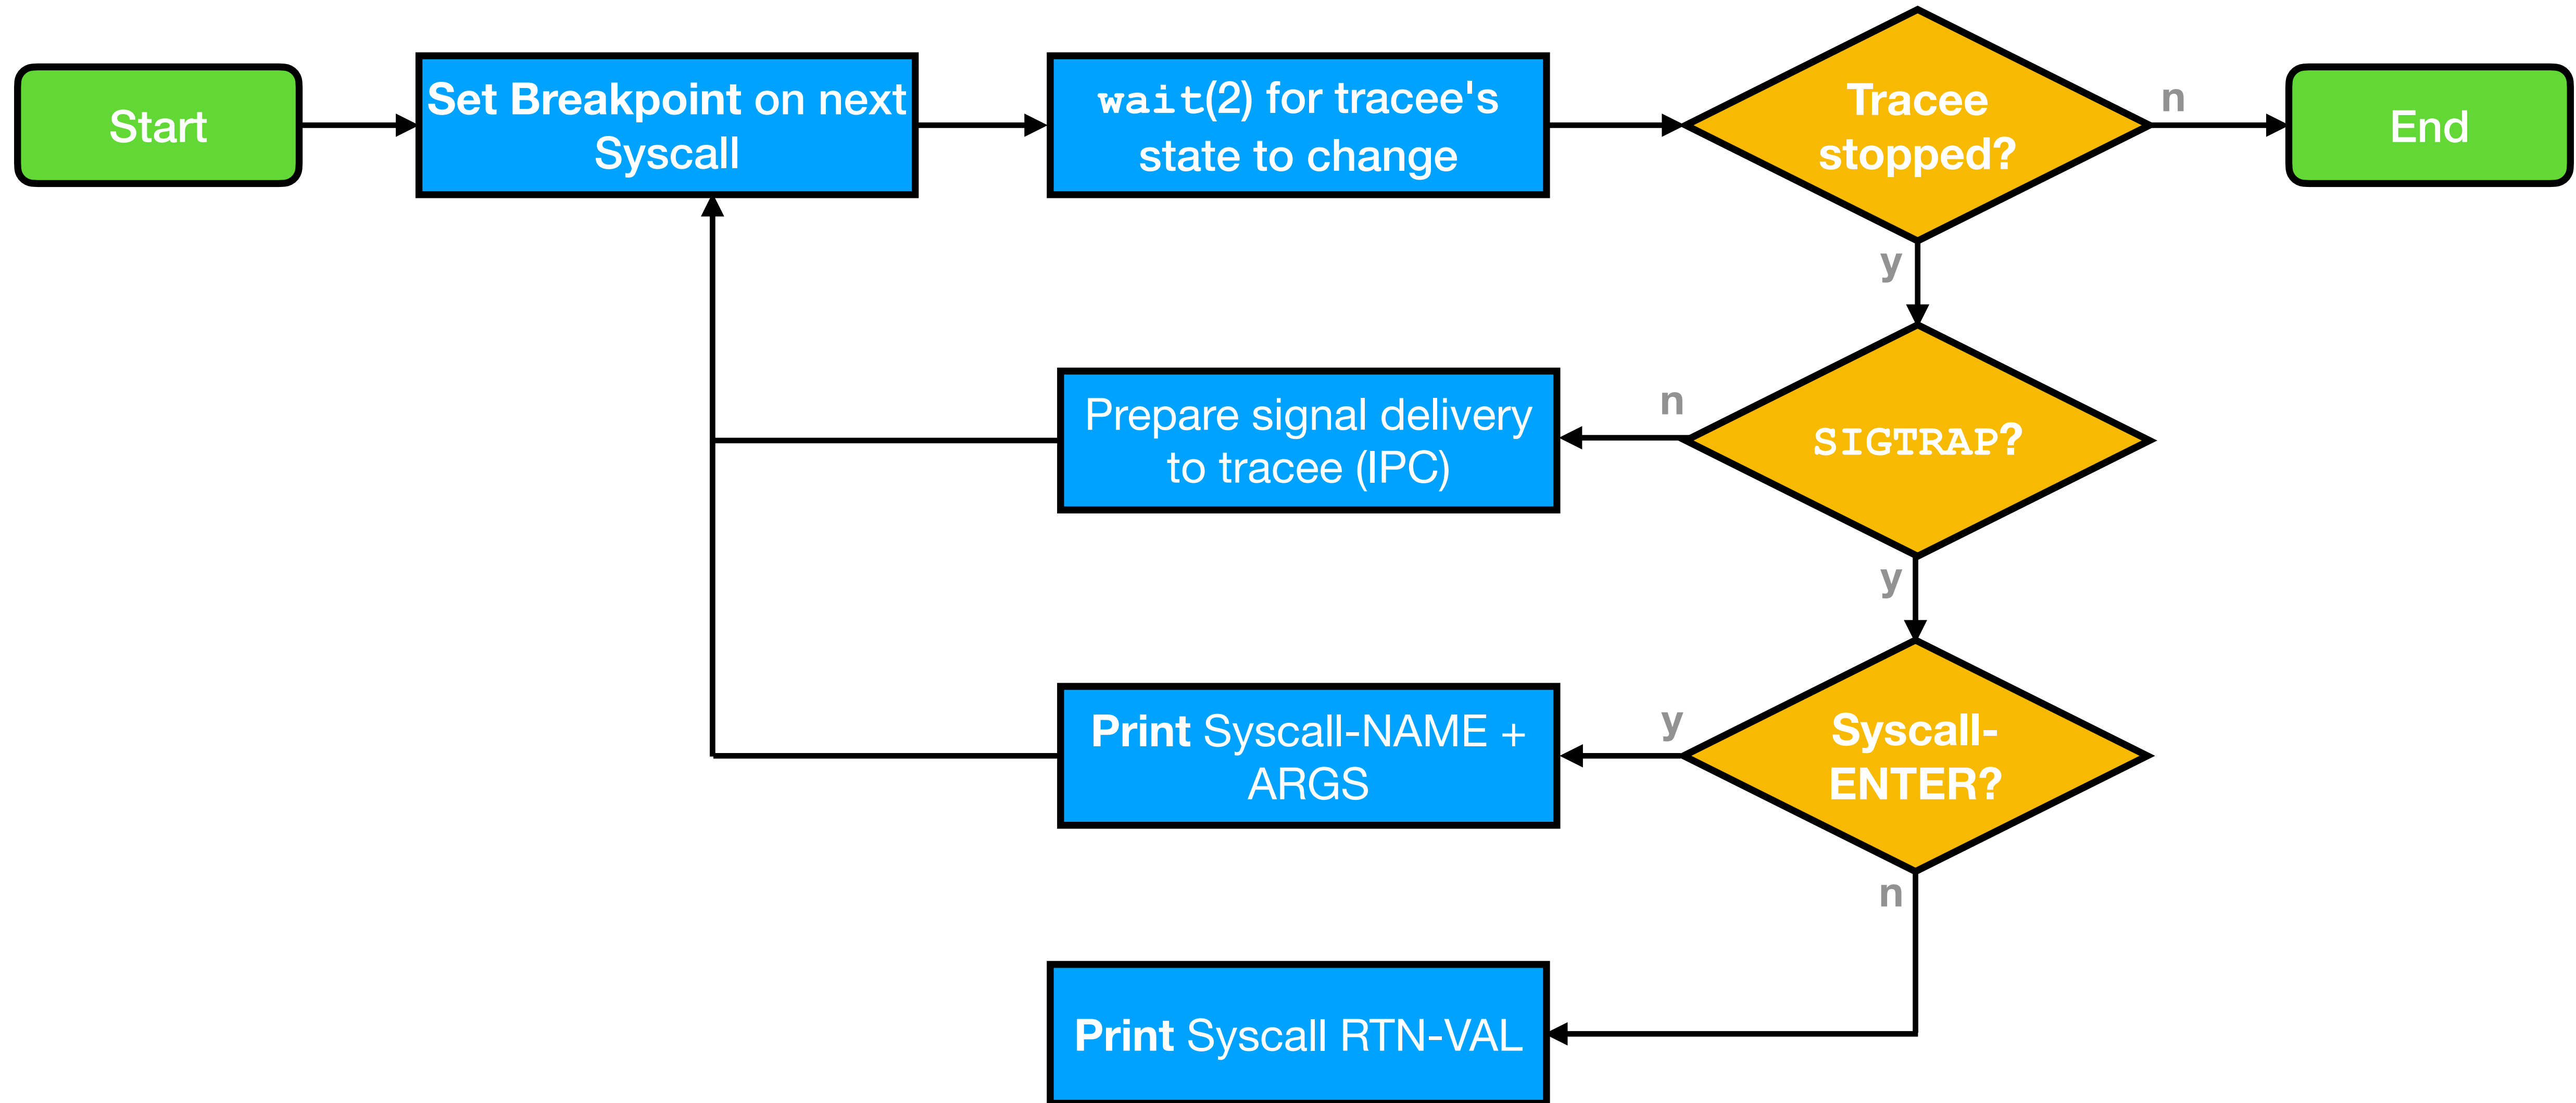
Tracing - Flow



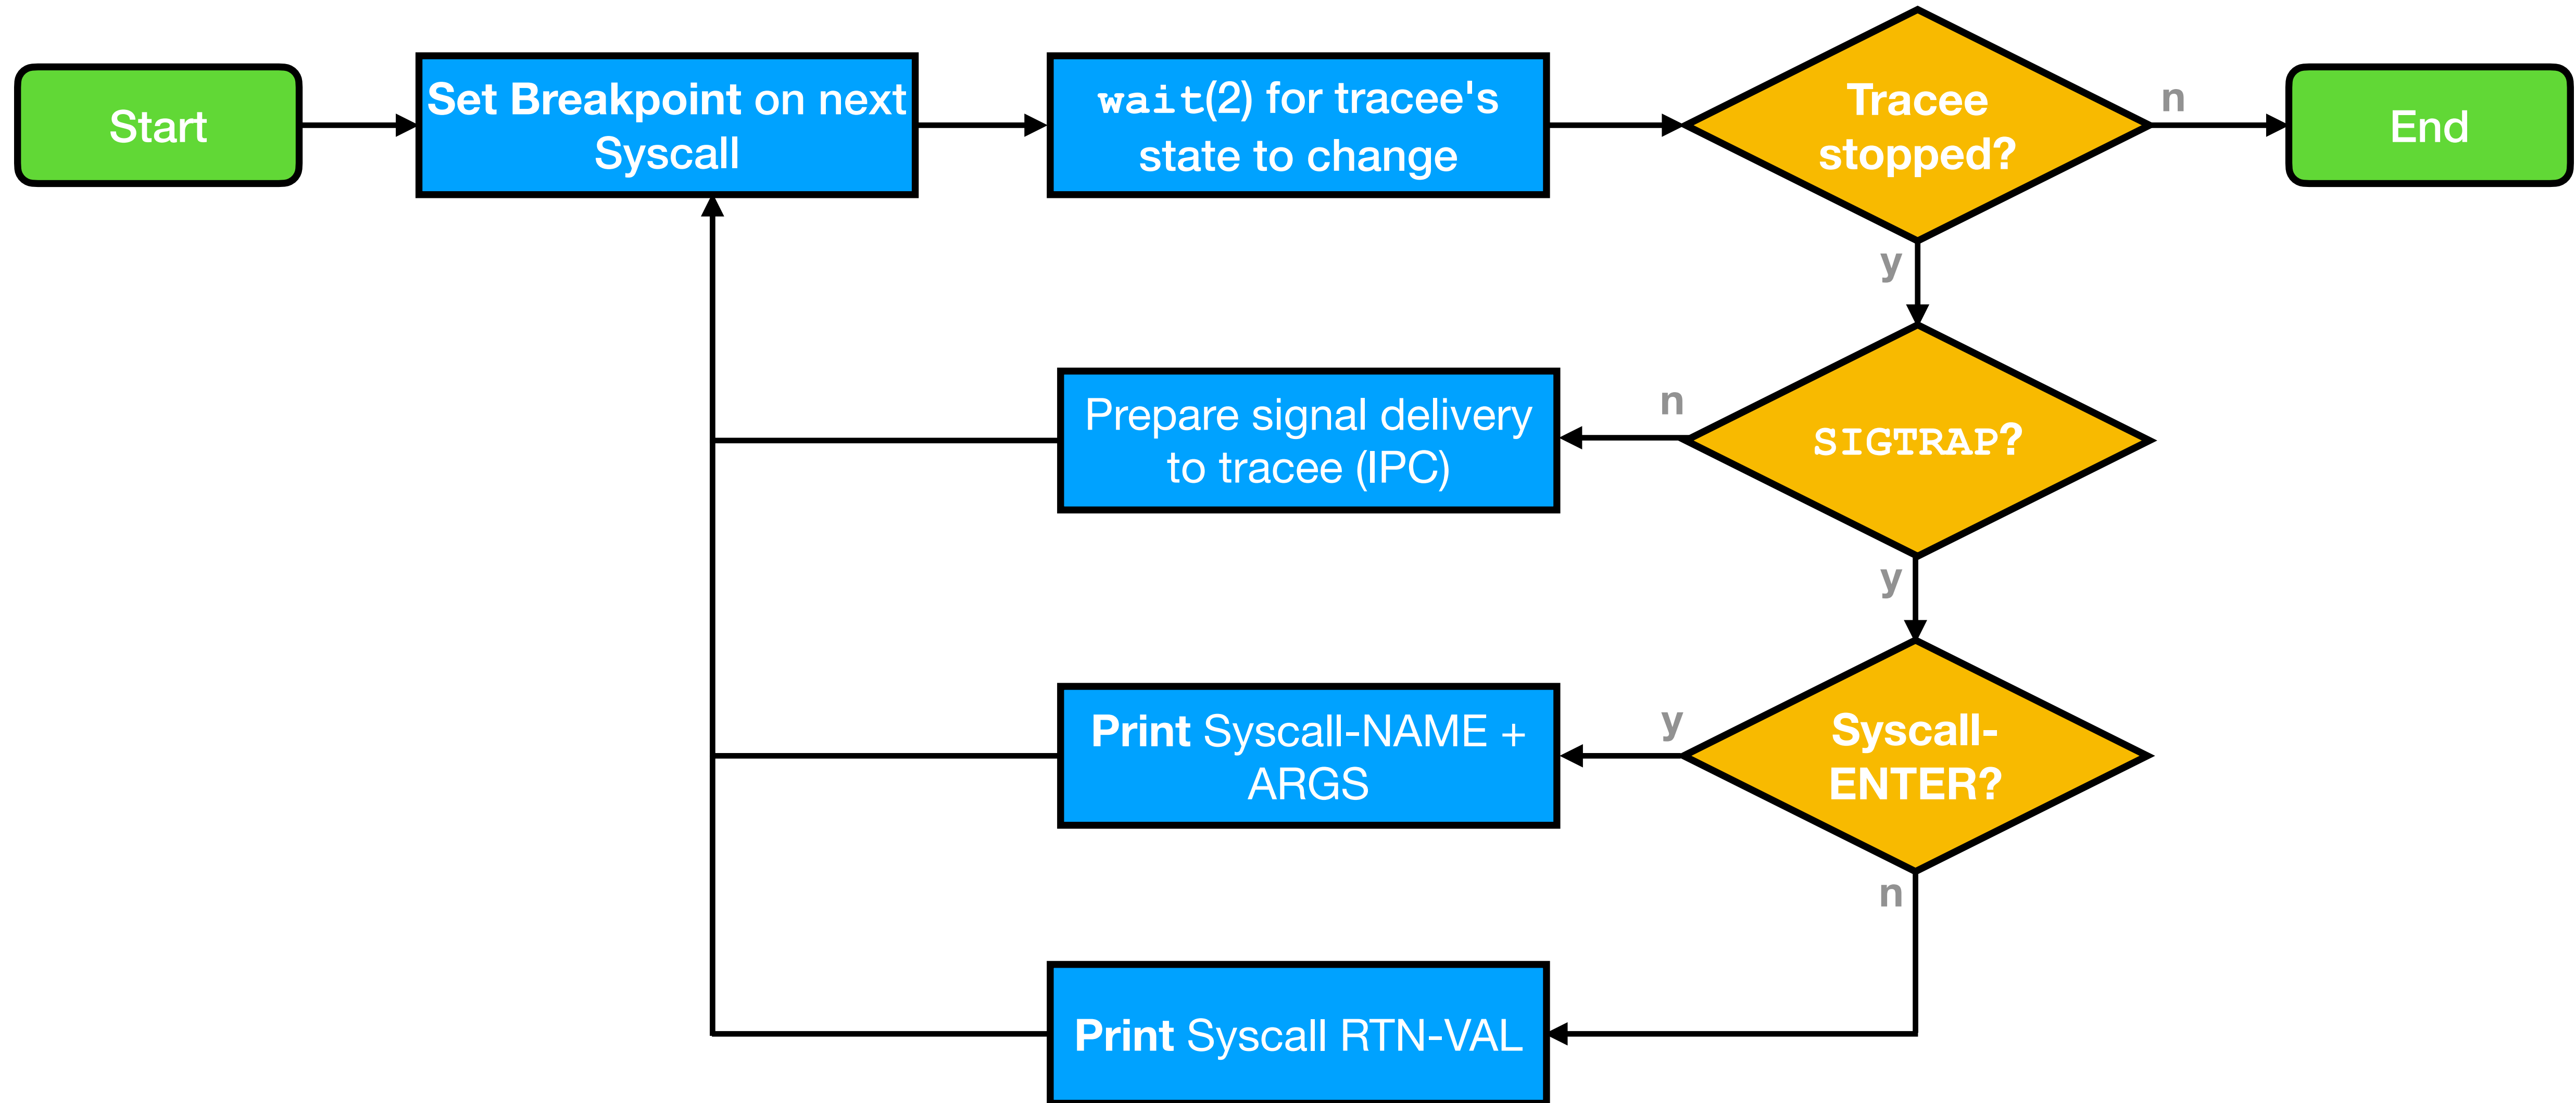
Tracing - Flow



Tracing - Flow



Tracing - Flow



Probleme

Probleme

- **aarch64 Support -> Falsche Syscall-NR**

Probleme

- **aarch64 Support -> Falsche Syscall-NR**
- **Tracing subset Syscalls**

Probleme

- **aarch64 Support** -> Falsche Syscall-NR
- **Tracing subset Syscalls**
- Workaround: `-e dup2,execve` -> Lookup table

Probleme

- **aarch64 Support** -> Falsche Syscall-NR
- **Tracing subset Syscalls**
 - Workaround: `-e dup2,execve` -> Lookup table
- **Performance**

Probleme

- **aarch64 Support** -> Falsche Syscall-NR
- **Tracing subset Syscalls**
 - Workaround: `-e dup2,execve` -> Lookup table
- **Performance**
- **OR'ed Flags** `open("test.txt", O_WRONLY | O_APPEND);`

Probleme

- **aarch64 Support** -> Falsche Syscall-NR
- **Tracing subset Syscalls**
 - Workaround: `-e dup2,execve` -> Lookup table
- **Performance**
- **OR'ed Flags**

```
open("test.txt", O_WRONLY | O_APPEND);
```

↓
1025

Todo: Integration

Ansatz

Ansatz

- **Shared VM** (Thread-like)

Ansatz

- **Shared VM** (Thread-like)
- Problem: Buffer Mutex => Deadlock

Ansatz

- **Shared VM** (Thread-like)
- Problem: Buffer Mutex => Deadlock
- Idee: Atomare Queue

Quellen

1. **How debuggers work: Part 2 - Breakpoints - Eli Bendersky's website.** (2011). Retrieved 6 February 2022, from <https://eli.thegreenplace.net/2011/01/27/how-debuggers-work-part-2-breakpoints>
2. **ptrace(2) - Linux manual page.** (2022). Retrieved 6 February 2022, from <https://man7.org/linux/man-pages/man2/ptrace.2.html>
3. **Playing with ptrace, Part I | Linux Journal.** (2002). Retrieved 6 February 2022, from <https://www.linuxjournal.com/article/6100?page=0,1>
4. **GitHub - nelhage/ministrace: A minimal toy implementation of strace(1).** (2022). Retrieved 6 February 2022, from <https://github.com/nelhage/ministrace>
5. Valsorda, F. (2022). **Searchable Linux Syscall Table for x86 and x86_64 | PyTux.** Retrieved 6 February 2022, from <https://filippo.io/linux-syscall-table/>